A New Framework for RFID Privacy

No Author Given

No Institute Given

Abstract. Formal RFID security and privacy frameworks are fundamental to the design and analysis of robust RFID systems. In this paper, we develop a new definitional framework for RFID privacy in a rigorous and precise manner. Our framework is based on a zero-knowledge (ZK) formulation [9, 7] and incorporates the notions of adaptive completeness and mutual authentication. We provide meticulous justification of the new framework and contrast it with existing ones in the literature. In particular, we prove that our framework is strictly stronger than the ind-privacy model of [17], which answers an open question posed in [17] for developing stronger RFID privacy models. We also clarify certain confusions and rectify several defects in the existing frameworks. Finally, based on the protocol of [19], we propose an efficient RFID mutual authentication protocol and analyze its security and privacy. The methodology used in our analysis can also be applied to analyze other RFID protocols within the new framework.

1 Introduction

Radio Frequency IDentification (RFID) tags are low-cost electronic devices, from which the stored information can be collected by an RFID reader efficiently (from tens to hundreds of tags per second) at a distance (from several centimeters to several meters) without the line of sight [24]. RFID technology has been widely used in numerous applications, ranging from manufacturing, logistics, transportation, warehouse inventory control, supermarket checkout counters, to many emerging applications [1]. As a key component of future ubiquitous computing environment, however, RFID technology has triggered significant concerns on its security and privacy as a tag's information can be read or traced by malicious readers from a distance without its owner's awareness [17, 14, 26, 16, 18, 6, 15].

It is critical to investigate formal RFID security and privacy frameworks that are fundamental to the design and analysis of robust RFID systems [17, 4, 25, 22, 11, 20, 19, 21]. However, due to high system complexity, it turns out to be full of subtleties in developing rigorous and precise RFID system models. By examining the existing RFID system models, in this paper we develop a new definitional framework for RFID security and privacy in a rigorous and precise manner. Our framework is based on a zero-knowledge formulation [9, 7], and incorporates the notions of adaptive completeness and mutual authentication. Compared to existing frameworks, our framework is more practical than those of [11, 19], and is stronger in terms of privacy than those of [17, 4]. Along the way, we also clarify certain confusions and rectify several defects in the existing frameworks.

To show how this new framework can be applied, we design an efficient RFID mutual authentication protocol based on the RFID protocol of [19] and analyze its security and privacy. The methodology used in our analysis is of independent interest and can be applied to analyze other RFID protocols within the new framework.

2 Preliminaries

If $A(\cdot, \cdot, ...)$ is a randomized algorithm, then $y \leftarrow A(x_1, x_2, ...; \rho)$ means that y is assigned with the unique output of the algorithm A on inputs $x_1, x_2, ...$ and coins ρ, while $y \leftarrow A(x_1, x_2, ...)$ is a shorthand for first picking ρ at random and then setting $y \leftarrow A(x_1, x_2, ...; \rho)$. Let $y \leftarrow A^{O_1, ..., O_n}(x_1, x_2, ...)$ denote that y is assigned with the output of the algorithm A which takes $x_1, x_2, ...$ as inputs and has oracle accesses to $O_1, ..., O_n$. If S is a set, then $s \in_R S$ indicates that s is chosen uniformly at random from S. If $x_1, x_2, ...$ are strings, then $x_1 || x_2 || \cdots$ denotes the concatenation of them. If x is a string, then $|x|$ denotes its bit length in binary code. If S is a set, then $|S|$ denotes its cardinality (i.e. the number of elements of S). Let $\Pr[E]$ denote the probability that an event E occurs, $\mathcal{N}$ denote the set of all integers, $\mathcal{R}$ denote the set of all real numbers.

A function $f : \mathcal{N} \rightarrow \mathcal{R}$ is said to be *negligible* if for every $c > 0$ there exits a number $m \in \mathcal{N}$ such that $f(\kappa) < \frac{1}{\kappa^c}$ holds for all $\kappa > m$. Two distribution ensembles $\{X(\kappa, z)\}_{\kappa \in N, z \in \{0,1\}^*}$ and $\{Y(\kappa, z)\}_{\kappa \in N, z \in \{0,1\}^*}$ are

computationally indistinguishable, if for any probabilistic polynomial-time (PPT) algorithm D, and for sufficiently large κ and any $z \in \{0,1\}^*$, it holds that $|\Pr[D(\kappa, z, X) = 1] - \Pr[D(\kappa, z, Y) = 1]|$ is negligible in κ.

Next, we review the definition of pseudorandom functions [8]. On a security parameter κ, let $m(\cdot)$ and $l(\cdot)$ be two positive polynomials in κ. We say that

$$\{F_k : \{0,1\}^{m(\kappa)} \longrightarrow \{0,1\}^{l(\kappa)}\}_{k \in_R \{0,1\}^\kappa}$$

is a PRF ensemble if the following two conditions hold:

1. Efficient evaluation: There exists a polynomial-time algorithm that on input k and $x \in \{0,1\}^{m(\kappa)}$ returns $F_k(x)$.
2. Pseudorandomness: A PPT oracle machine A (t, ϵ)-breaks the PRF ensemble, if

$$|\Pr[A^{F_\kappa}(\kappa) = 1] - \Pr[A^{H_\kappa}(\kappa) = 1]| \geq \epsilon$$

where F_κ is a random variable uniformly distributed over the multi-set $\{F_k\}_{k \in_R \{0,1\}^\kappa}$, H_κ is uniformly distributed among all functions mapping $m(\kappa)$-bit-long strings to $l(\kappa)$-bit-long strings, and the running time of A is at most t (here each oracle query accounts for one unit operation).

The PRF ensemble is (t, ϵ)-pseudorandom, if for all sufficiently large κ there exists no algorithm A that can (t, ϵ)-break the PRF ensemble. The PRF ensemble is pseudorandom, if for all sufficiently large κ's there exists no algorithm A that can (t, ϵ)-break the PRF ensemble, *for any t that is polynomial in κ and any ϵ that is non-negligible in κ.*

3 Model of RFID Systems

In this section, we first give a formal description of RFID system setting and adversary. We then define RFID systems to be "complete" in term of *adaptive completeness*, and "sound" in terms of *mutual authentication*.

3.1 RFID System Setting

Consider an RFID system comprising of a single[1] legitimate reader R and a set of ℓ tags $\mathcal{T} = \{\mathcal{T}_1, ..., \mathcal{T}_\ell\}$, where ℓ is a polynomial in a security parameter κ. The reader and the tags are probabilistic polynomial time interactive Turing machines. The RFID system $(R, \mathcal{T})$ is setup by a procedure, denoted $\mathsf{Setup}(\kappa, \ell)$. Specifically, on (κ, ℓ), this setup procedure generates the public system parameter σ_R, the reader secret-key k_R and initial internal state s_R^1 (if needed) for R. It may also setup an initial database DB^1 for R to store necessary information for identifying and authenticating tags. For each i, $1 \leq i \leq \ell$, this procedure generates the public parameter $\xi_{\mathcal{T}_i}$ and the initial secret-key $k_{\mathcal{T}_i}^1$ for a tag $\mathcal{T}_i$ and sets the tag's initial internal state $s_{\mathcal{T}_i}^1$ (typically, $s_{\mathcal{T}_i}^1$ includes the public parameters $\sigma_R, \xi_{\mathcal{T}_i}$). It may also associate the tag $\mathcal{T}_i$ with its unique ID, as well as other necessary information such as tag key and/or tag state information, as a record in the initial database DB^1 of R. Note that $\xi_{\mathcal{T}_i}$ or/and $s_{\mathcal{T}_i}^1$ can be empty strings.

We use $para = (\sigma_R, \xi_1, \cdots, \xi_\ell)$ to denote the public system parameters. We assume that in the RFID system, the reader is secure; in other words, the legitimate reader is a "black-box" to an adversary.

A tag $\mathcal{T}_i$, $1 \leq i \leq \ell$, exchanges messages with the reader R through a protocol $\pi(R, \mathcal{T}_i)$. Without loss of generality, we assume the protocol run of π is always initiated by R and π consists of $2\gamma + 1$ rounds[2] for some $\gamma \geq 1$. Each protocol run of π is called a session. We assume each tag interacts with the reader sequentially, but multiple tags can interact with the reader "concurrently" (with some anti-collision protocols [27]). To allow and distinguish concurrent sessions (at the side of the reader R), we associate each session of protocol π with a unique session identifier sid. In practice, sid is typically generated by the reader when it is invoked to send the first-round message. We assume each message from a tag to the reader always bears the corresponding session-identifier.

---

[1] It is straightforward to extend the model to include multiple legitimate readers. Notice that an adversary can use its own readers to interact with tags.

[2] For protocols of even 2γ rounds with the last-round message sent by the tag, we can define, by default, the $(2\gamma + 1)$-th round (from the reader to the tag) to be the output of R that indicates acceptation or rejection of the protocol run. Also, without loss of generality, we assume R and $\mathcal{T}_i$ exchange some system public parameters in the first two rounds.

Each tag $\mathcal{T}_i$, as well as the reader R, uses fresh and independent random coins (generated on the fly) in each session, *in case it is an randomized algorithm*. We assume that the random coins used in each session are erased once the session is completed (whether successfully finished or aborted). Also, in each session run, the tag may update its internal state and secret-key, and the reader may update its internal state and database. We assume that the update process of new internal state and secret-key by an uncorrupted tag automatically overwrites (i.e., erases) its old internal state and secret-key.

Given a security parameter κ, we assume that each tag $\mathcal{T}_i$ takes part in at most s (sequential) sessions in its life time[3] with R, and thus R involves at most $s\ell$ sessions, where s is some polynomial in κ. In practice, the value s can be a fixed constant (e.g., $s = 2^{28}$ [1]).

More precisely, for the j-th session (ordered by the session initiation time) where $1 \leq j \leq s\ell$, the reader R takes the input from the system parameters $para$, its secret-key k_R, current internal state s_R^j, database DB^j, random coins ρ_R^j, and a partial transcript T, where T is either an empty string (which indicates the starting of a new session) or a sequence of messages $(sid, c_1, \alpha_1, c_2, \alpha_2, \cdots, c_u, \alpha_u)$, $1 \leq u \leq \gamma$ (which indicates the on-going of session sid). The reader R outputs the next message c_{u+1}. In the case of $T = (sid, c_1, \alpha_1, c_2, \alpha_2, \cdots, c_\gamma, \alpha_\gamma)$, besides sending back the last-round message $c_{\gamma+1}$, the reader R also updates its internal state to s_R^{j+1}, its database to DB^{j+1}, and stops the session by additionally outputting a bit, denoted by o_R^{sid}. This output bit indicates either acceptance ($o_R^{sid} = 1$) or rejection ($o_R^{sid} = 0$) of the current session.

Without loss of generality, we assume that the j-th session run by the reader R corresponds to the v-th session (of session-identifier sid) run by tag $\mathcal{T}_i$, where $1 \leq v \leq s$ and $1 \leq i \leq \ell$. In this session, $\mathcal{T}_i$ takes the input from the system parameters $para$, its current secret-key $k_{\mathcal{T}_i}^v$, current internal state $s_{\mathcal{T}_i}^v$, random coins $\rho_{\mathcal{T}_i}^v$, and a partial transcript $T = (sid, c_1, \alpha_1, \cdots, \alpha_{u-1}, c_u)$, where $1 \leq u \leq \gamma$. The tag $\mathcal{T}_i$ outputs the next message (sid, α_u). In the case of $T = (sid, c_1, \alpha_1, \cdots, c_\gamma, \alpha_\gamma, c_{\gamma+1})$ (i.e., $\mathcal{T}_i$ has received the last-round message of the session sid), $\mathcal{T}_i$ updates its internal state to $s_{\mathcal{T}_i}^{v+1}$, its secret-key to $k_{\mathcal{T}_i}^{v+1}$, and stops the session by additionally outputting a bit, denoted by $o_{\mathcal{T}_i}^{sid}$. This output bit indicates either acceptance ($o_{\mathcal{T}_i}^{sid} = 1$) or rejection ($o_{\mathcal{T}_i}^{sid} = 0$) of the current session run by $\mathcal{T}_i$.

Note that in the above description, it is assumed that the reader and tags update their internal states, database, or keys *at the end of each protocol run*. In reality, this can be performed at any point of each protocol run. Also, for RFID protocol π with unidirectional authentication from tag to reader, the tag may not have a session output. In this case, the session output $o_{\mathcal{T}_i}^{sid}$ is set to "0" always.

3.2 Adversary

After an RFID system $(R, \mathcal{T})$ is setup by invoking $\mathsf{Setup}(\kappa, \ell)$, we model a probabilistic polynomial-time concurrent man-in-the-middle (CMIM) adversary $\mathcal{A}$ against $(R, \mathcal{T})$, with adaptive tag corruption. We use $\hat{m}$ to denote a message sent by adversary $\mathcal{A}$, and m to denote the actual message sent by reader R or an uncorrupted tag. The adversary is given access to the following oracles:

$\mathsf{InitReader}()$: $\mathcal{A}$ invokes R to start a session of protocol π and generate the first-round message c_1 which is also used as the session identifier sid. Supposing that the new session is the j-th session run by R, the reader R stores c_1 into its internal state s_R^j, and returns c_1 to the adversary.

$\mathsf{SendT}(\mathcal{T}_i, \hat{m})$: Adversary $\mathcal{A}$ sends $\hat{m}$ to $\mathcal{T}_i$[4]. After receiving $\hat{m}$, $\mathcal{T}_i$ works as follows: (1) If $\mathcal{T}_i$ currently does not run any existing session, $\mathcal{T}_i$ initiates a new session with the session-identifier sid set to $\hat{m}$, treats $\hat{m}$ as the first-round message of the new session, and returns the second-round message (sid, α_1). (2) If $\mathcal{T}_i$ is currently running an incomplete session with session-identifier $sid = \hat{c}$, and is waiting for the u-th message from R, where $u \geq 2$, $\mathcal{T}_i$ works as follows: If $2 \leq u \leq \gamma$, it treats $\hat{m}$ as the u-th message from the reader and returns the next round message (sid, α_u). If $u = \gamma + 1$ (i.e., $\mathcal{T}_i$ is waiting for the last-round message of the session sid), $\mathcal{T}_i$ returns its output $o_{\mathcal{T}_i}^{sid}$ to the adversary, and (internally) updates its internal state to $s_{\mathcal{T}_i}^{v+1}$, assuming that the session sid is the v-th session run by $\mathcal{T}_i$, where $1 \leq v \leq s$.

---

[3] It is assumed that s is large enough so that any tag can never run up to s sessions in its life time; otherwise, an adversary may distinguish two tags, thus violate their privacy, by running one tag more than s times while the other less than s times [17].

[4] For simplicity, we abuse the notation $\mathcal{T}_i$ to denote any virtual identity of a tag in $\mathcal{T}$ (not the tag's real identity) labeled by $\mathcal{A}$ when $\mathcal{A}$ selects the tag from $\mathcal{T}$.

SendR$(\widehat{sid}, \hat{\alpha})$: Adversary $\mathcal{A}$ sends $(\widehat{sid}, \hat{\alpha})$ to R. After receiving $(\widehat{sid}, \hat{\alpha})$, R checks from its internal state whether it is running a session of session identifier $sid = \widehat{sid}$, and works as follows: (1) If R is currently running an incomplete session with $sid = \widehat{sid}$ and is waiting for the u-th message from a tag, where $1 \leq u \leq \gamma$, R acts as follows: If $u < \gamma$, it treats $\hat{\alpha}$ as the u-th message from the tag, and returns the next round message c_{u+1} to $\mathcal{A}$. If $u = \gamma$, it returns the last-round message $c_{\gamma+1}$ and the output o_R^{sid} to $\mathcal{A}$, and internally updates its internal state to s_R^{j+1} and the database to DB^{j+1}, assuming that the session sid corresponds to the j-th session run by R. (2) In all other cases, R returns a special symbol $\perp$ (indicating invalid query).

Corrupt$(\mathcal{T}_i)$: Adversary $\mathcal{A}$ obtains the secret-key and internal state information (as well as the random coins) currently held by $\mathcal{T}_i$. Once a tag $\mathcal{T}_i$ is corrupted, all its actions are controlled and performed by the adversary $\mathcal{A}$.

Let O_1, O_2, O_3 and O_4 denote the above oracles, respectively. These oracles fully capture the capability of any PPT CMIM adversary with adaptive tag corruption.[5] For presentation simplicity, we denote by $\mathcal{O}$ the set of the four oracles $\{O_1, O_2, O_3, O_4\}$ specified above. *An adversary is a (t, n_1, n_2, n_3, n_4)-adversary, if it works in time t and makes oracle queries to O_μ without exceeding n_μ times, where $1 \leq \mu \leq 4$.* We treat each oracle call as a unit operation, and thus for a t-time adversary it holds that $\Sigma_{\mu=1}^4 n_\mu \leq t$. We denote by $A^{\mathcal{O}}(R, \mathcal{T}, para)$ a PPT algorithm A that, on input of some system public parameter $para$, concurrently interacts with R and the tags in $\mathcal{T}$ via the four oracles in $\mathcal{O}$, where $(R, \mathcal{T})$ is setup by Setup(κ, ℓ).

Note that in our formulation, the output bits of protocol participants (which indicate authentication success or failure) are *publicly* accessible to the adversary. The reason is that, in reality, such outputs can be publicly observed from the behaviors of protocol participants during/after the protocol run or can be learnt by some other side channels.

3.3 Adaptive Completeness and Mutual Authentication

Roughly speaking, adaptive completeness says that, after any attacks (*particularly the desynchronizing attacks*) made by the adversary $\mathcal{A}$, the protocol execution between the reader R and any honest uncorrupted tag is still complete (e.g., being able to recover from desynchronization). In other words, after undergoing arbitrary attacks, the uncorrupted parties of the RFID system still can recover *whenever the attacks stop*.

Definition 3.1 (adaptive completeness) *For an RFID system $(R, \mathcal{T})$ setup by Setup(κ, ℓ), denote by*

$$(sid, c_1^{sid}, \alpha_1^{sid}, \cdots, \alpha_\gamma^{sid}, c_{\gamma+1}^{sid}, o_R^{sid}, o_{\mathcal{T}_i}^{sid}) \leftarrow \pi(R, \mathcal{T}_i)$$

the running of a session with identifier sid of the protocol π between R and an uncorrupted tag $\mathcal{T}_i \in \mathcal{T}$. Suppose that the session sid corresponds to the v-th session at the side of $\mathcal{T}_i$ and the j-th session at the side of R, where $1 \leq v \leq s$ and $1 \leq j \leq s\ell$. Consider the case that the two sessions are of the same round messages, and that all the exchanged messages in these two (matching) sessions are all honestly generated by R and $\mathcal{T}_i$ respectively. Denote by E the event that $o_R^{sid} = 0$ holds (or $o_{\mathcal{T}_i}^{sid} = 0$ holds if the protocol π is for mutual authentication) or R identifies a different tag $\mathcal{T}_{i'} \neq \mathcal{T}_i$ in its j-th session.

A PPT CMIM adversary $\mathcal{A}$ $(t, \epsilon, n_1, n_2, n_3, n_4)$-breaks the adaptive completeness of the RFID system against the uncorrupted $\mathcal{T}_i$, if the probability that event E occurs is at least ϵ and $\mathcal{A}$ is a (t, n_1, n_2, n_3, n_4)-adversary. The probability is taken over the coins used by Setup(κ, ℓ), the coins of $\mathcal{A}$, the coins used by R (up to finishing the j-th session), and the coins used by $\mathcal{T}_i$ (up to finishing the v-th session). An RFID system $(R, \mathcal{T})$ satisfies adaptive completeness, if for all sufficiently large κ and for any uncorrupted tag $\mathcal{T}_i$, there exists no adversary $\mathcal{A}$ that can $(t, \epsilon, n_1, n_2, n_3, n_4)$-break the adaptive completeness against $\mathcal{T}_i$, for any (t, ϵ), where t is polynomial in κ and ϵ is non-negligible in κ.

Next, we define mutual authentication of RFID protocols. Roughly speaking, for a protocol π of the RFID system $(R, \mathcal{T})$, authentication from reader to tag (resp., from tag to reader) means that a CMIM adversary $\mathcal{A}$ cannot impersonate the reader R (resp., an uncorrupted tag $\mathcal{T}_i \in \mathcal{T}$) to an uncorrupted tag $\mathcal{T}_i \in \mathcal{T}$ (resp., reader R), unless $\mathcal{A}$ honestly relays messages actually generated and sent by R and the uncorrupted tag $\mathcal{T}_i$. Before we define mutual authentication for RFID protocols, we first clarify the notion of matching sessions.

---

[5] Here, for simpler definitional complexity, we assume all tags are always within the attack scope of adversary. In practice, some tags may be in or out from the attack scope of adversary at different time [25].

Definition 3.2 (matching sessions) *Denote by* $(sid, c_1^{sid}, \alpha_1^{sid}, \cdots, \alpha_\gamma^{sid}, c_{\gamma+1}^{sid})$ *the transcript of exchanged round messages (except the session outputs) of a successfully completed session sid of the protocol π run by a tag $\mathcal{T}_i$, where $1 \le i \le \ell$. This session has a matching session at the side of the reader R, if R ever successfully completed a session of the identical session transcript.*

Denote by $(sid', c_1^{sid'}, \alpha_1^{sid'}, \cdots, \alpha_\gamma^{sid'}, c_{\gamma+1}^{sid'})$ *the transcript of exchanged round messages (except the session outputs) of a successfully completed session sid' run by R. This session has a matching session at the side of some tag $\mathcal{T}_i$, where $1 \le i \le \ell$, if either of the following conditions holds:*

- *$\mathcal{T}_i$ ever completed, whether successfully finished or aborted, a session of the identical transcript prefix $(sid', c_1^{sid'}, \alpha_1^{sid'}, \cdots, \alpha_\gamma^{sid'})$;*
- *Or, $\mathcal{T}_i$ is now running a session with partial transcript $(sid', c_1^{sid'}, \alpha_1^{sid'}, \cdots, \alpha_\gamma^{sid'})$ and is waiting for the last-round message of the session sid'.*

The matching-session definition, for a successfully completed session run by the reader R, takes into account the following "cutting-last-message" attack: a CMIM adversary $\mathcal{A}$ relays the messages being exchanged by R and an uncorrupted tag $\mathcal{T}_i$ for a protocol run of π until receiving the last-round message $c_{\gamma+1}^{sid'}$ from R; after this, $\mathcal{A}$ sends an arbitrary message $\hat{c}_{\gamma+1}^{sid'} (\ne c_{\gamma+1}^{sid'})$ to $\mathcal{T}_i$ (which typically causes $\mathcal{T}_i$ to abort the session), or, just drops the session at the side of $\mathcal{T}_i$ without sending $\mathcal{T}_i$ the last-round message. Such "cutting-last-message" attacks are unpreventable.

Figure 1 shows the authentication experiment $\mathbf{Exp}_{\mathcal{A}}^{auth}[\kappa, \ell]$. A CMIM adversary $\mathcal{A}$ interacts with R and tags in $\mathcal{T}$ via the four oracles in $\mathcal{O}$; At the end of the experiment, $\mathcal{A}$ outputs the transcript, $trans$, of a session. Denote by E_1 the event that $trans$ corresponds to the transcript of a successfully completed session run by R in which R successfully identifies an *uncorrupted* tag $\mathcal{T}_i$, but this session has no matching session at the side of the uncorrupted tag $\mathcal{T}_i$. Denote by E_2 the event that $trans$ corresponds to the transcript of a successfully completed session run by some *uncorrupted* tag $\mathcal{T}_i \in \mathcal{T}$, and this session has no matching session at the side of R.

Experiment $\mathbf{Exp}_{\mathcal{A}}^{auth}[\kappa, \ell]$
  1. run $\mathsf{Setup}(\kappa, \ell)$ to setup the reader R and a set of tags $\mathcal{T}$;
      denote by $para$ the public system parameters;
  2. $trans \leftarrow \mathcal{A}^{\mathcal{O}}(R, \mathcal{T}, para)$.

Fig. 1. Authentication Experiment

Definition 3.3 (authentication) *On a security parameter κ, an adversary $\mathcal{A}$ $(\epsilon, t, n_1, n_2, n_3, n_4)$-breaks the authentication of an RFID system $(R, \mathcal{T})$ against the reader R (resp., an uncorrupted tag $\mathcal{T}_i \in \mathcal{T}$) if the probability that event E_1 (resp., E_2) occurs is at least ϵ and $\mathcal{A}$ is a (t, n_1, n_2, n_3, n_4)-adversary.*

The RFID system $(R, \mathcal{T})$ satisfies tag-to-reader authentication (resp., reader-to-tag authentication), if for all sufficiently large κ there exists no adversary $\mathcal{A}$ that can $(\epsilon, t, n_1, n_2, n_3, n_4)$-break the authentication of $(R, \mathcal{T})$ against the reader R (resp., any uncorrupted tag $\mathcal{T}_i \in \mathcal{T}$), for any (t, ϵ), where t is polynomial in κ and ϵ is non-negligible in κ. An RFID system is of mutual authentication, if it satisfies both tag-to-reader authentication and reader-to-tag authentication.

Adaptive completeness vs. (mutual) authentication. Adaptive completeness essentially says that, after any attacks (*particularly the desynchronizing attacks*) made by the adversary $\mathcal{A}$, the protocol execution between the reader R and any honest uncorrupted tag is still complete (e.g., being able to recover from desynchronization). In other words, after undergoing arbitrary attacks, the uncorrupted parties of the RFID system still can recover *whenever the attacks stop*.

We highlight some formulation differences between adaptive completeness and (mutual) authentication. On the one hand, adaptive completeness is formulated w.r.t. the session transcript between R and an honest uncorrupted tag $\mathcal{T}_i$ (that corresponds to the j-th (resp., v-th) session at the side of R (resp., $\mathcal{T}_i$)), while in the authentication experiment the session transcript (*output by the adversary $\mathcal{A}$*) is typically between the adversary $\mathcal{A}$ and the reader R or an honest tag; On the other hand, in the definition of adaptive completeness the probability is taken over the coins of R *up to finishing the j-th session* and the coins of $\mathcal{T}_i$ *up to finishing the v-th session*, while in the definition of (mutual) authentication the probability is taken over the coins of R and tags in all sessions (besides the coins used by $\mathcal{A}$ and $\mathsf{Setup}(\kappa, \ell)$).

4 Zero-Knowledge Based RFID Privacy

In this section, we present a zero-knowledge based definitional framework for RFID privacy. To make our definition formal, we need to clarify the notion of blind access to tags and the notion of clean tags.

Let $A^{\mathcal{O}}(R, \widehat{\mathcal{T}}, \mathcal{I}(\mathcal{T}_g), aux)$ be a PPT algorithm A that, on input $aux \in \{0,1\}^*$ (typically, aux includes the system parameters or some historical state information of A), concurrently interacts with R and a set of tags $\widehat{\mathcal{T}}$ via the four oracles $\mathcal{O} = \{O_1, O_2, O_3, O_4\}$. We say that A has *blind access* to a *challenge* tag $\mathcal{T}_g \notin \widehat{T}$ if A interacts with $\mathcal{T}_g$ via a special interface $\mathcal{I}$. Specifically, $\mathcal{I}$ is a PPT algorithm that runs $\mathcal{T}_g$ internally, and interacts with A externally. To send a message $\hat{c}$ to $\mathcal{T}_g$, A sends to $\mathcal{I}$ a special O_2 oracle query of the form $\mathsf{SendT}(challenge, \hat{c})$; after receiving this special O_2 query, $\mathcal{I}$ invokes $\mathcal{T}_g$ with $\mathsf{SendT}(\mathcal{T}_g, \hat{c})$, and returns to A the output by $\mathcal{T}_g$. From the viewpoint of A, it does not know which tag it is interacting with. It is also required that A interacts with $\mathcal{T}_g$ via O_2 queries only.

Next, we define the notion of clean tags. A tag $\mathcal{T}_i$ is called *clean*, if it is not corrupted (i.e., the adversary has not made any O_4 query to $\mathcal{T}_i$), and is not currently running an incomplete session with the reader (i.e., the last session of the tag has been either finished or aborted). In other words, a clean tag is an uncorrupted tag that is currently at the status of waiting for the first-round message from the reader to start a new session.

Now, we are ready to give a formal definition of zero-knowledge based RFID privacy (zk-privacy, for short). Figure 2 illustrates the real world of the zk-privacy experiment, $\mathbf{Exp}_{\mathcal{A}}^{zkp}[\kappa, \ell]$ ($\mathbf{Exp}_{\mathcal{A}}^{zkp}$, for simplicity), in which a PPT CMIM adversary $\mathcal{A}$ is comprised of a pair of algorithms $(\mathcal{A}_1, \mathcal{A}_2)$ and runs in two stages. In the *first stage*, algorithm $\mathcal{A}_1$ is concurrently interacting with R and all the tags in $\mathcal{T}$ via the four oracles in $\mathcal{O}$, and is required to output a set $\mathcal{C}$ of *clean* tags at the end of the first stage, where $\mathcal{C} \subseteq \mathcal{T}$ consists of δ *clean* tags, denoted as $\{\mathcal{T}_{i_1}, \cdots, \mathcal{T}_{i_\delta}\}$. The algorithm $\mathcal{A}_1$ also outputs a state information st, which will be transmitted to algorithm $\mathcal{A}_2$. Between the first stage and the second stage, a challenge tag, denoted as $\mathcal{T}_g$, is taken uniformly at random from $\mathcal{C}$. Note that if $\delta = 0$, then no challenge tag is selected, and $\mathcal{A}$ is reduced to $\mathcal{A}_1$ in this experiment. In the *second stage*, on input st, $\mathcal{A}_2$ concurrently interacts with the reader R and the tags in $\widehat{\mathcal{T}} = \mathcal{T} - \mathcal{C}$ via the four oracles in $\mathcal{O}$, and additionally has blind access to $\mathcal{T}_g$. Note that $\mathcal{A}$ cannot corrupt any tag (particularly $\mathcal{T}_g$) in $\mathcal{C}$, and $\mathcal{A}$ does not have access to tags in $\mathcal{C} - \{\mathcal{T}_g\}$ in the second stage. Finally, $\mathcal{A}_2$ outputs its view, denoted by $view_{\mathcal{A}}$, at the end of the second stage. Specifically, $view_{\mathcal{A}}$ is defined to include the system public parameters $para$, the random coins used by $\mathcal{A}$, $\rho_{\mathcal{A}}$, and the (ordered) list of all oracle answers to the queries made by $\mathcal{A}$ in the experiment $\mathbf{Exp}_{\mathcal{A}}^{zkp}$. Note that $view_{\mathcal{A}}$ does not explicitly include the oracle queries made by $\mathcal{A}$ and $\mathcal{A}$'s output at the first stage, as all these values are implicitly determined by the system public parameter $para$, $\mathcal{A}$'s coins and all oracle answers to $\mathcal{A}$'s queries. The output of experiment $\mathbf{Exp}_{\mathcal{A}}^{zkp}$ is defined to be $(g, view_{\mathcal{A}})$. Denote by $(g, view_{\mathcal{A}}(\kappa, \ell))$ the random variable describing the output of experiment $\mathbf{Exp}_{\mathcal{A}}^{zkp}[\kappa, \ell]$.

Experiment $\mathbf{Exp}_{\mathcal{A}}^{zkp}[\kappa, \ell]$
 1. run $\mathsf{Setup}(\kappa, \ell)$ to setup the reader R and a set of tags $\mathcal{T}$; denote by $para$ the public system parameter;
 2. $\{\mathcal{C}, st\} \leftarrow \mathcal{A}_1^{\mathcal{O}}(R, \mathcal{T}, para)$, where $\mathcal{C} = \{\mathcal{T}_{i_1}, \mathcal{T}_{i_2}, \cdots, \mathcal{T}_{i_\delta}\} \subseteq \mathcal{T}$ is a set of *clean* tags, $0 \leq \delta \leq \ell$;
 3. $g \in_R \{1, \cdots, \delta\}$, set $\mathcal{T}_g = \mathcal{T}_{i_g}$ and $\widehat{\mathcal{T}} = \mathcal{T} - \mathcal{C}$;
 4. $view_{\mathcal{A}} \leftarrow \mathcal{A}_2^{\mathcal{O}}(R, \widehat{\mathcal{T}}, \mathcal{I}(\mathcal{T}_g), st)$;
 5. output $(g, view_{\mathcal{A}})$.

Fig. 2. zk-privacy experiment: real world

Figure 3 illustrates the simulated world of zk-privacy experiment, $\mathbf{Exp}_{\mathcal{S}}^{zkp}[\kappa, \ell]$ ($\mathbf{Exp}_{\mathcal{S}}^{zkp}$, for simplicity), in which a PPT simulator $\mathcal{S}$ is comprised of a pair of algorithms $(\mathcal{S}_1, \mathcal{S}_2)$ and runs in two stages. In the *first stage*, algorithm $\mathcal{S}_1$ concurrently interacts with R and all the tags in $\mathcal{T}$ via the four oracles in $\mathcal{O}$, and outputs a set, denoted $\mathcal{C}$, of *clean* tags, where $|\mathcal{C}| = \delta$ and $0 \leq \delta \leq \ell$. It also outputs a state information st, which will be transmitted to algorithm $\mathcal{S}_2$. Between the two stages, a value g is taken uniformly at random from $\{1, \cdots, |\mathcal{C}|\}$ (which is unknown to $\mathcal{S}$). In the *second stage* of $\mathcal{S}$, on input st, $\mathcal{S}_2$ concurrently interacts with the reader R and the tags in $\widehat{\mathcal{T}} = \mathcal{T} - \mathcal{C}$, and outputs a

Experiment $\mathbf{Exp}_{\mathcal{S}}^{zkp}[\kappa, \ell]$

   1. run $\mathsf{Setup}(\kappa, \ell)$ to setup the reader R and a set of tags $\mathcal{T}$; denote by $para$ the public system parameter;

   2. $\{\mathcal{C}, st\} \leftarrow \mathcal{S}_1^{\mathcal{O}}(R, \mathcal{T}, para)$, where $\mathcal{C} = \{T_{i_1}, T_{i_2}, \cdots, T_{i_\delta}\} \subseteq \mathcal{T}$ is a set of *clean* tags, $0 \le \delta \le \ell$;

   3. $g \in_R \{1, \cdots, \delta\}$, and set $\widehat{\mathcal{T}} = \mathcal{T} - \mathcal{C}$;

   4. $sview \leftarrow \mathcal{S}_2^{\mathcal{O}}(R, \widehat{\mathcal{T}}, st)$, where $sview$ particularly includes all oracle answers to queries made by $\mathcal{S}$;

   5. output $(g, sview)$.

Fig. 3. zk-privacy experiment: simulated world

simulated view, denoted $sview$, at the end of the second stage. We require that all oracle answers to the queries made by $\mathcal{S}$ (in both the first stage and the second stage) in the experiment $\mathbf{Exp}_{\mathcal{S}}^{zkp}$ are included in $sview$. The output of the experiment $\mathbf{Exp}_{\mathcal{S}}^{zkp}$ is defined to be $(g, sview)$. Denote by $(g, sview(\kappa, \ell))$ the random variable describing the output of the experiment $\mathbf{Exp}_{\mathcal{S}}^{zkp}[\kappa, \ell]$.

Informally, an RFID protocol π satisfies zk-privacy, if what can be derived by interacting with the challenge tag $\mathcal{T}_g$ in the second-stage of $\mathcal{A}$ can actually be derived by $\mathcal{A}$ itself *without interacting with* $\mathcal{T}_g$. In this sense, the interaction between $\mathcal{A}_2$ and $\mathcal{T}_g$ leaks "zero knowledge" to $\mathcal{A}$. For this reason, our RFID privacy notion is named zk-privacy.

Definition 4.1 (zk-privacy) *An RFID protocol π satisfies computational (resp., statistical) zk-privacy, if for any PPT CMIM adversary $\mathcal{A}$ there exists a polynomial-time simulator $\mathcal{S}$ such that for all sufficiently large κ and any ℓ which is polynomials in κ (i.e., $\ell = poly(\kappa)$, where $poly(\cdot)$ is some positive polynomial), the following ensembles are computationally (resp., statistically) indistinguishable:*

- *$\{g, view_{\mathcal{A}}(\kappa, \ell)\}_{\kappa \in N, \ell \in poly(\kappa)}$*
- *$\{g, sview(\kappa, \ell)\}_{\kappa \in N, \ell \in poly(\kappa)}$*

That is, for any polynomial-time (resp., any power unlimited) algorithm D, it holds that $|Pr[D(\kappa, \ell, g, view_{\mathcal{A}}(\kappa, \ell)) = 1] - Pr[D(\kappa, \ell, g, sview(\kappa, \ell)) = 1]| = \varepsilon$, where ε is negligible in k. The probability is taken over the random coins used by $\mathsf{Setup}(\kappa, \ell)$, the random coins used by $\mathcal{A}$, $\mathcal{S}$, the reader R and all (uncorrupted) tags, the choice of g, and the coins used by the distinguisher algorithm D.

We now extend our definition to forward and backward zk-privacy. Denote by $(k_{\mathcal{T}_g}^f, s_{\mathcal{T}_g}^f)$ (resp., $(k_{\mathcal{T}_g}^1, s_{\mathcal{T}_g}^1)$) the final (resp., initial) secret-key and internal state of $\mathcal{T}_g$ at the end of (resp., beginning) of the experiment $\mathbf{Exp}_{\mathcal{A}}^{zkp}$. An RFID protocol π is of *forward* (resp., *backward*) *zk-privacy*, if for any PPT CMIM adversary $\mathcal{A}$ there exists a polynomial-time simulator $\mathcal{S}$ such that for all sufficiently large κ and any $\ell = poly(\kappa)$, the following distributions are indistinguishable: $\{k_{\mathcal{T}_g}^f, s_{\mathcal{T}_g}^f (resp., k_{\mathcal{T}_g}^1, s_{\mathcal{T}_g}^1), g, view_{\mathcal{A}}(\kappa, \ell)\}$ and $\{k_{\mathcal{T}_g}^f, s_{\mathcal{T}_g}^f (resp., k_{\mathcal{T}_g}^1, s_{\mathcal{T}_g}^1), g, sview(\kappa, \ell)\}$. For forward/backward zk-privacy, it is required that the challenge tag $\mathcal{T}_g$ should remain *clean* at the end of experiment $\mathbf{Exp}_{\mathcal{A}}^{zkp}$. Note that the adversary is allowed to corrupt the challenge tag after the end of $\mathbf{Exp}_{\mathcal{A}}^{zkp}$.

4.1 Discussions

Why allow $\mathcal{A}_1$ to output an arbitrary set $\mathcal{C}$ of tags, and limit $\mathcal{A}_2$ to blind access to a challenge tag chosen randomly *from $\mathcal{C}$?* The definition of zk-privacy implies that the adversary $\mathcal{A}$ cannot distinguish any challenge tag $\mathcal{T}_g$ from any set $\mathcal{C}$ of tags; otherwise, $\mathcal{A}$ can figure out the identity of $\mathcal{T}_g$ in $\mathcal{C}$ from its view $view_{\mathcal{A}}$, while this tag's identity cannot be derived from any simulator's view $sview$ (a formal proof of this in case of $|\mathcal{C}| = 2$ is provided in Section 5.1). If $\mathcal{C}$ is removed from the definition of zk-privacy, it is possible for the adversary to distinguish any two tags under its attack, even if each of the tags can be perfectly simulated by a simulator. A special case is that each tag has an upper-bound of sessions in its life time so that an adversary can distinguish any two tags by setting one tag to be run out of sessions in the learning stage [17]. In addition, we do not restrict $\mathcal{C}$ to two tags so as to take into account the case that any number of tags may be correlated.

Why limit $\mathcal{A}_1$ to output of clean *tags?* If $\mathcal{A}_1$ is allowed to output "unclean tags", $\mathcal{A}_2$ can trivially violate the zk-privacy. Consider that $\mathcal{A}_1$ selects two tags that are waiting for different round message (e.g., one tag is clean and the other is not), then $\mathcal{A}_2$ can trivially distinguish them by forwarding to $\mathcal{T}_g$ different round messages.

Why allow $\mathcal{S}$ to have access to oracles in $\mathcal{O}$? Suppose that $\mathcal{S}$ simulates a tag from scratch and $\mathcal{A}$ (run by $\mathcal{S}$ as a subroutine) requests to corrupt the tag in the middle of the simulation. Without oracle access, it is difficult or even impossible for $\mathcal{S}$ to continue its simulation and keep it consistent with its previous simulation for the same tag.

Why limit $sview$ to include all oracle answers to queries made by $\mathcal{S}$? This is to restrict $\mathcal{S}$ not to access the oracles in $\mathcal{O}$ more than $\mathcal{A}$ does. The indistinguishability between the simulated view $sview$ and the real view $view_{\mathcal{A}}$ of adversary $\mathcal{A}$ in zk-privacy implies that for any (t, n_1, n_2, n_3, n_4)-adversary $\mathcal{A}$, with overwhelming probability, $\mathcal{S}$ cannot query O_1, O_2, O_3, O_4 more than n_1, n_2, n_3, n_4 times, respectively.

Why require $\mathcal{T}_g$ to remain clean at the end of $\mathbf{Exp}_{\mathcal{A}}^{zkp}$ for forward/backward privacy? In general, forward/backward privacy cannot be achieved if the adversary is allowed to corrupt the challenge tag before the end of its sessions in $\mathbf{Exp}_{\mathcal{A}}^{zkp}$ (i.e., the tag is not clean at the moment of corruption); otherwise, the adversary is able to derive certain protocol messages from the tag's internal state, secret-key, random coins, and the partial session transcript.

More on backward privacy. In general, backward privacy means that even if $\mathcal{A}$ learns the internal state and secret-key of a tag for the v-th session, it still cannot distinguish the run of $(v+1)$-th session run by this tag from a simulated session run. Without loss of generality, we assume that the internal state and secret-key known to $\mathcal{A}$ are the initial ones (i.e., $k_{\mathcal{T}_g}^1$ and $s_{\mathcal{T}_g}^1$). For most RFID protocols in practice, the internal state and the secret-key of any tag at any time t can be determined by the tag's initial state, initial secret-key, and the session transcript related to the tag up to time t. In such a case, the indistinguishability between the simulated view $sview$ of $\mathcal{S}$ and the real view $view_{\mathcal{A}}$ of $\mathcal{A}$ relies upon the random coins used by $\mathcal{T}_g$ in experiment $\mathbf{Exp}_{\mathcal{A}}^{zkp}$. These random coins are not disclosed to $\mathcal{A}$ since the random coins used by an uncorrupted tag in any session are erased once the session is completed, and the challenge tag $\mathcal{T}_g$ is required to be clean at the end of $\mathbf{Exp}_{\mathcal{A}}^{zkp}$.

Why disallow $\mathcal{A}_2$ to corrupt tags in $\mathcal{C}$ in zk-privacy formulation? For any tag $\mathcal{T}_i \in \mathcal{C}$ corrupted by $\mathcal{A}_2$, it can distinguish whether $\mathcal{T}_i$ is the challenge tag $\mathcal{T}_g$ or not, which can nullify any polynomial-time successful simulation by the simulator $\mathcal{S}$ *unless $\mathcal{S}$ can also corrupt the corresponding tags in $\mathcal{C}$.* However, allowing the simulator $\mathcal{S}$ to corrupt (or just get access to) tags in $\mathcal{C}$ weakens simulatability (e.g., in case tags have correlated states) and can even make simulatability meaningless (e.g., in case $\mathcal{S}$ corrupts the challenge tag $\mathcal{T}_g$), as such a simulator may be too powerful. Recall that $\mathcal{A}_2$ and the simulator should, in particular, not know which tag in $\mathcal{C}$ is the challenging tag. For conceptual simplicity, we disallow $\mathcal{A}_2$ to have access to tags in $\mathcal{C}$ other than blind access to the challenge tag $\mathcal{T}_g$. As we shall see, the zk-privacy is still very poweful. We remind that $\mathcal{A}_2$ still can corrupt any tags in $\hat{\mathcal{T}} = \mathcal{T} - \mathcal{C}$.

On some special cases in zk-privacy experiments. One special case is that in the experiment $\mathbf{Exp}_{\mathcal{A}}^{zkp}$, $\mathcal{A}_1$ outputs $\mathcal{C} = \mathcal{T}$. In this case, the simulator $\mathcal{S}_2$ does not have oracle access to any tag. The zk-privacy is analogue to auxiliary-input zero-knowledge [7], where the view of $\mathcal{A}_1/\mathcal{S}_1$ corresponds to the auxiliary input. Another special case is that $\mathcal{A}_1$ outputs only a single tag in $\mathcal{C}$, and all other tags can be corrupted by $\mathcal{A}_1$ and $\mathcal{A}_2$. In this case, the forward/backward zk-privacy implies that both adversary $\mathcal{A}$ and simulator $\mathcal{S}$ have access to certain secret information of all tags.

Comparison with traditional formulation of zero-knowledge. The notion of zk-privacy is defined based on the traditional zero-knowledge formulation [9, 7] with the following differences. First, in zk-privacy, the simulator $\mathcal{S}$ is allowed to have access to oracles in $\mathcal{O}$ (where the actions of these oracles may depend upon some secret values such as secret-keys and internal states), while traditional zk-simulator is a polynomial-time algorithm without oracle access to players of secret values. Second, the zk-privacy is formulated against a structured adversary $\mathcal{A}$ which is divided into two phases, while the traditional zk is formulated against any polynomial-time adversary. Third, in zk-privacy, the random challenge g is unknown to $\mathcal{A}$, but is presented to the distinguisher, which renders extra power to the distinguisher; in comparison, in the traditional zero-knowledge formulation, the distinguisher and the adversary essentially have the same power and advantage. Lastly, for forward (resp., backward) zk-privacy, the final (resp., initial) secret-key and internal state of the challenge tag $\mathcal{T}_g$ are disclosed to $\mathcal{A}$, while for the traditional zero-knowledge formulation, no secret values of the knowledge prover are assumed to be leaked to the adversary.

5 Comparison with Existing Frameworks

In this section, we compare our RFID security and privacy framework with typical existing frameworks. We argue that our framework is more reasonable in practice than some frameworks, and it is stronger in terms of privacy than at least one of the existing frameworks. We also clarify some subtleties and confusions in the existing frameworks. The detailed comparisons, alone with subtlety clarifications, also further justify the zk-privacy formulation.

5.1 Comparison with Model in [17]

The RFID privacy model proposed in [17] describes the indistinguishability between any two tags by an adversary, which is referred to as "ind-privacy." It was mentioned in [17] that an important area for future research is to study stronger RFID privacy notions. We shall prove that zk-privacy is strictly stronger than a revised version of ind-privacy after some subtleties are clarified.

Figure 4 illustrates the ind-privacy experiment $\mathbf{Exp}_{\mathcal{A}}^{ind}[\kappa, \ell, n_1, n_2, n_3, n_4]$ ($\mathbf{Exp}_{\mathcal{A}}^{ind}$, for simplicity) in terms of the notion defined in this paper. In $\mathbf{Exp}_{\mathcal{A}}^{ind}$, an adversary $\mathcal{A}$ is comprised of a pair of algorithms $(\mathcal{A}_1, \mathcal{A}_2)$ and runs in two stages. Throughout the experiment, the adversary $\mathcal{A}$ is allowed to launch O_1, O_2, O_3 and O_4 oracle queries without exceeding n_1, n_2, n_3 and n_4 overall calls, respectively. The experiment proceeds as follows. At first, the experiment runs $\mathsf{Setup}(\kappa, \ell)$ to setup an RFID system $(R, \mathcal{T})$. Then, in the *learning stage*, algorithm $\mathcal{A}_1$ outputs a piece of state information st and a pair of uncorrupted tags $\{\mathcal{T}_{i_0}, \mathcal{T}_{i_1}\}$ to which it has not made $\mathsf{Corrupt}$ queries. Next, the experiment selects a random bit g and sets the challenge tag to be $\mathcal{T}_{i_g}$. Finally, in the *guess stage*, algorithm $\mathcal{A}_2$ is asked to guess the random bit g by outputting a bit b'. During the second stage, $\mathcal{A}_2$ can interact with R and the tags in $\widehat{T} = \mathcal{T} - \{\mathcal{T}_{i_0}, \mathcal{T}_{i_1}\}$, and gets blind access to (but cannot corrupt) the challenge tag $\mathcal{T}_{i_g}$ via the interface $\mathcal{I}$. (In [17], it is stated that $\mathcal{A}_2$ is still allowed to access the challenge tag $\mathcal{T}_{i_g}$ but cannot corrupt $\mathcal{T}_{i_g}$, without formally formulating the interface entity $\mathcal{I}$ as within our framework.)

Experiment $\mathbf{Exp}_{\mathcal{A}}^{ind}[\kappa, \ell, n_1, n_2, n_3, n_4]$
1. run $\mathsf{Setup}(\kappa, \ell)$ to setup the reader R and a set of tags $\mathcal{T}$; denote by $para$ the system public parameters;
2. $\{\mathcal{T}_{i_0}, \mathcal{T}_{i_1}, st\} \leftarrow \mathcal{A}_1^{\mathcal{O}}(R, \mathcal{T}, para)$; //*learning stage*
3. set $\widehat{\mathcal{T}} = \mathcal{T} - \{\mathcal{T}_{i_0}, \mathcal{T}_{i_1}\}$;
4. $g \in_R \{0, 1\}$;
5. $b' \leftarrow \mathcal{A}_2^{\mathcal{O}}(R, \widehat{\mathcal{T}}, \mathcal{I}(\mathcal{T}_{i_g}), st)$; //*guess stage*
6. the experiment outputs 1 if $b' = g$, 0 otherwise.

Fig. 4. Ind-Privacy Experiment

Definition 5.1 (ind-privacy) *The advantage of $\mathcal{A}$, denoted $Adv_{\mathcal{A}}^{ind}(\kappa, \ell, n_1, n_2, n_3, n_4)$, in the experiment* $\mathbf{Exp}_{\mathcal{A}}^{ind}[\kappa, \ell, n_1, n_2, n_3, n_4]$ *is defined to be :*

$$|Pr[\mathbf{Exp}_{\mathcal{A}}^{ind}[\kappa, \ell, n_1, n_2, n_3, n_4] = 1] - \frac{1}{2}|.$$

An adversary $\mathcal{A}$ $(\epsilon, t, n_1, n_2, n_3, n_4)$-breaks the ind-privacy of the RFID system $(R, \mathcal{T})$ if the advantage of $\mathcal{A}$ in the experiment $\mathbf{Exp}_{\mathcal{A}}^{ind}$, *i.e., $Adv_{\mathcal{A}}^{ind}(k, \ell, n_1, n_2, n_3, n_4)$, is at least ϵ and the running time of $\mathcal{A}$ is at most t.*

An RFID system $(R, \mathcal{T})$ is said to be of ind-privacy, if there exists no adversary who can $(\epsilon, t, n_1, n_2, n_3, n_4)$-break the RFID system for some non-negligible ϵ and some polynomials t, n_1, n_2, n_3, n_4 (all of them are in κ).

On some subtleties in ind-privacy. In the original definition of ind-privacy, it is not explicitly specified that the two tags output by $\mathcal{A}_1$ must be clean tags. In the definition of forward ind-privacy [17], it is not precisely specified the time point of tag corruption and the actions of adversary after tag corruption.

zk-privacy vs. ind-privacy for single-tag systems. We note that any RFID protocol, *even if it just reveals the tag's secret-key*, trivially satisfies ind-privacy for special RFID systems consisting only one tag (e.g., for a unique item of high value). The reason is that in this special scenario, the view of $\mathcal{A}$ is independent of the random bit g (as the challenge tag $\mathcal{T}_{i_g}$ is always the unique tag regardless of the choice of g), and thus $\Pr[b' = g]$ is just $\frac{1}{2}$ for any adversary. In comparison, in this special scenario the zk-privacy is essentially degenerated to the traditional zero-knowledge definition, which still provides very reasonable privacy guarantee.

Theorem 1 *zk-privacy is stronger than ind-privacy.*

Proof. First, we show that zk-privacy implies ind-privacy, which holds unconditionally. In other words, if an RFID system $(R, \mathcal{T})$ does not satisfy ind-privacy, then it also does not satisfy zk-privacy. To prove this, we show that if there exists a PPT adversary $\mathcal{A} = (\mathcal{A}_1, \mathcal{A}_2)$ which can $(\epsilon, t, n_1, n_2, n_3, n_4)$-break the ind-privacy of the RFID system $(R, \mathcal{T})$, then we can construct another PPT adversary $\mathcal{A}'$ such that no PPT simulator exists for $\mathcal{A}'$.

In the experiment $\mathsf{Exp}_{\mathcal{A}'}^{zkp}$, let $\mathcal{A}'$ run $\mathcal{A}$ and do whatever $\mathcal{A}$ does. In particular, $\mathcal{A}'$ and $\mathcal{A}$ are of the same parameters (t, n_1, n_2, n_3, n_4). Since $\mathcal{A}$ run by $\mathcal{A}'$ always outputs a *pair* of clean tags at the end of its first stage, $\mathsf{Exp}_{\mathcal{A}'}^{zkp}$ outputs $(g, view_{\mathcal{A}'})$, where $g \in \{0, 1\}$ is a random bit, and $view_{\mathcal{A}'}$ implicitly determines the output of $\mathcal{A}$ (i.e., the guessed bit b'). That is, the guessed bit b' can be computed out from $view_{\mathcal{A}'}$ in polynomial-time. As we assume $\mathcal{A}$ $(\epsilon, t, n_1, n_2, n_3, n_4)$-breaks ind-privacy, it holds that $\Pr[b' = g]$ is at least $\frac{1}{2} + \epsilon$ for the output of $\mathsf{Exp}_{\mathcal{A}'}^{zkp}$. However, the simulated view $sview$ in the output of the experiment $\mathsf{Exp}_{\mathcal{S}}^{zkp}$ is independent of g (recall that the random value g is unknown to the simulator $\mathcal{S}$). Therefore, for the guessed bit b' implied by $sview$ (which can be computed out from $sview$ in polynomial-time), it always holds that $Pr[b' = g] = \frac{1}{2}$. This shows that for the above $\mathcal{A}'$ and for any polynomial-time simulator, there exists a polynomial-time distinguisher that can distinguish the output of $\mathsf{Exp}_{\mathcal{A}}^{zkp}$ and that of $\mathsf{Exp}_{\mathcal{S}}^{zkp}$ with non-negligible probability at least ϵ.

Next, we present several protocol examples (based on one-time secure signatures or CPA-secure public-key encryption) that satisfy ind-privacy but dissatisfy zk-privacy.

Consider a special RFID system that consists of only one tag $\mathcal{T}_1$ (and a reader R). The secret-key of $\mathcal{T}_1$ is the signature of $\mathcal{T}_1$'s ID, denoted s_{ID}, signed by R under the public-key of R. Consider an RFID protocol π in which $\mathcal{T}_1$ just reveals its secret-key s_{ID} to R. As discussed above, any RFID protocol trivially satisfies ind-privacy for RFID systems consisting of only one tag, and thus the protocol π is of ind-privacy. But, π clearly does not satisfy zk-privacy. Specifically, considering an adversary $\mathcal{A} = (\mathcal{A}_1, \mathcal{A}_2)$ where $\mathcal{A}_1$ simply outputs $\mathcal{C} = \{\mathcal{T}_1\}$ and then $\mathcal{A}_2$ invokes $\mathcal{T}_g = \mathcal{T}_1$ to get the signature s_{ID}, no PPT simulator can output s_{ID} by the security of the underlying signature scheme. Note that one-time secure signature is sufficient to show this protocol example not satisfying zk-privacy, and one-time secure signatures can be based on any one-way function [23].

Given any ind-private two-round RFID protocol $\pi = (c, a)$ for an RFID system $(R, \mathcal{T})$, where $\mathcal{T}$ consists of polynomially many tags, c is the first-round message from the reader and a is the response from a tag, we transform π into a new protocol π' as follows: In the protocol π', besides their respective secret-keys all tags in $\mathcal{T}$ also share a unique pair of public-key PK and secret-key SK for a CPA-secure public-key encryption scheme. For a protocol run of π' between the reader R and a tag $\mathcal{T}_i$, R sends $c' = E_{PK}(c)$ in the first-round, and $\mathcal{T}_i$ decrypts c' to get c and then sends back $a' = c\|a$. The protocol π' could appear in the scenario of tag group authentication, where the ability of sending back c can demonstrate the membership of the group identified by the public-key PK. Furthermore, in the scenario of anonymizer-enabled RFID systems [10], the decryption operation can be performed by the anonymizer. As in the new protocol π' all tags share the same public-key PK, the ind-privacy of π' is inherited from that of π. Specifically, the session transcripts of π' can be computed in polynomial-time from the session transcripts of π and the public-key PK. However, π' does not satisfy zk-privacy. Specifically, consider an adversary $\mathcal{A} = (\mathcal{A}_1, \mathcal{A}_2)$, where $\mathcal{A}_1$ simply outputs the set of clean tags $\mathcal{C} = \mathcal{T}$ (in particular, $\mathcal{A}$ never corrupts tags) and then $\mathcal{A}_2$ blindly interacts with the challenge tag $\mathcal{T}_g$ for only one session. By the CPA-security of the underlying public-key encryption scheme, no PPT simulator can handle the $\mathsf{SendT}(challenge, \hat{c})$ queries made by $\mathcal{A}_2$, as such ability implies the ability of ciphertext decryption. Note that CPA security is sufficient here, as the adversary $\mathcal{A}$ involves only one session with the challenge tag $\mathcal{T}_g$. $\qquad\square$

We remark that though the above two protocol examples may not be very realistic, they do separate the zk-privacy notion and the ind-privacy notion. We leave it an interesting question to find more protocol examples that are ind-private but not zk-private

5.2 Comparison with Model in [25, 22]

In the framework of [25, 22], the adversary is categorized into the following classes: (1) Weak adversary, which cannot corrupt any tags; (2) Forward adversary, which can corrupt tags under the limitation that once the adversary corrupts a tag, it can do nothing subsequently except for corrupting more tags; (3) Destructive adversary, which can do anything after a tag corruption, but under the limitation that the adversary cannot reuse a tag after corrupting it. Specifically, once a tag is corrupted it will be virtually destroyed. In particular, a destructive adversary cannot observe or interact with a corrupted tag nor can the adversary impersonate a corrupted tag to the reader; (4) Strong adversary, which has no limitations on corrupting tags, and can do anything at its wish. For each category of adversary defined above, it is also defined a *narrow* variant, where a narrow adversary cannot access the outputs of the players for any protocol run.

Suppose that C is one of the adversary categories. Informally, an RFID protocol is called C-private if for any adversary $\mathcal{A} \in C$, there exists a simulator $\mathcal{S}$ such that $\mathcal{A}$ cannot distinguish its interactions with the actual RFID system or with the simulator (the reader is referred to [25, 22] for formal definitions).

The major differences between our framework and that of [25, 22] are summarized below.

In [25, 22], the simulator is not required to handle tag corruption queries by the adversary. In other words, the simulator works only for those adversaries which do not make tag corruption queries. It is not clear how such a simulator acts upon tag corruption queries made by an adversary. Suppose that $\mathcal{S}$ simulates a tag from scratch and $\mathcal{A}$ (typically run by $\mathcal{S}$ as a subroutine) requests to corrupt the tag in the middle of simulation (possibly in the middle of a session run). Without access to tag corruption queries, it is difficult or even impossible for $\mathcal{S}$ to continue its simulation for the tag and keep it consistent with its previous simulation for the same tag.

The adversary considered in our framework essentially corresponds to strong adversary in [25, 22], with the difference in that the adversary cannot corrupt any tag in set C before the end of zk-privacy experiment $\mathbf{Exp}_{\mathcal{A}}^{zkp}$. In comparison, the model in [25, 22] poses no restriction on tag corruption (though it is not clear how the simulator handles such adversaries), which implies that an adversary can corrupt any tag at any time (possibly in the middle of session). However, in such a case, forward/backward privacy may not be achievable if the challenge tag is corrupted in the middle of a session; this is the reason why we require that the challenge tag $\mathcal{T}_g$ must remain *clean* at the moment of corruption. Indeed, there are some confusions in [25, 22].

The matching session concept defined in [25, 22] is restricted to identical session transcript, without clarifying some subtleties such as the "last-round-message attacks" for defining authentication from tag to reader.

The notion of adaptive completeness is not defined in [25, 22]. The completeness notion in [25, 22] is defined for honest protocol execution only, with no adversarial desynchronizing attacks being taken into account.

The privacy notions proposed in [25, 22] and that proposed in [17] are incomparable *in general* (though for certain concrete adversarial strategies, the privacy notions of [25, 22] may imply the ind-privacy), while the privacy notion proposed in this work is strictly stronger than that of [17].

5.3 Comparison with Models in [11, 19]

The RFID privacy notion given in [11, 19] is formulated based on the unpredictability of protocol output. We refer to this privacy notion as "unp-privacy." The unp-privacy is formulated with respect to RFID protocols with a 3-round canonical form, denoted as $\pi = (c, r, f)$, where c, r, f stand for the first, second, and third round message, respectively. Note that our framework, as well as models in [17, 25, 22]), are not confined to this protocol structure.

The unp-privacy notion formulated in [11, 19] essentially says that the second-round message sent from a tag must be pseudorandom (i.e., indistinguishable from a truly random string). We observe that this requirement has certain limitations. First, given any unp-private RFID protocol $\pi = (c, r, f)$ between a reader and a tag, we can modify the protocol to $\pi' = (c, r||1, f)$, where "$||$" denotes the string concatenation operation. That is, the modified protocol π' is identical to π except that in the second-round the tag additionally concatenates a bit '1' to r. This modified RFID-protocol π' is not of unp-privacy, as the second-round message $r||1$ is clearly not pseudorandom. However, intuitively, the tags' privacy should be preserved since the same bit '1' is appended to all second-round messages for all tags. Notice that when RFID-protocols are implemented in practice, the messages being exchanged between reader and tags normally bear some non-random information such as version number of RFID standard. Another limitation is that the unp-privacy may exclude the use of public-key encryption in RFID-protocols, as public-key generated ciphertexts are typically *not* pseudorandom.

Another point is that the adversaries considered in the definition of unp-privacy [11, 19] is not allowed to access protocol outputs. Therefore, such adversaries are *narrow* ones as defined in [25, 22]. Informally, the unp-privacy experiment works as follows. Given a first-round message c (which could be generated by the adversary $\mathcal{A}$), the experiment selects a value r which could be either the actual second-round message generated by an uncorrupted tag in response to c or just a random value in a certain domain; then the experiment presents the value r to $\mathcal{A}$. The unp-privacy means that $\mathcal{A}$ cannot determine in which case the value r is. Note that if $\mathcal{A}$ has access to protocol outputs, it can simply distinguish between the two cases of r. What $\mathcal{A}$ needs to do is to forward r to the reader R as the second round message. If r is generated by an uncorrupted tag (and the value c was generated by the reader in a matching session), R will always output "accept." On the other hand, if r is just a random value, with overwhelming probability R will reject the message due to authentication soundness from tag to reader.

In summary, we argue that zk-privacy is more reasonable than unp-privacy in practice. It allows for more general protocol structure, more powerful adversary, and non-pseudorandom protocol messages.

5.4 Comparison with Model in [3]

In the model of [3], a tag can have multiple different pseudonyms in different sessions. Then, roughly speaking, an RFID scheme is private, if for any pair (t_i, t_j), no probabilistic polynomial-time algorithm can tell whether t_i, t_j correspond to the same tag or not, where t_i, t_j are two pseudonyms used in two sessions in chronological order. That is, any polynomial-time adversary $\mathcal{A}$ is not able to make the link between several authentications of a same tag [3]. The scheme is called *past* (resp., *future*) private, if $\mathcal{A}$ cannot link (t_i, t_j) even if t_j (resp., t_i) is corrupted. The reader is referred to [3] for details.

We note that, in the framework of [3], upon corruption of a tag, only the secret-key of this tag is revealed to the adversary. In comparison, within our framework, upon corruption of a tag both its secret-key and internal state and *random coins* are revealed to the adversary. We note the security analysis of the RFID protocols proposed in [3], based upon public-key cryptosystems, critically relies upon this assumption on tag corruption (i.e., no random coins are revealed upon tag corruption). Specifically, in the analysis of [3], the "dummy" adversary (corresponding to the simulator within our framework) simulates messages sent by the challenge tag (which are just ciphertexts of a public-key encryption scheme). But, suppose the challenge tag is corrupted just after the dummy adversary sent the simulated ciphertexts, and the adversary learns both the secret-key and *random coins* of the challenge tag, then the dummy adversary (i.e., the simulator) cannot give a consistent simulation further (i.e., unable to make the simulated ciphertext to be consistent with the private-key and random coins of the corrupted tag).

In the framework of [3], no internal state update mechanism is given. This implies that symmetric encryptions or MACs alone are almost useless for achieving past or future privacy (without appropriate internal state update mechanisms). Indeed, all the authentication protocols proposed in [3] are based upon public-key cryptosystems.

In the framework of [3], only authentication from tag to reader is formulated, while mutual (both tag-to-reader and reader-to-tag) authentication is formulated within our framework.

6 An RFID Protocol within Our Framework

Let $F_k: \{0,1\}^{2\kappa} \rightarrow \{0,1\}^{2\kappa}$ be a pre-specified keyed PRF and F_k^0 (resp., F_k^1) the κ-bit prefix (resp., suffix) of the output of F_k, where κ is the system security parameter. In practice, the PRF can be implemented based on some lightweight stream or block ciphers [13, 2, 12]. When a tag $\mathcal{T}_i$ with identity ID registers to the reader R, it is assigned a secret-key $k \in_R \{0,1\}^\kappa$, a counter ctr of length l_{ctr} with initial value 1. R pre-computes an initial index $I = F_k^0(1||pad_1)$ for the tag, where $pad_1 \in \{0,1\}^{2\kappa-l_{ctr}}$ is a fixed padding, and stores the tuple (I, k, ctr, ID) into its database.

At the start of a new protocol session, R sends a challenge string $c \in_R \{0,1\}^\kappa$ to $\mathcal{T}_i$, which also serves as the session identifier. To simplify the presentation, the session identifier and the corresponding verification of the identifier by protocol players are implicitly implied and will not be explicitly mentioned in the following.

Upon receiving c from R, $\mathcal{T}_i$ computes $I = F_k^0(ctr||pad_1)$, $(r_0, r_1) = F_k(c||I)$ (where $r_0 = F_k^0(c||I)$ and $r_1 = F_k^0(c||I)$), and $r_\mathcal{T} = r_0 \oplus (ctr||pad_2)$. $\mathcal{T}_i$ sends $(I, r_\mathcal{T})$ to R and then updates its counter $ctr = ctr + 1$, where $pad_2 \in \{0,1\}^{\kappa-l_{ctr}}$ is another predetermined padding string.

After receiving $(I, r_\mathcal{T})$, R searches its database to find a tuple indexed by I:

– If R finds such a tuple, say (I, k, ctr', ID), it computes $(r_0, r_1) = F_k(c||I)$, and checks whether $ctr'||pad_2 = r_0 \oplus r_{\mathcal{T}}$: If yes, R accepts $\mathcal{T}_i$ by outputting "1", sends $r_R = r_1$ to the tag, updates the tuple (I, k, ctr', ID) with $ctr' = ctr' + 1$ and $I = F_k^0(ctr'||pad_1)$; If not, R searches for the next tuple including I (to avoid potential collision of index I, i.e., two different tuples are of the same index I).

– If no tuple is found to have an index I (which indicates counter desynchronization between R and $\mathcal{T}_i$), for each tuple (I', k, ctr', ID) in its database, R computes $(r_0, r_1) = F_k(c||I)$ and $ctr||pad_2 = r_0 \oplus r_{\mathcal{T}}$, and checks whether $I = F_k^0(ctr||pad_1)$: If yes (which indicates ctr is the correct counter value at $\mathcal{T}_i$), R accepts $\mathcal{T}_i$, outputs "1", sends back $r_R = r_1$ as the third message, and updates the tuple (I', k, ctr', ID) with $ctr' = ctr + 1$ and $I' = F_k^0(ctr'||pad_1)$. In the case that R fails with all the tuples in its database, it rejects the tag and outputs "0".

Upon receiving r_R, $\mathcal{T}_i$ checks whether $r_R = r_1$: If yes, $\mathcal{T}_i$ accepts the reader and outputs "1"; otherwise it rejects the reader and outputs "0". This RFID protocol is also depicted in Figure 5.


Fig. 5. RFID Protocol with Mutual Authentication and ZK-Privacy

In comparison with the protocol proposed in [19], the above protocol adds mutual authentication (and is logically more precise), and we can formally prove that it is of adaptive completeness, mutual authentication, and zk-privacy within the new framework. Analysis of completeness and authentication was not conducted in [19], and as we shall see, the zk-privacy analysis of the new protocol is much more complicated than the unp-privacy analysis in [19]. We suggest that the methodology used in our analysis is of independent interest, which can be applied to analyze other RFID protocols (particularly those based on PRFs) within our new framework.

Theorem 2 *Assuming F_k is a pseudorandom function, the protocol specified above satisfies adaptive completeness, mutual authentication and zk-privacy.*

Below we provide a high level analysis of the zk-privacy property. The complete detailed proof of the theorem is given in Appendix A.

The core of the simulation by the simulator $\mathcal{S} = (\mathcal{S}_1, \mathcal{S}_2)$, who runs the underlying adversary $\mathcal{A} = (\mathcal{A}_1, \mathcal{A}_2)$ as a subroutine, lies in the actions of $\mathcal{S}_2$ in dealing with the following queries made by $\mathcal{A}_2$ to the reader R and the challenge tag $\mathcal{T}_g$. $\mathcal{S}_1$ just mimics $\mathcal{A}_1$ by using the PRF F_k.

1. On oracle query InitReader(), $\mathcal{S}_2$ makes the same oracle query to R, and gets back a random string $c \in \{0, 1\}^\kappa$ from R. Then, $\mathcal{S}_2$ relays back c to $\mathcal{A}_2$.

2. On oracle query $\mathsf{SendT}(challenge, \hat{c})$, where the challenge tag $\mathcal{T}_g$ (simulated by $\mathcal{S}_2$) currently does not run any session, $\mathcal{S}_2$ opens a session for $\mathcal{T}_g$ with $\hat{c}$ as the first-round message (that also serves as the session-identifier of this new session); Then, $\mathcal{S}_2$ randomly selects $I, r_{\mathcal{T}} \in_R \{0,1\}^\kappa$, and sends back $I||r_{\mathcal{T}}$ to $\mathcal{A}_2$ as the second-round message.

3. On oracle query $\mathsf{SendR}(\hat{c}, \hat{I}||\hat{r}_{\mathcal{T}})$, $\mathcal{S}_2$ works as follows:

   Case-3.1. If $\hat{I}||\hat{r}_{\mathcal{T}}$ was sent by $\mathcal{T}_g$ (simulated by $\mathcal{S}_2$) in a session of session-identifier $\hat{c}$, $\mathcal{S}_2$ simulates the responses of the reader R as follows:

   Case-3.1.1 If R is running an incomplete session of session-identifier $\hat{c}$ (i.e., $\hat{c}$ was sent by R upon an $\mathsf{InitReader}$ query and R is waiting for the second-round message), $\mathcal{S}_2$ just returns a random string $r_R \in_R \{0,1\}^\kappa$ to $\mathcal{A}_2$, and outputs "1" indicating "accept".

   Case-3.1.2. Otherwise, $\mathcal{S}_2$ simply returns a special symbol "$\perp$" indicating invalid query.

   Case-3.2. In all other cases, $\mathcal{S}_2$ makes the same oracle query $\mathsf{SendR}(\hat{c}, \hat{I}||\hat{r}_{\mathcal{T}})$ to the reader R, and relays back the answer from R to $\mathcal{A}_2$.

4. On oracle query $\mathsf{SendT}(challenge, \hat{r}_R)$, where the challenge tag $\mathcal{T}_g$ (simulated by $\mathcal{S}_2$) currently runs a session of partial session-transcript $(\hat{c}, I||r_{\mathcal{T}})$ and is waiting for the third-round message, $\mathcal{S}_2$ works as follows:

   Case-4.1. If there exists a matching session of the same session transcript $(\hat{c}, I||r_{\mathcal{T}}, \hat{r}_R)$ at the side of R (where $\hat{r}_R$ may be simulated by $\mathcal{S}_2$ as in the above Case-3.1), $\mathcal{S}_2$ outputs "1" indicating "accept".

   Case-4.2. Otherwise, $\mathcal{S}_2$ simply outputs "0" indicating "reject".

5. Output of $\mathcal{S}_2$: Finally, whenever $\mathcal{A}_2$ stops, $\mathcal{S}_2$ also stops and outputs the simulated view $sview$ as specified in the zk-privacy definition, which particularly consists of all oracle answers (including ones provided by the real oracles in $\mathcal{O}$ and *ones simulated by $\mathcal{S}_2$*) to queries made by $\mathcal{A}$.

It is easy to see that $\mathcal{S}$ works in polynomial-time. We investigate the differences between the simulated view $sview$ output by $\mathcal{S}$ and the real view $view_{\mathcal{A}}$ of $\mathcal{A}$:

Difference-1: In Case-4.1 (resp., Case-4.2) $\mathcal{S}_2$ always outputs "accept" (resp., "reject"), while the actual challenge tag $\mathcal{T}_g$ may output "reject" in Case-4.1 (resp., "accept" in Case-4.2) in the experiment $\mathbf{Exp}_{\mathcal{A}}^{zkp}$.

Difference-2: On oracle query $\mathsf{SendT}(challenge, \hat{c})$ or in Case-3.1 upon the oracle query $\mathsf{SendR}(\hat{c}, \hat{I}||\hat{r}_{\mathcal{T}})$, $\mathcal{S}_2$ always returns truly random strings, while the actual players (i.e., $\mathcal{T}_g$ and R) provide pseudorandom strings in the experiment $\mathbf{Exp}_{\mathcal{A}}^{zkp}$ by invoking the PRF F_k where k is the secret-key of $\mathcal{T}_g$.

Intuitively, Difference-1 can occur only with negligible probability, by the properties of adaptive completeness and mutual authentication. The subsequent analysis argues that the properties of adaptive completeness and mutual authentication indeed hold under the simulation of $\mathcal{S}$ in $\mathbf{Exp}_{\mathcal{S}}^{zkp}$.

Intuitively, Difference-2 should not constitute distinguishable gap between $sview$ and $view_{\mathcal{A}}$, due to the pseudo-randomness of F_k. However, the technical difficulty and subtlety here is that: the difference between pseudorandomness and real randomness only occurs in the second stages of both $\mathbf{Exp}_{\mathcal{A}}^{zkp}$ and $\mathbf{Exp}_{\mathcal{S}}^{zkp}$ (i.e., $\mathcal{A}_2$ and $\mathcal{S}_2$), while both $\mathcal{S}_1$ and $\mathcal{A}_1$ are w.r.t. the PRF F_k. In other words, to distinguish the PRF F_k from a truly random one in the second stage, the distinguisher has already accessed F_k for polynomially many times in the first stage. In general, the definition of PRF says nothing on the pseudorandomness in the second stage. To overcome this technical difficulty, we build a list of hybrid experiments.

In the first hybrid experiment, a polynomial-time algorithm $\hat{S}$ runs $\mathcal{A}$ as a subroutine and has oracle access to the PRF F_k or a truly random function H. $\hat{S}$ first randomly guesses the challenge tag $\mathcal{T}_g$ (by taking g uniformly at random from $\{1, \cdots, \ell\}$), and then setups the RFID system $(R, \mathcal{T})$ except for the challenge-tag $\mathcal{T}_g$. Note that $\hat{S}$ can perfectly handle all oracle queries made by $\mathcal{A}$ to the reader R and all tags in $\mathcal{T} - \{\mathcal{T}_g\}$. For oracle queries directed to $\mathcal{T}_g$, $\hat{S}$ mimics $\mathcal{T}_g$ with the aid of its oracle, i.e, the PRF F_k or a truly random function H. Denote by the view of $\mathcal{A}$ under the run of $\hat{S}$ with oracle access to F_k (resp., H) as $view_{\mathcal{A}}^{\hat{S}^{F_k}}$ (resp., $view_{\mathcal{A}}^{\hat{S}^H}$). By the pseudorandomness of F_k, we have that $view_{\mathcal{A}}^{\hat{S}^{F_k}}$ and $view_{\mathcal{A}}^{\hat{S}^H}$ are indistinguishable. Next, suppose $\hat{S}$ successfully guesses the challenge tag $\mathcal{T}_g$ (that occurs with probability $\frac{1}{\ell}$), $view_{\mathcal{A}}^{\hat{S}^{F_k}}$ is identical to $view_{\mathcal{A}}$. In particular, in this case, the properties of adaptive completeness and mutual authentication hold in $view_{\mathcal{A}}^{\hat{S}^{F_k}}$ and thus also in $view_{\mathcal{A}}^{\hat{S}^H}$ (as $view_{\mathcal{A}}^{\hat{S}^{F_k}}$ and $view_{\mathcal{A}}^{\hat{S}^H}$ are indistinguishable). Thus, to show the indistinguishability between $view_{\mathcal{A}}$ and $sview$, it is reduced to show the indistinguishability between $view_{\mathcal{A}}^{\hat{S}^H}$ (in case $\hat{S}$ successfully guesses the challenge tag $\mathcal{T}_g$) and $sview$.

In the second hybrid experiment, we consider another polynomial-time algorithm S' that mimics $\hat{S}$, with oracle access to F_k or H, but with the following modifications: in the second stage of this hybrid experiment, S' essentially mimics the original zk-privacy simulator $\mathcal{S}$. Denote by the view of $\mathcal{A}$ under the run of S' with oracle access to F_k (resp., H) as $view_{\mathcal{A}}^{S'^{F_k}}$ (resp., $view_{\mathcal{A}}^{S'^{H}}$). By the pseudorandomness of F_k, $view_{\mathcal{A}}^{S'^{F_k}}$ and $view_{\mathcal{A}}^{S'^{H}}$ are indistinguishable. We can show that $view_{\mathcal{A}}^{S'^{H}}$ and $view_{\mathcal{A}}^{\hat{S}^{H}}$ are also indistinguishable, and that $view_{\mathcal{A}}^{S'^{F_k}}$ and $sview$ are also indistinguishable (conditioned on S' successfully guesses the challenge tag $\mathcal{T}_g$), which particularly implies that the properties of adaptive completeness and mutual authentication hold also in $sview$. This establishes the indistinguishability between $sview$ and $view_{\mathcal{A}}$.

7 Future Work

One of our future research directions is to analyze existing RFID protocols and design new protocols within the new framework presented in this paper.

Since our framework is formulated w.r.t. the basic scenario of an RFID system consisting of a single uncompromised reader and multiple tags, where tags identify themselves to the reader individually and independently. A future research direction is to extend our RFID privacy framework to more sophisticated and practical scenarios which allow compromising of readers, tag group authentication, anonymizer-enabled RFID systems, and tag ownership transfer.

References

1. C. Berbain, O. Billet, J. Etrog and H. Gilbert. An Efficient Forward Private RFID Protocol. In *Conference on Computer and Communications Security – CCS'09*.

2. C. de Canniere and B. Preneel. Trivium. In M. Robshaw and O. Billet, editors, New Stream Cipher Designs: The eSTREAM Finalists, volume 4986 of LNCS, pages 244-266. Springer-Verlag, 2008.

3. S. Canard1, I. Coisel, J. Etrog, and M. Girault. Privacy-Preserving RFID Systems: Models and Constructions, Cryptology ePrint Archive, Report No. 2010/405.

4. I. Damgård and M. Ostergaard. RFID Security: Tradeoffs between Security and Efficiency. In *Topics in Cryptology–CT-RSA 2008*, volume 4964 of *Lecture Notes in Computer Science*, pages 318–332, 2008.

5. R. H. Deng, Y. Li, A. C. Yao, M. Yung and Y. Zhao. A New Framework for RFID Privacy. Cryptology ePrint Archive, Report No. 2010/059.

6. S. Garfinkel, A. Juels, and R. Pappu. RFID Privacy: An Overview of Problems and Proposed Solutions. *IEEE Security and Privacy*, 3(3):34–43, 2005.

7. O. Goldreich. *The Foundations of Cryptography*, volume I, Basic Tools. Cambridge University Press, 2001.

8. O. Goldreich, S. Goldwasser, and S. Micali. How to construct random functions. *J. ACM*, 33(4):792–807, 1986.

9. S. Goldwasser, S. Micali and C. Rackoff. The Knowledge Complexity of Interactive Proof-Systems In *ACM Symposium on Theory of Computing*, pages 291-304, 1985.

10. P. Golle, M. Jakobsson, A. Juels, and P. Syverson. Universal reencryption for mixnets. In *Topics in Cryptology–CT-RSA 2004*, LNCS 2964, pages 163-178.

11. J. Ha, S. Moon, J. Zhou, and J. Ha. A new formal proof model for RFID location privacy. In *European Symposium on Research in Computer Security (ESORICS) 2008*, volume 5283 of *Lecture Notes in Computer Science*.

12. M. Hell, T. Johansson, and W. Meier. The Grain Family of Stream Ciphers. In M. Robshaw and O. Billet, editors, New Stream Cipher Designs: The eSTREAM Finalists, volume 4986 of LNCS, pages 179-190. Springer-Verlag, 2008.

13. International Standard ISO/IEC 9798 Information technology−Security techniques−Entity authentication−Part 5: Mechanisms using Zero-Knowledge Techniques.

14. Nicholas J. Hopper and Manuel Blum. Secure human identification protocols. In *ASIACRYPT*, pages 52–66, 2001.

15. A. Juels. RFID Security and Privacy: A Research Survey. *IEEE Journal on Selected Areas in Communications*, 24(2):381–394, 2006.

16. A. Juels, R. L. Rivest, and M. Szydlo. The blocker tag: Selective blocking of RFID tags for consumer privacy. In *8th ACM Conference on Computer and Communications Security – ACM CCS*, pages 103–111. ACM Press, 2003.

17. A. Juels and S. Weis. Defining Strong Privacy for RFID. In *International Conference on Pervasive Computing and Communications – PerCom 2007*.

18. Ari Juels and Stephen A. Weis. Authenticating pervasive devices with human protocols. In *CRYPTO*, pages 293–308, 2005.
19. C. Ma, Y. Li, R. Deng, and T. Li. RFID Privacy: Relation Between Two Notions, Minimal Condition, and Efficient Construction. In *Conference on Computer and Communications Security – ACM CCS*, 2009.
20. C. Yu Ng, W. Susilo, Y. Mu, and R. Safavi-Naini. RFID privacy models revisited. In *European Symposium on Research in Computer Security (ESORICS) 2008*, volume 5283 of *Lecture Notes in Computer Science*.
21. C. Yu Ng, W. Susilo, Y. Mu, and R. Safavi-Naini. New Privacy Results on Synchronized RFID Authentication Protocols against Tag Tracing. In *European Symposium on Research in Computer Security (ESORICS) 2009*, pages 321-336, volume 5786 of *Lecture Notes in Computer Science*.
22. R.L.Paise and S. Vaudenay. Muthal Authentication in RFID: Security and Privacy. In *AsiaCCS 2008*, pages 292-299.
23. J. Rompel. One-Way Functions are Necessary and Sufficient for Digital Signatures. In *22nd ACM Symposium on Theory of Computing (STOC'90)*, pages 12C19, 1990.
24. A. Shamir. SQUASH: A New MAC with Provable Security Properties for Highly Constrained Devices Such as RFID Tags. FSE 2008, LNCS 5086, pages 144-157, 2008.
25. S. Vaudenay. On Privacy Models for RFID. In *Advances in Cryptology - Asiacrypt 2007*.
26. S. Weis, S. Sarma, R. Rivest, and D. Engels. Security and Privacy Aspects of Low-Cost Radio Frequency Identification Systems. In *International Conference on Security in Pervasive Computing – SPC 2003*.
27. 860 MHz - 930 MHz Class 1 Radio Frequency Identification Tag Radio Frequency and Logical Communication Interface Specification Candidate Recommendation Version 1.0.1, Auto-ID Center, 2002.

A   Proof Details of Theorem 1

In this section, we present the proof details of Theorem 2.

A.1   Adaptive Completeness

For any uncorrupted tag $\mathcal{T}_i$, denote by $(c, I||r_{\mathcal{T}}, r_R, o_R^c, o_{\mathcal{T}_i}^c)$ the run of a session between R and an uncorrupted tag $\mathcal{T}_i$. Suppose this session corresponds to the v-th session run by $\mathcal{T}_i$ and the j-th session run by R, where $1 \leq v \leq s$ and $1 \leq j \leq s\ell$. That is, after any potential (particularly, desynchronizing) attacks by a PPT CMIM adversary $\mathcal{A}$, finally R completes its j-th session and $\mathcal{T}_i$ completes its v-th session such that these two sessions are of identical round messages and identical session-identifier. Suppose $\mathcal{T}_i$ is of identity ID and secret-key k, and denote by ctr_i the counter value that was used by $\mathcal{T}_i$ to generate I, i.e., $I = F_k^0(ctr_i||pad_1)$, we consider the probability that the event E (defined in Definition 3.1) occurs.

Firstly, we observe that as R sends the third-round message, it implies $o_R^c = 1$. Secondly, we consider the probability that R identified a different tag $\mathcal{T}_{i'}(\neq \mathcal{T}_i)$ of identity $ID'(\neq ID)$ in its j-th session. This event occurs only with one of the following two cases:

Case-1. There exists a tuple (I, k', ctr', ID') in the database of R, where $ID' \neq ID$, such that $ctr'||pad_2 = r_0' \oplus r_{\mathcal{T}}$ and $(r_0', r_R) = F_{k'}(c||I)$, where $I = F_{k'}^0(ctr'||pad_1)$ was pre-computed by R.
Case-2. There exists a tuple (I', k', ctr', ID') in the database of R, where $ID' \neq ID$, such that $ctr||pad_2 = F_{k'}^0(c||I) \oplus r_{\mathcal{T}}$ and $I = F_{k'}^0(ctr||pad_1)$.

As $ID \neq ID'$ and the random secret-keys for different tags are independent, it holds that $\Pr[k = k'] = 2^{-\kappa}$ in the above two cases. Further note that both Case-1 and Case-2 imply that there exists a tuple in the database of R that determines a counter value, denoted $ctr_{i'}$ (i.e., ctr' or ctr), such that $I = F_k^0(ctr_i||pad_1) = F_{k'}^0(ctr_{i'}||pad_1)$. We consider the event that there *exists* an i', $1 \leq i' \neq i \leq \ell$, such that $I = F_k^0(ctr_i||pad_1) = F_{k'}^0(ctr_{i'}||pad_1)$, where k (resp., k') is the secret-key of $\mathcal{T}_i$ (resp., $\mathcal{T}_{i'}$ of identity ID'), and have the following observations:

- There exists an i', $1 \leq i' \neq i \leq \ell$, such that $k' = k$, which happens with probability $(\ell - 1)2^{-\kappa}$;
- Suppose for all i', $1 \leq i' \neq i \leq \ell$, $k' \neq k$ and F_k is a truly random function, we get that the probability that there *exists* an i', $1 \leq i' \neq i \leq \ell$, such that $I = F_k^0(ctr_i||pad_1) = F_{k'}^0(ctr_{i'}||pad_1)$, is also $(\ell - 1)2^{-\kappa}$. Due to the pseudorandomness of the underlying PRF F_k, by straightforward calculation, we get that the probability that there *exists* an i', $1 \leq i' \neq i \leq \ell$, such that $I = F_k^0(ctr_i||pad_1) = F_{k'}^0(ctr_{i'}||pad_1)$ and $k \neq k'$, is at most $(\ell - 1)2^{-\kappa} + \epsilon$, where ϵ is some negligible quantity in κ.

Putting all together, we conclude that the total probability that the above Case-1 or Case-2 occurs is *at most* $\hat{\epsilon} = 2(\ell - 1)2^{-\kappa} + \epsilon$ that is still negligible in κ.

Also note that conditioned on R correctly identifies $\mathcal{T}_i$ in its j-th session, the third-round message r_R will be $F_k^1(c\|I)$ upon which $\mathcal{T}_i$ will output "accept" (i.e., $o_{\mathcal{T}_i}^c = 1$). We conclude that the event E occurs with probability at most $\hat{\epsilon}$ that is negligible in κ.

A.2 Mutual Authentication

We show that suppose there exists a CMIM adversary $\mathcal{A}$ that breaks the tag-to-reader (resp., reader-to-tag) authentication of the RFID protocol depicted in Figure 5, we can construct another algorithm $\mathcal{A}'$ that breaks the pseudorandomness of the underlying PRF.

The algorithm $\mathcal{A}'$ has oracle access to a PRF F_k or a truly random function $H : \{0,1\}^{2\kappa} \to \{0,1\}^{2\kappa}$, where $k \in_R \{0,1\}^\kappa$ is the random seed, and works as follows:

1. Run $\mathsf{Setup}(\kappa, \ell)$ to setup an RFID system $(R, \mathcal{T})$;
2. Randomly take $i \in_R \{1, \cdots, \ell\}$; That is, $\mathcal{A}'$ randomly guesses the target tag $\mathcal{T}_i$, with respect to which the event E_1 defined for tag-to-reader authentication (resp., E_2 defined for reader-to-tag authentication) occurs.
3. Corrupt all tags in $\mathcal{T}$ other than $\mathcal{T}_i$; Note that, by this way, $\mathcal{A}'$ can perfectly mimic all the actions of the tags in $\mathcal{T} - \{\mathcal{T}_i\}$. Moreover, $\mathcal{A}'$ gets and maintains all the records of the database DB of the reader R *except the record corresponding to $\mathcal{T}_i$*.
4. Run $\mathcal{A}$ and perfectly answer all oracle queries (made by $\mathcal{A}$) directed to tags $\mathcal{T} - \{\mathcal{T}_i\}$, and handle oracle queries directed to the reader R and $\mathcal{T}_i$ as follows (in particular, $\mathcal{A}'$ runs a counter ctr for simulating $\mathcal{T}_i$ that is initiated to be 1):
   (a) On query $\mathsf{InitReader}()$, $\mathcal{A}'$ initiates a new session and returns back a random string $c \in_R \{0,1\}^\kappa$. Here, the random string c also serves as the session identifier.
   (b) On query $\mathsf{SendT}(\mathcal{T}_i, \hat{c})$ where $\hat{c} \in \{0,1\}^\kappa$ and $\mathcal{T}_i$ (simulated by $\mathcal{A}'$) is waiting for the first-round message of a new session, $\mathcal{A}'$ initiates a new session with $\hat{c}$ as the session identifier. Then, $\mathcal{A}'$ queries its oracle (i.e., F_k or H) with $ctr\|pad_1$ to get I and then queries its oracle with $\hat{c}\|I$ to get (r_0, r_1). Finally, $\mathcal{A}'$ computes $r_{\mathcal{T}} = r_0 \oplus ctr\|pad_2$, keeps r_1 for this incomplete session, returns back $(I, r_{\mathcal{T}})$ (as the second-round message) and updates $ctr = ctr + 1$.
   (c) On query $\mathsf{SendR}(\hat{c}, (\hat{I}, \hat{r}_{\mathcal{T}}))$ (where $\hat{c}$ is supposed to be a session identifier), $\mathcal{A}'$ first checks whether it is keeping a session (simulated for R) of the session identifier $\hat{c}$ and is waiting for the second-round message: If not, $\mathcal{A}'$ returns back a special symbol "$\perp$" indicating invalid query; If $\mathcal{A}'$ indeed runs a session of session-identifier $\hat{c}$ (which means that $\hat{c}$ is some random string ever sent by $\mathcal{A}'$ as the first-round message) and is waiting for the second-round message, $\mathcal{A}'$ mimics the actions of the reader R as follows: Any computation involving records of tags in $\mathcal{T} - \{\mathcal{T}_i\}$ in the database DB (actually maintained by $\mathcal{A}'$) is performed by $\mathcal{A}'$ itself; The computation involving the database record of $\mathcal{T}_i$ is performed by $\mathcal{A}'$ with the aid of its oracle (i.e., the PRF F_k or a truly random function H). Note that $\mathcal{A}'$ actually does not necessarily know the actual identities of the tags in $\mathcal{T}$ to simulate the reader R in this case.
   (d) On query $\mathsf{SendT}(\mathcal{T}_i, \hat{r}_R)$ where $\hat{r}_R \in \{0,1\}^\kappa$ and $\mathcal{T}_i$ (simulated by $\mathcal{A}'$) is currently running an incomplete session of session-identifier $\hat{c}$ and is waiting for the third-round message, $\mathcal{A}'$ retrieves the value r_1 (generated when generating the second-round message of the session $\hat{c}$) and checks whether $\hat{r}_R = r_1$. If $\hat{r}_R = r_1$ $\mathcal{A}'$ completes the session $\hat{c}$ and returns back "accept", otherwise, it completes the session and returns back "reject".
   (e) On query $\mathsf{Corrupt}(\mathcal{T}_i)$ (i.e., $\mathcal{A}$ tries to corrupt $\mathcal{T}_i$), $\mathcal{A}'$ stops and outputs "failure".
5. If $\mathcal{A}'$ did not output "failure" in the above step (e) due to request by $\mathcal{A}$ to corrupt $\mathcal{T}_i$, $\mathcal{A}'$ stops whenever $\mathcal{A}$ stops, and then investigates and outputs "1" if the event E_1 (resp., E_2) defined in Definition 3.3 (resp., Definition **??**) occurs.

We first note that the running time of $\mathcal{A}'$, denoted t', is polynomially related to the running time of $\mathcal{A}$. Furthermore, suppose the oracle accessed by $\mathcal{A}'$ is the PRF F_k and $\mathcal{A}'$ did not stop outputting "failure" at step (e) due to corruption query of $\mathcal{T}_i$, the view of $\mathcal{A}$ in its real attack and the view of $\mathcal{A}$ under the simulation of $\mathcal{A}'$ are identical. As $\mathcal{A}'$ uniformly

selects the uncorrupted tag $\mathcal{T}_i$ and the view of $\mathcal{A}$ is independent of the choice of i when $\mathcal{A}'$ gets oracle access to F_k, we have that: suppose $\mathcal{A}$ $(\epsilon, t, n_1, n_2, n_3, n_4)$-breaks the tag-to-reader (resp., reader-to-tag) authentication and $\mathcal{A}'$ gets oracle access to the PRF F_k, with probability $\ell^{-1}\epsilon$ the event E_1 (resp., E_2) occurs (under the simulation of $\mathcal{A}'$) w.r.t. *the uncorrupted tag $\mathcal{T}_i$ (selected by $\mathcal{A}'$), on which $\mathcal{A}'$ will output "1".*

Recall that the event E_1 is: $\mathcal{A}$ successfully finishes a session with the reader R of the session transcript $(c, I||r_\mathcal{T}, r_R)$, in which R identified $\mathcal{A}$ as some uncorrupted tag $\mathcal{T}_i$, but no matching session exists at the side of the tag $\mathcal{T}_i$; That is, no session of transcript prefix $(c, I||r_\mathcal{T})$ was ever run or is now running by $\mathcal{T}_i$, where $I = F_k^0(ctr||pad_1)$ and $r_\mathcal{T} = F_k^0(c||I) \oplus (ctr||pad_2)$ for some counter value ctr. The event E_2 is: $\mathcal{A}$ successfully finishes a session with an uncorrupted tag $\mathcal{T}_i$ of the session transcript $(c, I||r_\mathcal{T}, r_R)$, in which $\mathcal{T}_i$ outputs "accept", but no matching session *of the identical session transcript* exists at the side of the reader R.

Now, we consider the probability that $\mathcal{A}'$ outputs "1" when having oracle access to a truly random function H, i.e., the probability that the event E_1 (resp., E_2) occurs w.r.t. $\mathcal{T}_i$ under the simulation of $\mathcal{A}'$ with oracle access to the truly random function H.

We first consider the probability the event E_1 occurs w.r.t. $\mathcal{T}_i$ under the simulation of $\mathcal{A}'$ with oracle access to a truly random function H. Recall that, in this case, $I = H^0(ctr||pad_1)$ and $r_\mathcal{T} = H^0(c||I) \oplus (ctr||pad_2)$ for some ctr, where H^0 (resp., H^1) stands for the κ-bit prefix (resp., suffix) of the output of the truly random function H.

– There exists no session of partial session transcript $(c, I||r'_\mathcal{T})$ at the side of $\mathcal{T}_i$ (in the simulation of $\mathcal{A}'$), where $r'_\mathcal{T} = H^0(c||I) \oplus (ctr'||pad_2)$ for any counter value ctr'. That is, $\mathcal{T}_i$ did not ever compute the value $H(c||I)$ before the event E_1 occurs. In this case, the view of $\mathcal{A}$ under the simulation of $\mathcal{A}'$ is independent of $r_\mathcal{T} = H^0(c||I) \oplus (ctr||pad_2)$ that is a truly random string in $\{0,1\}^\kappa$. Thus, in this case, the event E_1 occurs w.r.t. $\mathcal{T}_i$ with probability at most $n_3 2^{-\kappa}$, where n_3 is the upper-bound of O_2 (i.e., SendR) oracle queries made by $\mathcal{A}$.

– For each session of partial session transcript $(c, I||r'_\mathcal{T})$ at the side of $\mathcal{T}_i$ (in the simulation of $\mathcal{A}'$), where $r'_\mathcal{T} = H^0(c||I) \oplus (ctr'||pad_2)$ and $I = H^0(ctr'||pad_1)$ for some counter value ctr', we observe that: (1) As we assume no matching session exists at the side of $\mathcal{T}_i$, it must be that $r'_\mathcal{T} \neq r_\mathcal{T} = H^0(c||I) \oplus (ctr||pad_2)$ and thus $ctr' \neq ctr$; (2) But, the fact that $I = H^0(ctr||pad_1) = H^0(ctr'||pad_1)$ means that such a session exists with probability $2^{-\kappa}$. As we assume each tag involves at most s sessions, where $s \leq n_2$ and n_2 is the upper-bound of O_2 (i.e., SendT) oracle queries made by $\mathcal{A}$, we conclude that the probability that there *exists* a session of partial session transcript $(c, I||r'_\mathcal{T})$ at the side of $\mathcal{T}_i$ (in the simulation of $\mathcal{A}'$), where $r'_\mathcal{T} = H^0(c||I) \oplus (ctr'||pad_2)$ and $I = H^0(ctr'||pad_1)$ for some counter value ctr', is at most $s \cdot 2^{-\kappa}$.

Putting all together, we get that the probability that event E_1 occurs w.r.t. $\mathcal{T}_i$ under the simulation of $\mathcal{A}'$ when oracle accessing a truly random function, on which $\mathcal{A}'$ outputs "1", is at most $(n_3 + s)2^{-\kappa} \leq (n_3 + n_2)2^{-\kappa}$.

Now, we consider the probability that event E_2 occurs w.r.t. $\mathcal{T}_i$ under the simulation of $\mathcal{A}'$ with oracle access to a truly random function H. Recall that, in this case, the event E_2 w.r.t. $\mathcal{T}_i$ is: $\mathcal{A}$ successfully finishes a session with the uncorrupted tag $\mathcal{T}_i$ of the session transcript $(c, I||r_\mathcal{T}, r_R)$, in which $\mathcal{T}_i$ outputs "accept", but no matching session *of the identical session transcript* exists at the side of the reader R, where $I = H^0(ctr||pad_1)$, $r_\mathcal{T} = H^0(c||I) \oplus (ctr||pad_2)$ and $r_R = H^1(c||I)$. We investigate several cases:

– The reader R (simulated by $\mathcal{A}'$) did not send $r_R = H^1(c||I)$ (before the event E_2 occurs). In this case, the view of $\mathcal{A}$ (under the simulation of $\mathcal{A}'$) is independent of $r_R = H^1(c||I)$ that is a random string in $\{0,1\}^\kappa$. Thus, the probability that event E_2 occurs w.r.t. $\mathcal{T}_i$ is at most $s \cdot 2^{-\kappa}$ in this case, where $s(\leq n_2)$ is the upper-bound of sessions involved by $\mathcal{T}_i$.

– For each session run by R in which R sends $r_R = H^1(c||I)$ to $\mathcal{T}_i$, we observe that: This session is of session transcript $(c, I||r'_\mathcal{T}, r_R))$, where $r'_\mathcal{T} = H^0(c||I) \oplus (ctr'||pad_2)$ and $I = H^0(ctr'||pad_1)$ for some counter value ctr'; As we assume no matching session exists at the side of R, we get $r'_\mathcal{T} = H^0(c||r) \oplus (ctr'||pad_2) \neq r_\mathcal{T} = H^0(c||r) \oplus (ctr||pad_2)$, and thus $ctr' \neq ctr$; But, as $\Pr[I = H^0(ctr||pad_1) = H^0(ctr'||pad_1)|ctr' \neq ctr] = 2^{-\kappa}$, this session happens with probability $2^{-\kappa}$. As $\mathcal{A}$ makes at most n_3 SendR (i.e., O_3) queries, we have that the probability that there exists a session at the side of R in which R sends $H^1(c||r)$ is at most $n_3 \cdot 2^{-\kappa}$.

Putting all together, we get that the probability E_2 occurs w.r.t. $\mathcal{T}_i$ under the simulation of $\mathcal{A}'$ when oracle accessing a truly random function, on which $\mathcal{A}'$ outputs "1", is also at most $(s + n_3)2^{-\kappa} \leq (n_2 + n_3)2^{-\kappa}$.

Recall that, suppose $\mathcal{A}$ $(\epsilon, t, n_1, n_2, n_3, n_4)$-breaks the tag-to-reader authentication, we have shown that $\mathcal{A}'$ outputs "1" with probability at least $\ell^{-1}\epsilon$ when oracle accessing the PRF F_k. Denote by $\epsilon_1 = |\ell^{-1}\epsilon - (s + n_3)2^{-\kappa}|$, suppose $\mathcal{A}$ $(\epsilon, t, n_1, n_2, n_3, n_4)$-breaks the tag-to-reader or reader-to-tag authentication, we conclude that $\mathcal{A}'$ can (t', ϵ_1)-breaks the pseudorandomness of the underlying PRF, where t' is polynomially related to t. As we assume the underlying PRF is secure, i.e., ϵ_1 is negligible for any t' that is polynomial in κ, we have that ϵ must also be negligible for any t-time adversary $\mathcal{A}$ where t is polynomial in κ. This establishes both the tag-to-reader and the reader-to-tag authentication of the protocol depicted in Figure 5.

A.3 zk-Privacy

According to the zk-privacy formulation presented in Section 4, after an RFID system $(R, \mathcal{T})$ is setup by $\mathsf{Setup}(\kappa, \ell)$, for any PPT adversary $\mathcal{A} = (\mathcal{A}_1, \mathcal{A}_2)$ the simulator $\mathcal{S} = (\mathcal{S}_1, \mathcal{S}_2)$ works as follows: In the *first stage* of $\mathcal{S}$, $\mathcal{S}_1$ runs $\mathcal{A}_1$ as a subroutine, and concurrently interacts with R and all the tags in $\mathcal{T}$ (via the four oracles in $\mathcal{O} = \{O_1, O_2, O_3, O_4\}$). For all oracle queries made by $\mathcal{A}_1$, $\mathcal{S}_1$ makes the same oracle queries, and relays back the oracle answers to $\mathcal{A}_1$. Finally, $\mathcal{S}_1$ outputs whatever $\mathcal{A}_1$ outputs at the end of the first stage, say $(\mathcal{C}, st)$, where $\mathcal{C} = \{T_{i_1}, T_{i_2}, \cdots, T_{i_\delta}\} \subseteq \mathcal{T}$ is a set of *clean* tags, $0 \le \delta \le \ell$, and st is some state information to be transmitted to the second stage. As for any adversary $\mathcal{A}$ who outputs an empty set $\mathcal{C}$ of clean tags (i.e., no challenge tag will be chosen), the view of $\mathcal{A}$ can be trivially perfectly simulated by the simulator with oracle access to all the tags and the reader R. In the following analysis, we focus on the case that $|\mathcal{C}| \ge 1$, i.e., $1 \le \delta \le \ell$.

Then, a value g is selected uniformly at random from $\{1, \cdots, \delta\}$, which specifies the challenge tag $\mathcal{T}_g = \mathcal{T}_{i_g}$. In the *second stage* of $\mathcal{S}$, $\mathcal{S}_2$ runs $\mathcal{A}_2(st)$ as a subroutine and concurrently interacts with the reader R and the tags in $\widehat{\mathcal{T}} = \mathcal{T} - \mathcal{C}$. For all oracle queries made by $\mathcal{A}_2$ directed to tags in $\widehat{\mathcal{T}} = \mathcal{T} - \mathcal{C}$, $\mathcal{S}_2$ makes the same oracle queries, and relays back oracle answers to $\mathcal{A}_2$. But, for oracle queries made by $\mathcal{A}_2$ directed to the reader R and to the challenge tag $\mathcal{T}_g = \mathcal{T}_{i_g}$ (blindly accessed by $\mathcal{A}_2$), $\mathcal{S}_2$ works as follows: (Recall that, according to the zk-privacy formulation in Section 4, the oracle queries made by $\mathcal{A}_2$ to the challenge tag is of the form $\mathsf{SendT}(challenge, \cdot)$.)

1. On oracle query $\mathsf{InitReader}()$ made by $\mathcal{A}_2$, $\mathcal{S}_2$ makes the same oracle query to R, and gets back a random string $c \in \{0,1\}^\kappa$ from R. Then, $\mathcal{S}_2$ relays back c to $\mathcal{A}_2$.
2. On oracle query $\mathsf{SendT}(challenge, \hat{c})$, where the challenge tag $\mathcal{T}_g$ (simulated by $\mathcal{S}_2$) currently does not run any session (which means $\hat{c}$ will be treated as the first-round message), $\mathcal{S}_2$ opens a session for $\mathcal{T}_g$ with $\hat{c}$ as the first-round message (that also serves as the session-identifier of this new session); Then, $\mathcal{S}_2$ randomly selects $I, r_\mathcal{T} \in_R \{0,1\}^\kappa$, and sends back $I||r_\mathcal{T}$ to $\mathcal{A}_2$ as the second-round message.
3. On oracle query $\mathsf{SendR}(\hat{c}, \hat{I}||\hat{r}_\mathcal{T})$, $\mathcal{S}_2$ works as follows:

    Case-3.1. If $\hat{I}||\hat{r}_\mathcal{T}$ was sent by the challenge tag $\mathcal{T}_g$ (simulated by $\mathcal{S}_2$) in a session of session-identifier $\hat{c}$ (i.e., the first-round message), in this case $\mathcal{S}_2$ simulates the responses of the reader R according to the following two cases. (In particular, $\mathcal{S}_2$ does not make the oracle query $\mathsf{SendR}(\hat{c}, \hat{I}||\hat{r}_\mathcal{T})$ to R in this case. Note that, as $\hat{I}$ and $\hat{r}_\mathcal{T}$ are truly random strings in this case, if $\mathcal{S}_2$ makes the same oracle query $\mathsf{SendR}(\hat{c}, \hat{I}||\hat{r}_\mathcal{T})$ to the reader R, with overwhelming probability R will abort this session and outputs "0" indicating "reject".)

    Case-3.1.1 If R is running an incomplete session of session-identifier $\hat{c}$ (i.e., $\hat{c}$ was sent by R upon an InitReader query and R is waiting for the second-round message), $\mathcal{S}_2$ just returns back a random string $r_R \in_R \{0,1\}^\kappa$ to $\mathcal{A}_2$, and outputs "1" indicating "accept".

    Case-3.1.2. Otherwise, $\mathcal{S}_2$ simply returns back a special symbol "$\perp$" indicating invalid query.

    Case-3.2. In all other cases, $\mathcal{S}_2$ makes the same oracle query $\mathsf{SendR}(\hat{c}, \hat{I}||\hat{r}_\mathcal{T})$ to the reader R, and relays back the answer from R to $\mathcal{A}_2$.

4. On oracle query $\mathsf{SendT}(challenge, \hat{r}_R)$, where the challenge tag $\mathcal{T}_g$ (simulated by $\mathcal{S}_2$) currently runs a session of partial session-transcript $(\hat{c}, I||r_\mathcal{T})$ and is waiting for the third-round message, $\mathcal{S}_2$ works as follows:

    Case-4.1. If there exists a matching session of the same session transcript $(\hat{c}, I||r_\mathcal{T}, \hat{r}_R)$ at the side of R (where $\hat{r}_R$ may be simulated by $\mathcal{S}_2$ as in the above Case-3.1), $\mathcal{S}_2$ outputs "1" indicating "accept";

    Case-4.2. Otherwise, $\mathcal{S}_2$ simply outputs "0" indicating "reject".

5. Output of $\mathcal{S}_2$: Finally, whenever $\mathcal{A}_2$ stops, $\mathcal{S}_2$ also stops and outputs the simulated view, denoted $sview(\kappa, \ell)$, which consists of: the system public parameters $para$ (output by $\mathsf{Setup}(\kappa, \ell)$), the random coins used by $\mathcal{A}$ (actually set by $\mathcal{S}$), all oracle answers (including ones provided by the real protocol participants and ones simulated by $\mathcal{S}_2$) to queries made by $\mathcal{A}$. Recall that the output of the experiment $\mathbf{Exp}_{\mathcal{S}}^{zkp}[\kappa, \ell]$ is defined to be $(g, sview(\kappa, \ell))$.

It is clear that if $\mathcal{A}$ works in polynomial-time, $\mathcal{S}$ works also in polynomial-time. For all sufficiently large κ, any ℓ (that is polynomial in κ), any PPT adversary $\mathcal{A}$, and any polynomial-time distinguisher D, denote by $p_{\mathcal{S}} = \Pr[D(\kappa, \ell, g, sview(\kappa, \ell)) = 1]$ and by $p_{\mathcal{A}} = \Pr[D(\kappa, \ell, g, view_{\mathcal{A}}(\kappa, \ell)) = 1]$. Below, we show that $|p_{\mathcal{S}} - p_{\mathcal{A}}|$ is negligible, which establishes the zk-privacy property.

For any (κ, ℓ) and any PPT adversary $\mathcal{A} = (\mathcal{A}_1, \mathcal{A}_2)$, we first consider a mental experiment $\mathbf{Exp}_{\hat{S}}^{prf}(\kappa, \ell)$ w.r.t. a PPT algorithm $\hat{S}$. The algorithm $\hat{S}$ has oracle access to a PRF F_k or a truly random function $H : \{0,1\}^{2\kappa} \to \{0,1\}^{2\kappa}$, and works as follows:

Step-1. Select $i \in_R \{1, \cdots, \ell\}$ uniformly at random.

Step-2. Setup $\ell - 1$ tags, denoted $\mathcal{T}' = \{\mathcal{T}_1, \cdots, \mathcal{T}_{i-1}, \mathcal{T}_{i+1}, \cdots, \mathcal{T}_\ell\}$, with secret-keys $k_1, \cdots, k_{i-1}, k_{i+1}, \cdots, k_\ell$ and counters $ctr_1, \cdots, ctr_{i-1}, ctr_{i+1}, \cdots, ctr_\ell$ respectively, where each secret-key is a random string of length κ and each counter is initiated to be "1"; $\hat{S}$ also maintains a counter ctr_i to simulate $\mathcal{T}_i$ with the aid of its oracle (i.e., the PRF F_k or a truly random function); $\hat{S}$ also setups and maintains all the database records (of the reader R) w.r.t. the tags in $\mathcal{T}'$.

Step-3. Run $\mathcal{A}_1$ as a subroutine, and answer the oracle queries made by $\mathcal{A}_1$ as follows:
   – For oracle queries directed to tags in $\mathcal{T}'$, $\hat{S}$ computes and gives the answers by itself. Note that $\hat{S}$ can perfectly simulate the actions of tags in $\mathcal{T}'$.
   – For oracle queries directed to the tag $\mathcal{T}_i$, $\hat{S}$ mimics the actions of $\mathcal{T}_i$, and provides the oracle answer, with the aid of its oracle (i.e., the PRF F_k or a truly random function H).
   – For any oracle query directed to the reader R, $\hat{S}$ works as follows: If the oracle query is $\mathsf{InitReader}()$, $\hat{S}$ simply returns back a random string $c \in_R \{0,1\}^\kappa$. If the oracle query is $\mathsf{SendR}(\hat{c}, \hat{I} || \hat{r}_\mathcal{T})$, $\hat{S}$ mimics the actions of R with the following modifications: the computation involving database records w.r.t. tags in $\mathcal{T}'$ is performed by $\hat{S}$ itself; the computation involving the record w.r.t. $\mathcal{T}_i$ is performed by $\hat{S}$ with the aid of its oracle (i.e., the PRF or a truly random function).

Step-4. When $\mathcal{A}_1$ stops outputting $(st, \mathcal{C})$, where $\mathcal{C} = \{\mathcal{T}_{i_1}, \cdots, \mathcal{T}_{i_\delta}\}$ is a set of clean tags, $1 \le \delta \le \ell$, $\hat{S}$ selects $g \in_R \{1, \cdots, \delta\}$ uniformly at random. If the challenge-tag $\mathcal{T}_g(= \mathcal{T}_{i_g}) \ne \mathcal{T}_i$, $\hat{S}$ stops and outputs "$\perp$" indicating "failure"; otherwise (i.e., $\mathcal{T}_g = \mathcal{T}_i$), $\hat{S}$ continues to do the next step.

Step-5. In case $\mathcal{T}_g = \mathcal{T}_i$, $\hat{S}$ further runs $\mathcal{A}_2(st)$ as a subroutine and mimics the actions of $\mathcal{S}_2$, but with the following modifications: (Recall that $\mathcal{T}_i$ is just the challenge tag $\mathcal{T}_g$ in this case.)
   – For oracle queries made by $\mathcal{A}_2$ to tags in $\hat{\mathcal{T}} = \mathcal{T} - \mathcal{C}$, $\hat{S}$ computes and gives the answers by itself. Note that $\hat{S}$ can perfectly simulate the actions of tags in $\mathcal{T}' \supseteq \hat{\mathcal{T}}$.
   – For oracle queries directed to the challenge tag $\mathcal{T}_g = \mathcal{T}_i$ of the form $\mathsf{SendT}(challenge, \cdot)$, $\hat{S}$ provides the answer, mimicking the actions of $\mathcal{T}_i$, with the aid of its oracle (i.e., the PRF F_k or a truly random function H).
   – For oracle queries directed to the reader R, $\hat{S}$ works in the same way as in its first stage.
   – Output of $\hat{S}$: Whenever $\mathcal{A}_2$ stops, $\hat{S}$ also stops and then gets a view, denoted $view_{\mathcal{A}}^{\hat{S}}(\kappa, \ell)$, which consists of: the system public parameters $para$ (generated by $\hat{S}$ itself), the random coins used by $\mathcal{A}$ (actually set by $\hat{S}$), all oracle answers (provided by $\hat{S}$ itself to queries made by $\mathcal{A}$). Finally, $\hat{S}$ runs the assumed distinguisher D on inputs $(\kappa, \ell, g, view_{\mathcal{A}}^{\hat{S}}(\kappa, \ell))$, and outputs whatever $D(\kappa, \ell, g, view_{\mathcal{A}}^{\hat{S}}(\kappa, \ell))$ outputs.

Denote by $p_{\hat{S}}^{F_k}$ (resp., $p_{\hat{S}}^{H}$) the probability that $\hat{S}$ outputs "1", with oracle access to the PRF F_k (resp., a truly random function H). Note that $\hat{S}$ outputting "1" implies $\mathcal{T}_g = \mathcal{T}_i$ (as otherwise $\hat{S}$ will output "$\perp$" indicating failure). By the pseudorandomness of F_k, we get that for some quantity $\epsilon_{\hat{S}}$ that is negligible in κ, it holds that (for sufficiently large κ):

$$|p_{\hat{S}}^{F_k} - p_{\hat{S}}^{H}| = \epsilon_{\hat{S}} \tag{1}$$

Note also that the real view of $\mathcal{A}$, i.e., $view_{\mathcal{A}}$, and the view of $\mathcal{A}$ under the simulation of $\hat{S}$ in the experiment $\mathbf{Exp}_{\hat{S}}^{prf}(\kappa, \ell)$, i.e., $view_{\mathcal{A}}^{\hat{S}}$, are identical *conditioned on $\hat{S}$ gets oracle access to the PRF F_k and $\hat{S}$ did not output "$\perp$" at Step-4 (i.e., $\mathcal{T}_g = \mathcal{T}_i$)*. In particular, the view (particularly, the output of $\mathcal{C}$) of $\mathcal{A}_1$ is independent of the choices of i in this case, and thus $\Pr[\mathcal{T}_i = \mathcal{T}_g] = \frac{1}{\ell}$, and we get:

$$p_{\hat{S}}^{F_k} = \frac{1}{\ell} p_{\mathcal{A}} \tag{2}$$

The above analysis also implies that the adaptive completeness and mutual authentication properties hold also in the experiment $\mathbf{Exp}_{\hat{S}}^{prf}(\kappa, \ell)$, whether $\hat{S}$ gets oracle access to the PRF F_k or a truly random function H; Otherwise, the pseudorandomness property of the underlying PRF F_k will be violated.

Now, we consider another experiment $\mathbf{Exp}_{S'}^{prf}(\kappa, \ell)$, in which a PPT algorithm S', with oracle access to the PRF F_k or a truly random function H, mimics the actions of $\hat{S}$ until Step-5 in the experiment $\mathbf{Exp}_{\hat{S}}^{prf}(\kappa, \ell)$, but with the following modifications at Step-5 (recall that Step-5 actions imply $\mathcal{T}_g = \mathcal{T}_i$):

D1. For any oracle query directed to the challenge tag $\mathcal{T}_g = \mathcal{T}_i$ of the form $\mathsf{SendT}(challenge, \hat{c})$, $\hat{S}$ mimics the actions of the simulator $\mathcal{S}$ and returns back $I, r_{\mathcal{T}} \in_R \{0,1\}^{\kappa}$, where $I, r_{\mathcal{T}}$ are taken randomly by S' itself.

D2. For any oracle query directed to the reader R of the form $\mathsf{SendR}(\hat{c}, \hat{I}||\hat{r}_{\mathcal{T}})$, S' works as follows:

Case-D2.1 In Case-3.1 (of the simulation of $\mathcal{S}$), S' just mimics the actions of the simulator $\mathcal{S}$.

Case-D2.2 In Case-3.2 (of the simulation of $\mathcal{S}$), S' simulates the actions of R as follows:

Case-D.2.2.1. The computation involving database records w.r.t. tags in $\mathcal{T}'$ is performed by S' itself (just as $\hat{S}$ does), which perfectly mimics the actions of the reader R in this case.

Case-D.2.2.2. But, the computation involving the record w.r.t. $\mathcal{T}_i$ is just *ignored* by S'.

D3. For any oracle query directed to $\mathcal{T}_g = \mathcal{T}_i$ of the form $\mathsf{SendT}(challenge, \hat{r}_R)$, S' just mimics the actions of the simulator $\mathcal{S}$ (in Case-4.1 and Case-4.2 of the simulation of $\mathcal{S}$).

Note that S' only queries its oracle (i.e., the PRF F_k or a truly random function H) in its first-stage of simulation (in dealing with $\mathcal{A}_1$), but never queries its oracle in the second-stage simulation (in dealing with $\mathcal{A}_2$).

Denote by $p_{S'}^{F_k}$ (resp., $p_{S'}^{H}$) the probability that S' outputs "1", with oracle access to the PRF F_k (resp., a truly random function H) in the experiment $\mathbf{Exp}_{S'}^{prf}$. By the pseudorandomness of F_k, we get that for some quantity $\epsilon_{S'}$ that is negligible in κ, it holds that (for sufficiently large κ):

$$|p_{S'}^{F_k} - p_{S'}^{H}| = \epsilon_{S'} \tag{3}$$

Now, we investigate the differences between the view of $\mathcal{A}$ under the simulation of $\mathcal{S}$ and the view of $\mathcal{A}$ under the simulation of S', *conditioned on S' gets oracle access to the PRF F_k and S' does not output "$\perp$" at Step-4 (i.e., $\mathcal{T}_g = \mathcal{T}_i$)*. The only difference is: in Case-D.2.2.2 above, the computation involving the record w.r.t. $\mathcal{T}_i = \mathcal{T}_g$ is always ignored by S'; but in the simulation of $\mathcal{S}$ this computation will be performed by the reader R. The observation here is: Case-D.2.2 (that corresponds to Case-3.2 of the simulation of $\mathcal{S}$) means that, when $\mathcal{A}_2$ makes the oracle query $\mathsf{SendR}(\hat{c}, \hat{I}||\hat{r}_{\mathcal{T}})$, $\mathcal{T}_g = \mathcal{T}_i$ did not run or is running a session of the (partial) session transcript $(\hat{c}, \hat{I}||\hat{r}_{\mathcal{T}})$. Just analogue to the analysis of tag-to-reader authentication (presented in Section A.2), it is straightforward to calculate that the probability that the reader R successfully identifies $\mathcal{T}_i$ in Case-3.2 of the simulation of $\mathcal{S}$ is negligible. Under this observation and by noting that the view of $\mathcal{A}_1$ is independent of the choice of i when S' gets oracle access to F_k (and thus $\Pr[\mathcal{T}_g = \mathcal{T}_i] = \frac{1}{\ell}$), we have that for some quantity ϵ_S that is negligible in κ, it holds that (for sufficiently large κ):

$$p_{S'}^{F_k} = \frac{1}{\ell} p_{\mathcal{S}} + \epsilon_S \tag{4}$$

Now, we investigate the differences between the view of $\mathcal{A}$ under the simulation of $\hat{S}$ and that under the simulation of S', when both $\hat{S}$ and S' get access to a truly random function H. We have the following observations:

– In Case-D.1 (resp., Case-D2.1), S' always returns back random and independent $I, r_{\mathcal{T}}$ (resp., r_R). As $\mathcal{T}_i$ always generates I with the updated counter value, the value I return back by $\hat{S}$ in Case-D.1, with the aid of the truly

random function H, is also always uniformly distributed over $\{0,1\}^\kappa$ and is independent of what previously sent by $\hat{S}$. But the value $r_\mathcal{T}$ (resp., r_R) returned back by $\hat{S}$ in Case-D.1 (resp., Case-D.2.1), with the aid of the truly random function H, might not be independent of the values previously sent by $\hat{S}$, for the following reasons: (1) In case I is equal to some I-values previously sent by $\hat{S}$, the value $H(c||I) = (r_0, r_1)$ (that determines $r_\mathcal{T} = r_0 \oplus ctr||pad_2$ and $r_R = r_1$) can be previously defined. Note that the adversary $\mathcal{A}$ can use the same $\hat{c}$ for all sessions with the tags. This event occurs with probability at most $s \cdot 2^{-\kappa}$, where s is the upper-bound of sessions involved by $\mathcal{T}_i$; (2) For Case-D.2.1, upon the request $\mathsf{SendR}(\hat{c}, \hat{I}||\hat{r}_\mathcal{T})$ where $\hat{I}||\hat{r}_\mathcal{T}$ was sent by $\mathcal{T}_g = \mathcal{T}_i$ (simulated by $\hat{S}$) in session $\hat{c}$, the reader R (simulated by $\hat{S}$) may identify a tag $\mathcal{T}_{i'} \neq \mathcal{T}_i$, and thus the value r_R returned by $\hat{S}$ may be $F^1_{k'}(\hat{c}||I)$ (rather than an independent and random string in $\{0,1\}^\kappa$), where k' is the secret-key of $\mathcal{T}_{i'}$. But, as the adaptive completeness property holds in the experiment $\mathbf{Exp}^{prf}_{\hat{S}}(\kappa, \ell)$, this event occurs also with negligible probability.

- By the mutual authentication (specifically, the tag-to-reader authentication in this case) of the experiment $\mathbf{Exp}^{prf}_{\hat{S}}(\kappa, \ell)$, the difference caused by Case-D.2.2.2 in $\mathbf{Exp}^{prf}_{S'}(\kappa, \ell)$ occurs also with negligible probability.

- In Case-D.3 in $\mathbf{Exp}^{prf}_{S'}(\kappa, \ell)$ (that corresponds to Case-4.1 and Case-4.2 of the simulation of $\mathcal{S}$) , $\mathcal{T}_g = \mathcal{T}_i$ (simulated by S') always outputs "accept" (resp., "reject") in Case-4.1 when existing matching session at the side of R (resp., Case-4.2 when existing no matching session at R). But, in the simulation of $\hat{S}$ in $\mathbf{Exp}^{prf}_{\hat{S}}(\kappa, \ell)$, with oracle access to a truly random function H, $\mathcal{T}_g = \mathcal{T}_i$ may output "reject" in Case-4.1 (resp., "accept" in Case-4.2).
  We first observe that, by the mutual authentication (specifically, the reader-to-tag authentication in this case) of the experiment $\mathbf{Exp}^{prf}_{\hat{S}}(\kappa, \ell)$, the probability that $\mathcal{T}_g = \mathcal{T}_i$ simulated by $\hat{S}$, with oracle access to the truly random function H, outputs "accept" in Case-4.2 is negligible.
  Next, note that $\mathcal{T}_g = \mathcal{T}_i$ (simulated by $\hat{S}$ in $\mathbf{Exp}^{prf}_{\hat{S}}(\kappa, \ell)$) outputs "reject" in Case-4.1, only if the reader R (simulated by $\hat{S}$) identified a tag $\mathcal{T}_{i'} \neq \mathcal{T}_i$ upon the query $\mathsf{SendR}(\hat{c}, I||r_\mathcal{T})$, where $I||r_\mathcal{T}$ was sent by $\mathcal{T}_i$ in the session $\hat{c}$. But, by the adaptive completeness of the experiment $\mathbf{Exp}^{prf}_{\hat{S}}(\kappa, \ell)$, this event occurs also with negligible probability.

From above discussions, we have that for some quantity ϵ_H that is negligible in κ, it holds (for sufficiently large κ):

$$|p^H_{S'} - p^H_{\hat{S}}| = \epsilon_H \tag{5}$$

Combining all the equations (1), (2), (3), (4), (5), we have: $|p^{F_k}_{\hat{S}} - p^H_{\hat{S}}| = |\ell^{-1}p_\mathcal{A} - p^H_{\hat{S}}| = |\ell^{-1}p_\mathcal{A} - p^H_{S'} \pm \epsilon_H| = |\ell^{-1}p_\mathcal{A} - p^{F_k}_{S'} \pm \epsilon_{S'} \pm \epsilon_H| = |\ell^{-1}p_\mathcal{A} - \ell^{-1}p_S \pm \epsilon_S \pm \epsilon_{S'} \pm \epsilon_H| = \epsilon_{\hat{S}}$. From this equation, we conclude that $|p_\mathcal{A} - p_S|$ must be negligible, which then establishes the zk-privacy property of the protocol depicted in Figure 5.