On the Complexity of Scrypt and Proofs of Space in the Parallel Random Oracle Model

Joël Alwen[1], Binyi Chen[2], Chethan Kamath[1], Vladimir Kolmogorov[1],
Krzysztof Pietrzak[1], and Stefano Tessaro[2]

[1] IST Austria
[2] UC Santa Barbara

Abstract. We investigate lower bounds in terms of time and memory on the *parallel* complexity of an adversary $\mathcal{A}$ computing labels of randomly selected challenge nodes in direct acyclic graphs, where the w-bit label of a node is the hash h (modelled as a random oracle with w-bit output) of the labels of its parents. Specific instances of this general problem underlie both proofs-of-space protocols [Dziembowski et al. CRYPTO'15] as well as memory-hardness proofs including $\mathtt{scrypt}$, a widely deployed password hashing and key-derivation function which is e.g. used within Proofs-of-Work for digital currencies like Litecoin.

Current lower bound proofs for these problems only consider *restricted* algorithms $\mathcal{A}$ which perform only a single h query at a time and which only store individual labels (but not arbitrary functions thereof). This paper substantially improves this state of affairs;

Our first set of results shows that even when allowing multiple parallel h queries, the "cumulative memory complexity" (CMC), as recently considered by Alwen and Serbinenko [STOC '15], of $\mathtt{scrypt}$ is at least $w \cdot (n/\log(n))^2$, when $\mathtt{scrypt}$ invokes h n times. Our lower bound holds for adversaries which can store (1) Arbitrary labels (i.e., random oracle outputs) and (2) Certain natural functions of these labels, e.g., linear combinations. The exact power of such adversaries is captured via the combinatorial abstraction of parallel "entangled" pebbling games on graphs, which we introduce and study.

We introduce a combinatorial quantity γ_n and under the conjecture that it is upper bounded by some constant, we show that the above lower bound on the CMC also holds for arbitrary algorithms $\mathcal{A}$, storing in particular arbitrary functions of their labels. We also show that under the same conjecture, the *time complexity* of computing the label of a random node in a graph on n nodes (given an initial kw-bit state) reduces tightly to the time complexity for entangled pebbling on the same graph (given an initial k-node pebbling). Under the conjecture, this solves the main open problem from the work of Dziembowski et al.

In fact, we note that every non-trivial upper bound on γ_n will lead to the first non-trivial bounds for general adversaries for this problem.

0   On Conjectures 13 and 16

In this Section we track the current state on our Conjecture 13 and the weaker Conjecture 16. A counterexample to Conjecture 13 (which states that for some

constant γ we have $\gamma_n \leq \gamma$ for all $n \in \mathbb{N}$) has been found by Michal Rolinek (IST Austria) who showed $\gamma_n \geq \log(n)$, independently Daniel Malinowski and Karol Żebrowski (University of Warsaw) proved gave an even stronger $\gamma_n \geq \sqrt{n}/polylog(n)$ counterexample.

More interestingly, Daniel and Karol give a counterexample to the weaker Conjecture 16 (via fractional chromatic numbers of Kneser graphs) showing that $\Gamma_n \geq \sqrt{\log(n)}$. This is only slightly superconstant, but already shows that our approach will not give a "lossless" transition from pebbling to the parallel ROM model. This still leaves the possibility of proving a $polylog(n)$ $upper$ bound on Γ_n, which would be highly interesting.

On the positive side, several groups have reported progress giving direct proofs (i.e., not using any of the Conjectures) for the security of `scrypt` and proofs of space.

1   Introduction

The common denominator of *password hashing* and *proofs-of-work* is that they both require a certain computation to be sufficiently expensive, while still remaining feasible. In this domain, "expensive" has traditionally meant high *time complexity*, but recent hardware advances have shown this requirement to be too weak, with fairly inexpensive tailored-made ASIC devices for Bitcoin mining and password cracking gaining increasingly widespread usage.

In view of this, a much better requirement is *memory-hardness*, i.e., the *product* of the memory (a.k.a. space) *and* the time required to solve the task at hand should be large. This is often referred to as the *space-time* (ST) complexity which in turn has been used as an estimate of the product of the area and the time (AT) complexity of a circuit solving the task. Thus increasing ST complexity of a task appears to incur a higher dollar cost for building custom circuits to solve the task compared to increase in dollar cost obtained by simply increasing the required raw computing power alone. This observation motivated Percival to introduce the concept of memory-hardness along with a candidate memory-hard function `scrypt` in [Per09]. In particular, memory-hardness was one of the main desiderata in the recent password-hashing competition, and its winner, Argon2 [BDK15], has been designed to be memory-hard. Dziembowski *et al* [DFKP15] introduce the concept of *proofs of space* (PoSpace), where the worker (or miner) can either dedicate a large amount of storage space, and then generate proofs extremely efficiently, or otherwise must pay a large time cost for every proof generated. The PoSpace protocol has also found its way into a recent proposal for digital currency [PPA$^+$15].

Our contributions, in a nutshell. Cryptanalytic attacks [BK15,AS15,CJM14] targeting candidates for memory-hard functions [FLW13,AABS14] have motivated the need for developing constructions with *provable security guarantees*. Moreover security proofs complement cryptanalytic results by highlighting which classes of attacks remain worth considering and which can be ruled out. With the exception of [AS15], most candidate memory-hard functions come without

security proofs and those that do (e.g. [Per09,FLW13]) only consider a severely restricted class of algorithms while using non-amortized notions of memory-hardness which don't capture the memory requirments of computing these functions on multiple inputs; a key requirment when considering their intended use for, say, brute-force resistant password hashing. A primary goal of this paper is to advance the foundations of memory-hardness, and we make progress along several fronts.

We develop a new class of randomized pebbling games on graphs – called *entangled pebbling games* – which are used to prove results on the memory-hardness of tasks such as computing `scrypt` for large non-trivial classes of adversaries. Moreover, we show how to boost these results to hold against *arbitrary* adversaries in the parallel random oracle model (pROM) of Alwen and Serbinenko [AS15] under the conjecture that a new combinatorial quantity which we introduce is (sufficiently) bounded.

A second application of the techniques introduced in this paper considers Proofs of Space. We show that time lower bounds on the pebbling complexity of a graph imply time lower bounds in the pROM model agains *any* adversary. The quantitative bounds we get depend on the combinatorial value we introduce, and assuming our conjecture, are basically tight. This solves, modulo the conjecture, the main problem left open in the Proofs of Space paper [DFKP15]. As we'll discuss in more detail below, pebbling lower bounds have been used extensively for lower bounding time/space complexity in the random oracle model before, but in most cases the adversary did not have any state depending on the random oracle, and the underlying graph was fixed and known. In this case, a standard "ex-post facto" argument suffices to translate pebbling lower bounds to time/memory lower bounds in the random oracle model. Once a state depending on the RO is available (as in proofs of space), or the graph is data dependent (as in scrypt), this standard technique fails, and currently in these cases bounds in the RO are only known against very restricted classes of adversaries which just store plain labels.

Sequentially memory-hard functions. Percival [Per09] gave the first proposal of a memory-hard hash function, called `scrypt`. The function has been well received in practice. For example it is currently in the process of being standerdized by the IETF [PJ] as a Key Derivation Function and also underlies the Proof of Work protocols of several cryptocurrencies e.g. Lightcoin [Lee11], one of the currently most prevelant cryptocurrencies in terms of market capitalization [Mar]. `scrypt`[3] uses a hash function $h : \{0,1\}^* \rightarrow \{0,1\}^w$ (e.g., SHA-256), and proceeds in two phases, given an input X. It first computes $X_i = h^i(X)$ for all $i \in [n]$, and with

---

[3] In fact, what we describe here is only a subset of the whole `scrypt` function, called ROMix. ROMix is the actual core of the `scrypt` function, and we will use the generic name "scrypt" for in the following. ROMix (with some minor modification and extensions) also underlies one of the two variants of the winner Argon [BDK15] of the recent password hashing competition `https://password-hashing.net/`, namely the data-dependent variant Argon2d.

$S_0 = X_n$, it then computes $S_1, \ldots, S_n$ where

$$S_i = \mathsf{h}(S_{i-1} \oplus X_{\mathsf{int}(S_{i-1})})$$

where $\mathsf{int}(S)$ reduces an w-bit string S to an integer in $[n]$. The final output is S_n. Note that is possible to evaluate $\mathtt{scrypt}$ on input X using $n \cdot w$ bits of memory and in linear time in n, by keeping the values $X_1, \ldots, X_n$ stored in memory once they are computed. However, the crucial point is that there is no apparent way to save memory – for example, to compute S_i, we need to know $X_{\mathsf{int}(S_{i-1})}$, and under the assumption that $\mathsf{int}(S_{i-1})$ is (roughly) random in $[n]$, an evaluator without memory needs to do linear work (in n) to recover this value before continuing with the execution. This gives a constant-memory, $O(n^2)$ time algorithm to evaluate $\mathtt{scrypt}$. In fact, as pointed out already in Percival [Per09], the actual expectation is that no matter how much time $T(n)$ and how much memory $S(n)$ an adversarial evaluator invests, we always have $S(n){\cdot}T(n) \geq n^{2-\epsilon}$ for all $\epsilon > 0$, *even if the evaluator can parallelize its computation arbitrarily.*

Percival's analysis of $\mathtt{scrypt}$ assumes that h is a random oracle. The analysis is limited in two ways: (1) It only considers adversaries which can only store random oracle outputs in their memory. (2) The bound measures memory complexity in terms of the *maximum* memory resources $S(n)$. The latter is undesirable, since the ultimate goal of an adversary performing a brute-force attack is to evaluate $\mathtt{scrypt}$ on *as many inputs as possible*, and if the large memory usage is limited to a small fraction of the computing time, a much higher *amortized complexity* can be achieved.

Alwen and Serbinenko [AS15] recently addressed these shortcomings, and delivered provably sequentially memory-hard functions in the so-called *parallel random oracle model* (pROM), developing new and better complexity metrics tailored to capturing amortized hardness. While their work falls short of delivering guarantees for $\mathtt{scrypt}$-like functions, it serves as an important starting point for our work, and we give a brief overview.

From sequential memory-hardness to pebbling. Their work explicitly consider adversaries attempting to evaluate a function $\mathcal{H}^{\mathsf{h}}$ (which make calls to some underlying hash function h, modeled as a random oracle) which proceed in multiple rounds, and in each round i, the adversary can make an *unbounded* number of *parallel* queries to h, and then pass on a state σ_i to the next round. The complexity of the adversary is captured by its *cumulative memory complexity* (CMC) given by $\sum_i |\sigma_i|$. One then denotes as $\mathsf{cmc}^{\mathsf{pROM}}(\mathcal{H})$ the expected CMC of the best adversary where the expectation is over the choice of RO h and coins of the adversary. We stress that CMC exhibits some very important features: *First,* a lower bound appears to yield a reasonable lower bound on the AT complexity metric. *Second,* In contrast to the ST complexity the CMC of a task also gives us a lower-bound on the electricity consumption of performing the task. This is because storing data in volatile memory for, say, the time it takes to evaluate h consumes a significant amount of electricity. Thus CMC tells us something not only about the dollar cost of building a custom circuit for computing a task

but also about the dollar cost of actually running it. While the former can be amortized over the life of the device, the later represents a recurring fee.

A natural candidate for a sequentially memory-hard function studied in [AS15] is defined by a single source and sink directed acyclic graph (DAG) $G = (V, E)$ and a hash function h. The label of a vertex $i \in V$ with parents $\{p_1, \ldots, p_d\}$ (i.e., $(p_j, v) \in E$ for $i = 1, \ldots, d$) is defined as $\ell_i = \mathsf{h}(i, \ell_{p_1}, \ldots, \ell_{p_d})$. Note that the labels of all vertices can be recursively computed starting with the sources. The function $\mathsf{label}(G, \mathsf{h})$ is now simply the label ℓ_v of the sink v. There is a natural connection between $\mathsf{cmc}^{\mathsf{pROM}}(\mathsf{label}(G, \mathsf{h}))$ for a randomly chosen h and the *cumulative pebbling complexity* (CC) of the graph G. This connection was first exploited in the to construct functions for which evaluation requires many cache memory in [DNW05] and more recently to build one-time computable functions [DKW11] as well as in the security proof the memory-hard function in [FLW13]. CC is defined in a game where one can place pebbles on the vertices of V, according to the following rules: In every step of the game, a new pebble can be placed on any vertex for which all parents of v have pebbles on them (in particular, pebbles can always be placed on sources), and pebbles can always be removed. The game is won when a pebble has been placed on the sink. The CC of a strategy for pebbling G is defined as $\sum_i |S_i|$, where S_i is the set of vertices on which a pebble is placed at the end of the i^{th} step, and the cc of G – denoted $\mathsf{cc}(G)$ – is the CC of the best strategy.

Indeed, $\mathsf{cc}(G)$ naturally captures the CMC of *restricted* pROM adversaries computing $\mathsf{label}(G, \mathsf{h})$ for which every state σ_i *only consists of random oracle outputs*, i.e., of vertex labels. A pebble on v is equivalent to the fact that σ_i contains ℓ_v. However, a full-fledged pROM adversary *has no reason to be restricted to such a strategy* – it could for example store as part of its state σ_i a particular encoding of the information accumulated so far. Nonetheless, the main result of [AS15] shows that (up to a negligible extent) such additional freedom does *not* help in computing $\mathsf{label}(G, \mathsf{h})$.

> **Informal theorem.** For every DAG $G = (V, E)$, with overwhelming potability (over the choice of h from the family of all functions) we have (w denotes the length of the labels ℓ_i)
>
> $$\mathsf{a\text{-}cmc}^{\mathsf{pROM}}(\mathsf{label}(G, \mathsf{h})) \approx w \cdot \mathsf{cc}(G)$$
>
> where $\mathsf{a\text{-}cmc}^{\mathsf{pROM}}$ is the CMC per labeling computed amortized across multiple labelings being computed simultanously.

This result is then complemented with a concrete class of constant-degree DAGs G_n on n vertices (which is efficiently constructible for every n) such that $\mathsf{cc}(G_n) = \Omega(n^2/\mathrm{polylog}(n))$.

Unfortunately however, the framework of [AS15] does not extend to `scrypt` like functions, for which we would hope to prove a similar $\Omega(n^2)$ lower bound on the CMC – indeed, the crucial point in functions like `scrypt` is that the values which need to be input to h are determined *at run-time* in a *data-dependent* fashion. While this makes the design far more intuitive, the proofs in [AS15]

crucially rely on the layout of the computation being laid out a priori in a *data-independent* fashion.

Our contributions. This paper supports the security of `scrypt`-like functions via two types of results – results for *restricted* adversaries, as well as results for arbitrary adversaries under a combinatorial conjecture. Our results also have direct implication on proofs of space, but we postpone this discussion below to ease presentation.

1) RANDOMIZED PEBBLING GAMES. As our first step, we introduce a generalization `pebble` of pebbling games on a DAG $G = (V, E)$ with dynamic challenges randomly sampled from a set $C \subseteq V$. With the same pebbling rules as before, we now proceed over n rounds, and at every round, a challenge c_i is drawn uniformly at random from C. The player's goal is to place a pebble on c_i, before moving to the next round, and learning the next challenge c_{i+1}. The game terminates when the last challenge has been covered by a pebble. One can similarly associate with G a *labeling game* `computeLabel` in the pROM, where the goal is instead to compute the *label* ℓ_{c_i} of c_i, rather than placing a pebble on it. For instance, the computation of `scrypt` is tightly connected to the `computeLabel` played on the chain graph L_n with vertices $[n] = \{1, 2, \ldots, n\}$, with edges $(i, i + 1)$ for $i \in [n - 1]$, and challenges $C = [n]$ (as detailed in Section 2.5). The labels to be computed in this game are those needed to advance the computation in the second half of the `scrypt` computation, and the challenges (in the actual `scrypt` function) are computed from hash-function outputs.

In fact, it is not hard to see that in `computeLabel` for some graph G a pROM adversary that only stores random-oracle generated outputs can easily be turned into a player for the `pebble` for graph G. This is particular true for $G = L_n$, and thus lower bounding the CC of an adversary playing `pebble` on L_n also yields a lower bound on the CMC of computing (the second half of) `scrypt`. Our first result provides such a lower bound.

> **Theorem 1.** For any constant $\delta > 0$, the CC of an adversary playing `pebble` on the chain graph L_n with challenges $[n]$ is $\Omega_\delta(n^2 / \log^2(n))$ with probability $1 - \delta$ over the choice of all challenges.[4]

To appreciate this result, it should be noted that it inherently relies on the choice of the challenges being *independent* of the adversary playing the game – indeed, *if the challenges are known a priori*, techniques from [AS15] directly give a strategy with CC $O(n^{1.5})$ for the above game. Also this result already improves on Percival's analysis (which, implictely, place similar restrictions on class of pROM algorithms considered), as the lower bound is uses the CC of (simple) pebbling of a graph, and thus it actually generalized to a lower bound on the *amortized* complexity of computing multiple `scrypt` instances in the pROM.[5]

---

[4] The subscript δ in Ω_δ denotes that the hidden constant depends on δ.

[5] This follows from a special case of the Lemma in [AS15] showing that CC of a graph is equal to the sum of the CCs the graphs disconnected components.

2) ENTANGLED PEBBLING. The above result is an important first step – to the best of our knowledge all known evaluation attacks against memory-hard functions indeed *only store hash labels directly or not at all* and thus fit into this model – but we ask the question whether the model can be strengthened. For example, an adversary could store the XOR $\ell_i \oplus \ell_j$ of two labels (which only takes w bits) and depending on possible futures of the game, recover both labels given any one of them. As we will see, this *can* help. As a middle ground between capturing pROM security for arbitrary adversaries and the above pebbling adversaries, we introduce a new class of pebbling games, called *entanglement pebbling games*, which constitutes a combinatorial abstraction for such adversaries.

In such games, an adversary can place on a set $\mathcal{Y} \subseteq V$ an "entangled pebble" $\langle \mathcal{Y} \rangle_t$ for some integer $0 \leq t \leq |\mathcal{Y}|$. The understanding here is that placing an individual pebble on any t vertices $v \in \mathcal{Y}$ – which we see as a special case of $\langle v \rangle_0$ entangled pebble – is equivalent to having individual pebbles on all vertices in $\mathcal{Y}$. The key point is that keeping an entangled pebble $\langle \mathcal{Y} \rangle_t$ costs *only* $|\mathcal{Y}| - t$, and depending on challenges, we may take different choices as to which t pebbles we use to "disentangle" $\langle \mathcal{Y} \rangle_t$. Also, note that in order to *create* such an entangled pebble, on all elements of $\mathcal{Y}$ there must be either an individual pebble, or such pebble can easily be obtained by disentangling existing entangled pebbles.

In the pROM labeling game, an entangled pebble $\langle \mathcal{Y} \rangle_t$ corresponds to an encoding of length $w \cdot (|\mathcal{Y}| - t)$ of the w-bit labels $\{\ell_i \; : \; i \in \mathcal{Y}\}$ such that given any t of those labels, we can recover all the remaining ones. Such an encoding can be obtained as follows: Fix $2d - t$ elements $x_1, \ldots, x_{2d-t}$ in the finite field $\mathbb{F}_{2^w}$. Let $\mathcal{Y} = \{y_1, \ldots, y_d\}$, and consider the (unique) degree $d - 1$ polynomial $p(.)$ over the finite field $\mathbb{F}_{2^w}$ (whose element are represented as w-bit strings) such that

$$\forall i \in [d] \; : \; p(x_i) = \ell_{y_i} \; .$$

The encoding now simply contains $\{p(x_{d+1}), \ldots, p(x_{2d-t})\}$, i.e., the evaluation of this polynomial on $d - t$ points. Note that given this encoding and any t labels $\ell_i, i \in \mathcal{Y}$, we have the evaluation of $p(.)$ on d points, and thus can reconstruct $p(.)$. Once we know $p(.)$, we can compute all the labels $\ell_{y_i} = p(i)$ in $\mathcal{Y}$.

We prove (in Appendix A) that in general, entangled pebbling is *strictly more powerful* (in terms of minimizing the expected CC) than regular pebbling by giving a concrete graph and set of challenge nodes for which no pebbling strategy using only unentangled pebbles can match a strategy we describe which uses entangled pebbles. Fortunately, we will also show that for the randomized pebbling game on the chain graph L_n entangled pebbling cannot outperform regular ones. We show:

Theorem 2. For any constant $\delta > 0$, the CC of an entangled pebbling adversary playing pebble on graph L_n is $\Omega(n^2 / \log^2(n))$ with probability $1 - \delta$ over the choice of all challenges.

Interestingly, the proof is a simple adaptation of the proof of for the non-entangled case. This result can again be interpreted as providing a guarantee

in the label game in the pROM for L_n for the class of adversaries that can be abstracted by entangled pebbling algorithms.

3) ARBITRARY ADVERSARIES. So far we have only discussed (entangled) pebbling lower bounds, which then imply lower bounds for restricted adversaries in the pROM model. In Section 4 we consider security against arbitrary adversaries. Our main results there show that there is a tight connection between the complexity of playinsolg computeLabel and a combinatorial quantity γ_n that we introduce. We show two results. The first lower-bounds the *time* complexity of playing computeLabel for *any* graph G while the second lower-bounds the *CMC* of playing computeLabel for L_n (and thus scrypt).

1. For any DAG $G = (V, E)$ with $|V| = n$, with high probability over the choice of the random hash function h, the pROM *time* complexity to play computeLabel for graph G, for any number of challenges, using h and when starting with any state of size $k \cdot w$ is (roughly) at least the time complexity needed to play pebble on G with the same number of challenges and starting with an initial pebbling of size roughly $\gamma_n \cdot k$.
2. The pROM CMC for pebble for L_n is $\Omega(n^2 / \log^2(n) \cdot \gamma_n)$.

At this point, we do not have any non-trivial upper bound on γ_n but we *conjecture* that γ_n is upper bounded by a constant γ (i.e., $\forall n \in \mathbb{N} : \gamma_n \leq \gamma$). The best lower bound we have is $\gamma_5 > 3/2$. Still, we note that our results are *concrete*, and remain meaningful for other values of γ_n. Indeed, we would get non-trivial statements even if γ_n were to *grow* moderately as a function of n, i.e. $\gamma_n = \text{polylog}(n)$ or $\gamma_n = n^\epsilon$ for some small $\epsilon > 0$.

Therefore, assuming our conjecture on γ_n, the first result in fact solves the main open problem from the work of Dziembowski et al [DFKP15] on proofs of space. The second result yields, in particular, a near-quadratic lower bound on the cc of evaluating scrypt for arbitrary pROM adversaries.

2 Pebbling, Entanglement, and the pROM

In this section, we first introduce different notions of pebbling graphs with challenges. In particular, we present both a notion of parallel pebbling of graphs with probabilistic challenges, and then extend this to our new notion of entangled pebbling games. Next, we discuss some generic relations between entangled and regular pebbling, before finally turning to defining the parallel random-oracle model (pROM), and associated complexity metrics.

Notation We use the following notation for common sets $\mathbb{N} := \{0, 1, 2, \ldots\}, \mathbb{N}^+ := \mathbb{N} \backslash \{0\}, \mathbb{N}_{\leq c} := \{0, 1, \ldots, c\}$ and $[c] := \{1, 2, \ldots, c\}$. For a distribution $\mathcal{D}$ we write $x \in \mathcal{D}$ to denote sampling x according to $\mathcal{D}$ in a random experiment.

2.1 Pebbling Graphs

Throughout, let $G = (V, E)$ denote a directed acyclic graph (DAG) with vertex set $V = [n]$. For a vertex $i \in V$, we denote by $\mathsf{parent}(i) = \{j \in V : (j, i) \in E\}$

the parents of i. The *m-round, probabilistic parallel pebbling game* between a player T on a graph $G = (V, E)$ with challenge nodes $C \subseteq V$ is defined as follows.

$\mathsf{pebble}(G, C, m, \mathsf{T}, P_{\mathsf{init}})$: The m-round parallel pebbling game on a DAG $G = (V, E)$ with challenge set $C \subseteq V$ and an initial pebbling configuration $P_{\mathsf{init}} \subseteq V$ is played between a challenger and a pebbler T.

1. Initialise $\mathsf{cnt} := 0$, $\mathsf{round} := 0$, $P_{\mathsf{cnt}} := P_{\mathsf{init}}$ and $\mathsf{cost} := 0$.
2. A challenge $c \leftarrow C$ is chosen uniformly from C and passed to T.
3. $\mathsf{cost} := \mathsf{cost} + |P_{\mathsf{cnt}}|$.
4. T choses a new pebbling configuration $P_{\mathsf{cnt}+1}$ which must satisfy

$$\forall i \in P_{\mathsf{cnt}+1} \setminus P_{\mathsf{cnt}} \;:\; \mathsf{parent}(i) \in P_{\mathsf{cnt}} \tag{1}$$

5. $\mathsf{cnt} := \mathsf{cnt} + 1$.
6. If $c \notin P_{\mathsf{cnt}}$ go to step 3.   *c not yet pebbled*
7. $\mathsf{round} := \mathsf{round} + 1$. If $\mathsf{round} < m$ go to step 2, otherwise if $\mathsf{round} = m$ the experiment is over, the output is the final count cnt and the cumulative cost cost.

The *cumulative black pebbling complexity* which is now defined as

$$\mathsf{cc}(G, C, m, \mathsf{T}, P_{\mathsf{init}}) := \mathop{\mathbb{E}}_{\mathsf{pebble}(G,C,m,\mathsf{T},P_{\mathsf{init}})} [\mathsf{cost}]$$

$$\mathsf{cc}(G, C, m, k) := \min_{\substack{\mathsf{T}, P_{\mathsf{init}} \subseteq V \\ |P_{\mathsf{init}}| \leq k}} \{\mathsf{cc}(G, C, m, \mathsf{T}, P_{\mathsf{init}})\}$$

We sometimes use the shorthand $\mathsf{cc}(G, C, m) = \max_k \mathsf{cc}(G, C, m, k)$. Similary, the time cost is defined as

$$\mathsf{time}(G, C, m, \mathsf{T}, P_{\mathsf{init}}) := \mathop{\mathbb{E}}_{\mathsf{pebble}(G,C,m,\mathsf{T},P_{\mathsf{init}})} [\mathsf{cnt}]$$

$$\mathsf{time}(G, C, m, k) := \min_{\substack{\mathsf{T}, P_{\mathsf{init}} \subseteq V \\ |P_{\mathsf{init}}| \leq k}} \{\mathsf{time}(G, C, m, \mathsf{T}, P_{\mathsf{init}})\}$$

The above notions consider the expected cost of a pebbling, thus even if, say $\mathsf{cc}(G, C, m, k)$, is very large, this could be due to the fact that for a tiny fraction of challenge sequences the complexity is very high, while for all other sequences it is very low. To get more robust security notions, we will define a more fine-grained notion which will guarantee that the complexity is high on all but some

ϵ fraction on the runs.

$$\mathsf{cc}_\epsilon(G, C, m, \mathtt{T}, P_{\mathsf{init}}) := \inf\left\{\gamma \;\middle|\; \mathop{\mathbb{P}}_{\mathsf{pebble}(G,C,m,\mathtt{T},P_{\mathsf{init}})}[\mathsf{cost} \geq \gamma] \geq 1 - \epsilon\right\}$$

$$\mathsf{cc}_\epsilon(G, C, m, k) := \min_{\substack{\mathtt{T}, P_{\mathsf{init}} \subseteq V \\ |P_{\mathsf{init}}| \leq k}} \left\{\mathsf{cc}_\epsilon(G, C, m, \mathtt{T}, P_{\mathsf{init}}\right\}$$

$$\mathsf{cc}_\epsilon(G, C, m) := \max_k \mathsf{cc}_\epsilon(G, C, m, k)$$

$$\mathsf{time}_\epsilon(G, C, m, \mathtt{T}, P_{\mathsf{init}}) := \inf\left\{\gamma \;\middle|\; \mathop{\mathbb{P}}_{\mathsf{pebble}(G,C,m,\mathtt{T},P_{\mathsf{init}})}[\mathsf{cnt} \geq \gamma] \geq 1 - \epsilon\right\}$$

$$\mathsf{time}_\epsilon(G, C, m, k) := \min_{\substack{\mathtt{T}, P_{\mathsf{init}} \subseteq V \\ |P_{\mathsf{init}}| \leq k}} \left\{\mathsf{time}_\epsilon(G, C, m, \mathtt{T}, P_{\mathsf{init}}\right\}$$

In general, we cannot upper bound cc in terms of cc_ϵ if $\epsilon > 0$ (same for time in terms of time_ϵ), but in the other direction it is easy to show that

$$\mathsf{cc}(G, C, m, \mathtt{T}, P_{\mathsf{init}}) \geq \mathsf{cc}_\epsilon(G, C, m, \mathtt{T}, P_{\mathsf{init}})(1 - \epsilon)$$

2.2 Entangled Pebbling

In the above pebbling game, a node is always either pebbled or not and there is only one type of pebble which we will hence forth refer to as a "black" pebble. We will now introduce a more general game, where $\mathtt{T}$ can put "entangled" pebbles.

A *t-entangled pebble*, denoted $\langle \mathcal{Y} \rangle_t$, is defined by a subset of nodes $\mathcal{Y} \subseteq [n]$ together with an integer $t \in \mathbb{N}_{\leq |\mathcal{Y}|}$. Having black pebble on all nodes $\mathcal{Y}$ now corresponds to the special case $\langle \mathcal{Y} \rangle_0$. Entangled pebbles $\langle \mathcal{Y} \rangle_t$ now have the following behaviour. Once any subset of $\mathcal{Y}$ of size (at least) t contains black pebbles then all $v \in \mathcal{Y}$ immediatly recieve a black pebble (regardless of whether their parents already contained black pebbles or not). We define the *weight* of an entangled pebble as:
$$w(\langle \mathcal{Y} \rangle_t) := |\mathcal{Y}| - t.$$

More generally, an *(entangled) pebbling configuration* is defined as a set $P = \{\langle \mathcal{Y}_1 \rangle_{t_1}, \ldots, \langle \mathcal{Y}_z \rangle_{t_s}\}$ of entangled pebbles and its weight is

$$w(P) := \sum_{i \in [s]} w(\langle \mathcal{Y}_i \rangle_{t_i}).$$

The rule governing how a pebbling configuration P_{cnt} can be updated to configuration $P_{\mathsf{cnt}+1}$ – which previously was the simple property eq.(1) – are now a bit more involved. To describe them formally we need the following definition.

Definition 1 (Closure). *The* closure *of an entangled pebbling configuration* $P = \{\langle \mathcal{Y}_1 \rangle_{t_1}, \ldots, \langle \mathcal{Y}_s \rangle_{t_s}\}$ *– denoted* $\mathsf{closure}(S)$ *– is defined recursively as follows: initialise* $\Lambda = \emptyset$ *and then*

$$\text{while } \exists j \in [s] \;:\; (\mathcal{Y}_j \not\subseteq \Lambda) \wedge (\Lambda \cap \mathcal{Y}_j \geq t_j) \text{ set } \Lambda := \Lambda \cup \mathcal{Y}_j$$

once Λ *cannot be further extended using the rule above we define* $\mathsf{closure}(S) = \Lambda$.

Note that $\mathsf{closure}(S)$ is non-empty iff there's at least one set of t-entangled pebbles $\langle \mathcal{Y} \rangle_t$ in P with $t = 0$. Equipped with this notion we can now specify how a given pebbling confugration can be updated.

Definition 2 (Valid Update). *Let $P = \{\langle \mathcal{Y}_1 \rangle_{t_1}, \ldots, \langle \mathcal{Y}_m \rangle_{t_s}\}$ be an entangled pebbling configuration. Further,*

- *Let $\mathcal{V}_1 := \mathsf{closure}(P)$.*
- *Let $\mathcal{V}_2 := \{i \; : \; \mathsf{parent}(i) \subseteq \mathcal{V}_1\}$. These are the nodes that can be pebbled using the black pebbling rules (eq.1).*

Now

$$P' = \{\langle \mathcal{Y}_1' \rangle_{t_1'}, \ldots, \langle \mathcal{Y}_{s'}' \rangle_{t_{s'}'}\}$$

is a valid update *of P if for every $\langle \mathcal{Y}_{j'}' \rangle_{t_{j'}'}$, one of the two conditions is satisfied*

1. *$\mathcal{Y}_{j'}' \subseteq (\mathcal{V}_1 \cup \mathcal{V}_2)$.*
2. *$\exists i$ with $\mathcal{Y}_{j'}' = \mathcal{Y}_i$ and $t_j' \geq t_i$. That is, $\langle \mathcal{Y}_{j'}' \rangle_{t_{j'}'}$ is an entangled pebble $\langle \mathcal{Y}_i \rangle_{t_i}$ that is already in P, but where we potentially have increased the threshold from t_i to $t_{j'}'$.*

The entangled pebbling game $\mathsf{pebble}^{\updownarrow}(G, C, m, \mathtt{T})$ is now defined like the game $\mathsf{pebble}(G, C, m, \mathtt{T})$ above, except that $\mathtt{T}$ is allowed to choose entangled pebblings:

$\mathsf{pebble}^{\updownarrow}(G, C, m, \mathtt{T}, P_{\mathsf{init}})$: The m-round parallel, entangled pebbling game on a DAG $G = (V, E)$ with challenge set $C \subseteq V$ and an initial entagled pebbling configuration P_{init}

1. Initialise $\mathsf{cnt} := 0$, $\mathsf{round} := 0$, $P_{\mathsf{cnt}} := P_{\mathsf{init}}$ and $\mathsf{cost} := 0$.
2. A challenge $c \leftarrow C$ is chosen uniformly from C and passed to $\mathtt{T}$.
3. $\mathsf{cost} := \mathsf{cost} + w(P_{\mathsf{cnt}})$.
4. $\mathtt{T}$ choses a new pebbling configuration $P_{\mathsf{cnt}+1}$ which must be a valid update of P_{cnt}.
5. $\mathsf{cnt} := \mathsf{cnt} + 1$.
6. If $c \notin \mathsf{closure}(P_{\mathsf{cnt}})$ go to step 3. \hfill *c not yet pebbled*
7. $\mathsf{round} := \mathsf{round}+1$. If $\mathsf{round} < m$ go to step 2, otherwise if $\mathsf{round} = m$ the experiment is over, the output is the final count cnt and the cumulative cost cost.

The *cumulative entangled pebbling complexity* and the *entangled time complexity* of this game are defined analogously to those of the simple pebbling

game.

$$\mathsf{cc}^{\updownarrow}(G,C,m,\mathsf{T},P_{\mathsf{init}}) := \mathop{\mathbb{E}}_{\mathsf{pebble}^{\updownarrow}(G,C,m,\mathsf{T},P_{\mathsf{init}})}[\mathsf{cost}]$$

$$\mathsf{cc}^{\updownarrow}(G,C,m,k) := \min_{\substack{\mathsf{T},P_{\mathsf{init}}\\ w(P_{\mathsf{init}})\leq k}} \{\mathsf{cc}(G,C,m,\mathsf{T},P_{\mathsf{init}})\}$$

$$\mathsf{cc}^{\updownarrow}\epsilon(G,C,m) := \max_k \mathsf{cc}^{\updownarrow}(G,C,m,k)$$

$$\mathsf{time}^{\updownarrow}(G,C,m,\mathsf{T},P_{\mathsf{init}}) := \mathop{\mathbb{E}}_{\mathsf{pebble}^{\updownarrow}(G,C,m,\mathsf{T},P_{\mathsf{init}})}[\mathsf{cnt}]$$

$$\mathsf{time}^{\updownarrow}(G,C,m,k) := \min_{\substack{\mathsf{T},P_{\mathsf{init}}\\ w(P_{\mathsf{init}})\leq k}} \{\mathsf{time}(G,C,m,\mathsf{T},P_{\mathsf{init}})\}$$

The more fine-grained versions are again defined much as before.

$$\mathsf{cc}^{\updownarrow}_{\epsilon}(\mathsf{G},C,m,\mathsf{T},P_{\mathsf{init}}) := \inf\left\{\gamma \,\middle|\, \mathop{\mathbb{P}}_{\mathsf{pebble}^{\updownarrow}(G,C,m,\mathsf{T},P_{\mathsf{init}})}[\mathsf{cost}\geq\gamma]\geq 1-\epsilon\right\}$$

$$\mathsf{cc}^{\updownarrow}_{\epsilon}(G,C,m,k) := \min_{\substack{\mathsf{T},P_{\mathsf{init}}\\ w(P_{\mathsf{init}})\leq k}} \left\{\mathsf{cc}^{\updownarrow}_{\epsilon}(G,C,m,\mathsf{T},P_{\mathsf{init}})\right\}$$

$$\mathsf{cc}^{\updownarrow}_{\epsilon}(G,C,m) := \max_k \mathsf{cc}^{\updownarrow}_{\epsilon}(G,C,m,k)$$

$$\mathsf{time}^{\updownarrow}_{\epsilon}(\mathsf{G},C,m,\mathsf{T},P_{\mathsf{init}}) := \inf\left\{\gamma \,\middle|\, \mathop{\mathbb{P}}_{\mathsf{pebble}^{\updownarrow}(G,C,m,\mathsf{T},P_{\mathsf{init}})}[\mathsf{cnt}\geq\gamma]\geq 1-\epsilon\right\}$$

$$\mathsf{time}^{\updownarrow}_{\epsilon}(G,C,m,k) := \min_{\substack{\mathsf{T},P_{\mathsf{init}}\\ w(P_{\mathsf{init}})\leq k}} \left\{\mathsf{time}^{\updownarrow}_{\epsilon}(G,C,m,\mathsf{T},P_{\mathsf{init}})\right\}$$

In Appendix A, we show that entanglement can indeed improve the cumulative complexity with respect to unentangled pebbling. However, in the next section, we will show that this is not true with respect to *time* complexity. This fact will be useful below for several reasons.

2.3 Entanglement Does not Improve Time Complexity

In terms of time complexity, entangled pebbling are no more efficient than normal pebbles.

Lemma 3 (Entangled Time = Simple Time). *For any* $G,C,m,\mathsf{T}^{\updownarrow},P_{\mathsf{init}}{}^{\updownarrow}$ *and* $\epsilon\geq 0$ *there exist a* $\mathsf{T},P_{\mathsf{init}}$ *such that* $|P_{\mathsf{init}}|\leq w(P_{\mathsf{init}}{}^{\updownarrow})$ *and*

$$\mathsf{time}(G,C,m,\mathsf{T},P_{\mathsf{init}}) \leq \mathsf{time}^{\updownarrow}(G,C,m,\mathsf{T}^{\updownarrow},P_{\mathsf{init}}{}^{\updownarrow}) \tag{2}$$

$$\mathsf{time}_{\epsilon}(G,C,m,\mathsf{T},P_{\mathsf{init}}) \leq \mathsf{time}^{\updownarrow}_{\epsilon}(G,C,m,\mathsf{T}^{\updownarrow},P_{\mathsf{init}}{}^{\updownarrow}) \tag{3}$$

in particular

$$\mathsf{time}^{\updownarrow}(G,C,m,k) = \mathsf{time}(G,C,m,k) \qquad \mathsf{time}^{\updownarrow}_{\epsilon}(G,C,m,k) = \mathsf{time}_{\epsilon}(G,C,m,k) \tag{4}$$

Proof. The $\geq$ directions in eq.(4) follows directly from the fact that a black pebbling is a special case of an entangled pebbling. The $\leq$ direction follows from eq.(2) and eq.(3). Below we prove eq.(2), the proof for eq.(3) is almost analogous.

We say that a player $\mathsf{A_{greedy}}$ for a normal or entangled pebbling is "greedy", if its strategy is simply to pebble everything possible in every round and never remove pebbles. Clearly, $\mathsf{A_{greedy}}$ is optimal for time complexity, i.e.,

$$\forall G, C, m, P_{\mathsf{init}} \; : \; \min_{\mathtt{T}} \mathsf{time}(G, C, m, \mathtt{T}, P_{\mathsf{init}}) = \mathsf{time}(G, C, m, \mathsf{A_{greedy}}, P_{\mathsf{init}}) \quad (5)$$

$$\forall G, C, m, P_{\mathsf{init}}^{\updownarrow} \; : \; \min_{\mathtt{T}} \mathsf{time}^{\updownarrow}(G, C, m, \mathtt{T}, P_{\mathsf{init}}^{\updownarrow}) = \mathsf{time}^{\updownarrow}(G, C, m, \mathsf{A_{greedy}}, P_{\mathsf{init}}^{\updownarrow}) \quad (6)$$

We next describe how to derive an initial black pebbling P_{init}^{*} from an entangled pebbling $P_{\mathsf{init}}^{\updownarrow}$ of cost $|P_{\mathsf{init}}^{*}| \leq w(P_{\mathsf{init}}^{\updownarrow})$ such that

$$\mathsf{time}(G, C, m, \mathsf{A_{greedy}}, P_{\mathsf{init}}^{*}) \leq \mathsf{time}^{\updownarrow}(G, C, m, \mathsf{A_{greedy}}, P_{\mathsf{init}}^{\updownarrow}) \quad (7)$$

Note that this then proves eq.(2) (with $\mathsf{A_{greedy}}, P_{\mathsf{init}}^{*}$ being $\mathtt{T}, P_{\mathsf{init}}$ in the statement of the lemma) as

$$\mathsf{time}^{\updownarrow}(G, C, m, \mathtt{T}^{\updownarrow}, P_{\mathsf{init}}^{\updownarrow}) \geq \mathsf{time}^{\updownarrow}(G, C, m, \mathsf{A_{greedy}}, P_{\mathsf{init}}^{\updownarrow}) \quad (8)$$

$$\geq \mathsf{time}(G, C, m, \mathsf{A_{greedy}}, P_{\mathsf{init}}^{*}) \quad (9)$$

It remains to prove eq.(7). For every share $\langle \mathcal{Y} \rangle_t \in P_{\mathsf{init}}^{\updownarrow}$ we observe which $|\mathcal{Y}| - t$ pebbles are the last ones to become available[6] in the random experiment $\mathsf{pebble}^{\updownarrow}(G, C, m, \mathtt{T}^{\updownarrow}, P_{\mathsf{init}}^{\updownarrow})$, and we add these pebbles to P_{init} if they're not already in there.

Note that then $|P_{\mathsf{init}}| \leq w(P_{\mathsf{init}}^{\updownarrow})$ as required. Moreover eq.(7) holds as at any timestep, the nodes available in $\mathsf{pebble}^{\updownarrow}(G, C, m, \mathsf{A_{greedy}}, P_{\mathsf{init}}^{\updownarrow})$ are nodes already pebbled in $\mathsf{pebble}(G, C, m, \mathsf{A_{greedy}}, P_{\mathsf{init}}^{*})$ at the same timestep.

2.4 The Parallel Random Oracle Model (pROM)

We now turn to an analogue of the above pebbling games but this time in the parallel random oracle model (pROM) of [AS15].

Let $G = (V, E)$ be a DAG with a dedicated set $C \subseteq V$ of challenge edges, we identify the vertices with $V = [n]$. A labelling $\ell_1, \ldots, \ell_n$ of G's vertices using a hash functiotn $\mathsf{h} : \{0,1\}^* \to \{0,1\}^w$ is defined as follows. Let $\mathsf{parent}(i) = \{j \in V : (j, i) \in E\}$ denote the parents of i, then

$$\ell_i = \mathsf{h}(i, \ell_{p_1}, \ldots, \ell_{p_d}) \quad \text{where} \quad (p_1, \ldots, p_d) = \mathsf{parent}(i) \quad (10)$$

Note that if i is a source, then its label is simply $\ell_i = \mathsf{h}(i)$.

---

[6] A pebble is available if it's in the closure of the current entangled pebbling configuration, also note that $\mathsf{A_{greedy}}$'s strategy is deterministic and independent of the challenges it gets, so the "last nodes to become available" is well defined.

---

computeLabel$(G, C, m, \mathtt{A}, \sigma_{\mathsf{init}}, \mathsf{h} : \{0,1\}^* \to \{0,1\}^w)$:

1. Initialise $\mathsf{cnt} := 0$, $\mathsf{round} := 0$, $\sigma_{\mathsf{cnt}} := \sigma_{\mathsf{init}}$ and $\mathsf{cost} := 0$.
2. A challenge $c \leftarrow C$ is chosen uniformly from C.
3. $(q_1, \ldots, q_s, \ell) \leftarrow \mathtt{A}(c, \sigma_{\mathsf{cnt}})$ $\mathtt{A}$ *choses parallel* h *queries and (optionally) a guess for* ℓ_c
4. $\mathsf{cost} := \mathsf{cost} + |\sigma_{\mathsf{cnt}}| + s \cdot w$.
5. $(\sigma_{\mathsf{cnt}+1}) \leftarrow \mathtt{A}(c, \sigma_{\mathsf{cnt}}, \mathsf{h}(q_1), \ldots, \mathsf{h}(q_s))$ $\mathtt{A}$ *outputs next state*
6. $\mathsf{cnt} := \mathsf{cnt} + 1$
7. If $\ell = \bot$ (no guess in this round) go to step 3.
8. If $\ell \neq \ell_c$ (wrong guess) set $\mathsf{cost} = \infty$ and abort.
9. $\mathsf{round} := \mathsf{round} + 1$. If $\mathsf{round} = m$ the experiment is over, otherwise go to step 2.
10. $\mathsf{round} := \mathsf{round}+1$. If $\mathsf{round} < m$ go to step 2, otherwise if $\mathsf{round} = m$ the experiment is over, the output is the final count cnt and the cumulative cost cost.

We consider a game computeLabel$(G, C, m, \mathtt{A}, \sigma_{\mathsf{init}}, \mathsf{h})$ where an algorithm $\mathtt{A}$ must m times consecutively compute the label of a node chosen at random from C. $\mathtt{A}$ gets an initial state $\sigma_0 = \sigma_{\mathsf{init}}$. The *cumulative memory complexity* is defined as follows.

$$\mathsf{cmc}^{\mathrm{pROM}}(\mathsf{G}, C, m, \mathtt{A}, \sigma_{\mathsf{init}}, \mathsf{h}) = \mathop{\mathbb{E}}_{\mathsf{computeLabel}(G,C,m,\mathtt{A},\sigma_{\mathsf{init}},\mathsf{h})}[\mathsf{cost}]$$

$$\mathsf{cmc}^{\mathrm{pROM}}(G, C, m, \sigma_{\mathsf{init}}) = \min_{\mathtt{A}} \mathop{\mathbb{E}}_{\mathsf{h} \leftarrow \mathcal{H}} \mathsf{cmc}^{\mathrm{pROM}}(G, C, m, \mathtt{A}, \sigma_{\mathsf{init}}, \mathsf{h})$$

The *time complexity* of a given adversary is

$$\mathsf{time}^{\mathrm{pROM}}(G, C, m, \mathtt{A}, \sigma_{\mathsf{init}}, \mathsf{h}) = \mathop{\mathbb{E}}_{\mathsf{computeLabel}(G,C,m,\mathtt{A},\sigma_{\mathsf{init}},\mathsf{h})}[\mathsf{cnt}]$$

We will also consider this notion against the best adversaries from some restricted class of adversaries, in this case we put the class as subscript, like

$$\mathsf{cmc}_{\mathcal{A}}^{\mathrm{pROM}}(G, C, m, \sigma_{\mathsf{init}}) = \min_{\mathtt{A} \in \mathcal{A}} \mathop{\mathbb{E}}_{\mathsf{h} \leftarrow \mathcal{H}} \mathsf{cmc}^{\mathrm{pROM}}(G, C, m, \mathtt{A}, \sigma_{\mathsf{init}}, \mathsf{h})$$

As for pebbling, also here we will consider the more meaningful ϵ variants of these notions

$$\mathsf{cmc}_{\epsilon}^{\mathrm{pROM}}(\mathsf{G}, C, m, \mathtt{A}, \sigma_{\mathsf{init}}, \mathsf{h}) = \inf \left\{ \gamma \ \middle| \ \mathop{\mathbb{P}}_{\mathsf{computeLabel}(G,C,m,\mathtt{A},\sigma_{\mathsf{init}},\mathsf{h})}[\mathsf{cost} \geq \gamma] \geq 1 - \epsilon \right\}$$

$$\mathsf{cmc}_{\epsilon}^{\mathrm{pROM}}(G, C, m, \sigma_{\mathsf{init}}) = \min_{\mathtt{A}} \mathop{\mathbb{E}}_{\mathsf{h} \leftarrow \mathcal{H}} \mathsf{cmc}_{\epsilon}^{\mathrm{pROM}}(G, C, m, \mathtt{A}, \sigma_{\mathsf{init}}, \mathsf{h})$$

$$\mathsf{time}_{\epsilon}^{\mathrm{pROM}}(\mathsf{G}, C, m, \mathtt{A}, \sigma_{\mathsf{init}}, \mathsf{h}) = \inf \left\{ \gamma \ \middle| \ \mathop{\mathbb{P}}_{\mathsf{computeLabel}(G,C,m,\mathtt{A},\sigma_{\mathsf{init}},\mathsf{h})}[\mathsf{cnt} \geq \gamma] \geq 1 - \epsilon \right\}$$

2.5 scrypt and the computeLabel Game

We briefly discuss the relation between evaluating the memory-hard function candidate scrypt in the pROM and the computeLabel game for the line graph (described bellow) and, and explain why we will focus on the latter.

Recall that Percival [Per09] gave the first proposal of a memory-hard hash function, called scrypt. scrypt[7] uses a hash function $h : \{0,1\}^* \to \{0,1\}^w$ (e.g., SHA-256), and proceeds in two phases, given an input X. In the first phase it computes $X_i = h^i(X)$ for all $i \in [n]$,[8] and in the second phase, setting $S_0 = X_n$, it computes $S_1, \ldots, S_n$ defined recursively to be

$$S_i = h(S_{i-1} \oplus X_{\mathsf{int}(S_{i-1})})$$

where $\mathsf{int}(S)$ reduces a w-bit string S to an integer in $[n]$ such that if S is uniform random then $\mathsf{int}(S)$ is (close to) uniform over $[n]$. The final output of $\mathsf{scrypt}_n^h(X) = S_n$.

To show that scrypt is memory-hard we would like to lower-bound the cumulative memory complexity (CMC) of computing scrypt in the pROM.[9]

We argue that to obtain a bound on the CMC of scrypt it suffices to restrict our attention to the minimal final value of cost in $\mathsf{cmc}^{\mathsf{pROM}}(L_n, [n], n)$ where $L_n = (V, E)$ is the *line graph* where $V = [n]$ and $E = \{(i, i+1) \; : \; i \in [n-1]\}$. Intuitively this is rather easy to see. Clearly any algorithm which hopes to evaluate scrypt with more than negligble probability must, at some point, compute all X_i values and all S_j values since guessing them is almost impossible. Moreover until S_{i-1} has been computed the value of $\mathsf{int}(S_{i-1})$ – i.e. the challenge label needed to compute S_i – is uniform random and independent, just like the distribution of i^{th} challenge $c \leftarrow C$ in the computeLabel game. In other words once an algorithm has computed the values $X_1, \ldots, X_n$ computing the values of $S_1, \ldots, S_n$ corresponds exactly to playing the computeLabel game on graph L_n with challenge set $[n]$ for n rounds. The initial state is exactly the state given to the algorithm as input in the step where it first computes X_n. It is immediate that, when restricted to strategies which don't simply guess relevant outputs of h, then any strategy for computing the values $S_1, \ldots, S_n$ corresponds to a strategy for playing $\mathsf{computeLabel}(L_n, [n], n)$.

In summary, once A has finished the first phase of evaluating scrypt, the second phase essentially corresponds to playing the computeLabel game on the graph L_n with challenge set $[n]$ for n rounds. The initial state σ_{init} in computeLabel is the state given to A as input in the first step of round 1 (i.e. in the step

---

[7] In fact, what we describe here is only a subroutine of the whole scrypt function, called ROMix. However ROMix is the actual core of the scrypt function, and we will use the generic name "scrypt" for it in the following.

[8] Here $h^i(X)$ denotes iteratively applying h i times to the input X.

[9] For a given execution in the pROM, the CMC is the sum of the size of each intermediary state between batches of calls to the ROM. In particular the cost variable in the computeLabel experiment is a special case of CMC when computing a particular task. For the details of the definition of CMC for a *arbitrary* computation we refer the interested reader to [AS15].

when A first computes X_n). It is now immediate that (when restricted to strategies which don't simply guess relevant outputs of h) then any strategy A for computing the second phase of scrypt is essentially a strategy for playing computeLabel($L_n, [n], n$). Clearly the total CMC of A when computing both phases of scrypt is at least the CMC of computing just the second. Thus our lowerbound on $\mathsf{cmc}^{\mathsf{pROM}}(L_n, [n], n)$ in Theorem 15 also gives us a lower bound on the CMC of scrypt_n.

Simple Algorithms. Theorem 15 is very general in that it makes no restrictions on the algorithm playing computeLabel. However this comes at the cost of relying on Conjecture 13. If we are willing to restrict our class of algorithms evaluating scrypt to *simple* algorithms $\mathcal{A}_{SA}$ then we obtain an unconditional lower-bound on the CMC of scrypt by using Theorem 4. Intuitively a simple algorithms $\mathtt{A} \in \mathcal{A}_{SA}$ is one which either stores a value X_i directly in its intermediary states[10] or stores nothing about the value of X_i at all. (They are however permitted to store arbitrary other information in their states.) For example a simple algorithm may not store, say, $X_i \oplus X_j$ or just the first 20 bits of X_i. We note that, to the best of our knowledge, all algorithms in the literature for computing scrypt are indeed of this form.

Much as in the more general case above, for the set of algorithms $\mathcal{A}_{SA}$ we can now draw a parallel between computing phase two of scrypt in the pROM and playing the game pebble on the graph L_n with challenge set $[n]$ for n rounds. In particular, having a pebble on a node $v \in [n]$ in some pebbling configuration P_{cnt} corresponds to A having stored the value of X_v in state σ_{cnt} which it recieved as input in that step cnt. Given this analogy, any $\mathtt{A} \in \mathcal{A}_{SA}$ which computes scrypt_n (without guessing some intermediary value X_i or S_j) implements a strategy T for playing pebble such that Equation 1 is satisfied at each step of each round of pebble and the CMC of A when computing phase two is at least $\mathsf{cc}(L_n, [n], n, \mathtt{T}, \sigma_{\mathsf{init}})$ where σ_{init} is the state given to A at the first step of phase two. Therefor Theorem 4 immediatly gives us a lower-bound on the CMC of scrypt_n for all algorithms in $\mathcal{A}_{SA}$.

Entangled Adversaries. In fact we can even relax our restrictions on algorithms computing scrypt to the class $\mathcal{A}_{EA}$ of *entangled* algorithms while still obtaining an unconditional lower-bound on the CMC of scrypt. In adition to what is permitted for simple algorithms we also allow storing "entangled" information about the values of $X_1, \ldots, X_n$ of the following form. For any subset $L \subseteq [n]$ and integer $t \in [|L|]$ an algorithm can store an encoding of $X_L = \{X_i\}_{i \in L}$ such that if it obtains any t values in L then it can immediately output all remaining $|L| - t$ values in L with no further information or queries to h. One such encoding uses polynomial interpolation as described in the introduction. Indeed, this motivates our definition of entangled pebbles above.

As demonstrated in Appendix A, the class $\mathcal{A}_{EA}$ is (in general) strictly more powerful $\mathcal{A}_{SA}$ when it comes to minimizing CMC. Thus we obtain a more general

---

[10] or at least an equivalent encoding of X_i

unconditional lower-bound on the CMC of `scrypt` using Theorem 9 which lower-bounds $cc^{\updownarrow}(L_n, [n], n)$ the entangled cumulative pebbling complexity of L_n.

3   Pebbling Lower Bounds for the Chain Graph

In this section, we provide lower bounds for the cumulative complexity of the n-round probabilistic pebbling game on the line graph L_n with challenges from $[n]$. We will start with the case without entanglement (i.e., dealing only with black pebbles) which captures the essence of our proof, and then below, explain how our proof approach can be extended to the entangled case.

Theorem 4 (Pebbling Complexity of the line graph). *For all $0 \le k \le n$, and constant $\delta > 0$,*

$$cc_{\delta}(L_n, C = [n], n, k) = \Omega_{\delta}\left(\frac{n^2}{\log^2(n)}\right) .$$

We note in passing that the above theorem can be extended to handle a different number of challenges $t \ne n$, as it will be clear in the proof. We dispense with the more general theorem, and stick with the simpler statement for the common case $t = n$ motivated by `scrypt`. The notation Ω_{δ} indicates that the constant hidden in the Ω depends on δ.

In fact, we also note that our proof allows for more concrete statements as a function of δ, which may be constant. However, not surprisingly, the bound becomes weaker the smaller δ is, but note that if we are only interested in the expectation $cc(L_n, C = [n], n, k)$, then applying the result with $\delta = O(1)$ (e.g., $\frac{1}{2}$ is sufficient to obtain a lower bound of $\Omega\left(\frac{n^2}{\log^2 n}\right)$.

Proof intuition – the expectation game. Before we turn to the formal proof, we give some high level intuition of how the bound is proved. It turns out that most of the proof is going to in fact lower bound the cc of a much simpler game, where the goal is far simpler than covering challenges from $[n]$ with a pebble. In fact, the game will be completely deterministic.

The key observation is that every time a new challenge c_i is drawn, and the player has reached a certain pebbling configuration P, then there is a well-defined expected number $\Phi(P)$ of steps the adversary needs to take *at least* in order to cover the random challenge. We refer to $\Phi(P)$ as the *potential* of P. In particular, the best strategy is the greedy one, which looks at the largest $j = j(c_i) \le c_i$ on which a pebble is placed, i.e., $j \in P$, and then needs to output a valid sequence of *at least* $c_i - j$ further pebbling configurations, such that the last configuration contains c_i. Note if $j = c_i$, we still need to perform one step to output a valid configuration. Therefore, $\Phi(P)$ is the expected value of $\max(1, c_i - j(c_i))$. We will consider a new game – called the *expectation game* – which has the property that at the beginning of every stage, the challenger just computes $\Phi(P)$, and expects the player `T` to take $\Phi(P)$ legal steps until `T` can move to the next stage.

Note that these steps can be totally arbitrary – there is no actual challenge any more to cover. Still, we will be interested in lower bounding the *cumulative complexity* of such a strategy for the expectation game, and it is not obvious how T can keep the cc low. Indeed:

- If the potential is high, say $\Phi(P) = \Omega(n)$, then this means that linearly many steps must be taken to move to the next stage, and since every configuration contains at least one pebble, we pay a cumulative cost of $\Omega(n)$ for the present stage.
- Conversely, if the potential $\Phi(P)$ is *low* (e.g., $O(1)$), then we can expect to be faster. However we will show that this implies that there are *many* pebbles in P (at least $\Omega(n/\Phi(P))$), and thus one can expect high cumulative cost again, i.e., linear $\Omega(n)$.

However, there is a catch – the above statements refer to the *initial configurations*. The fact that we have many pebbles at the beginning of a stage and at its end, does not mean we have many pebbles *throughout the whole stage*. Even though the strategy T is forced to pay $\Phi(P)$ steps, the strategy may try to drop as many pebbles as possible for a while, and then adding them back again. *Excluding that this can happen is the crux of our proof.* We will indeed show that for the expectation game, any strategy incurs cumulative complexity $\Omega(n^2/\log^2(n))$ roughly. The core of the analysis will be understanding the behavior of the potential function throughout a stage.

Now, we can expect that a low-cc strategy T for the original parallel pebbling game on L_n gives us one for the expectation game too – after all, for every challenge, the strategy T needs to perform roughly $\Phi(P)$ steps from the initial pebbling configuration when learning the challenge. This is *almost* correct, but again, there is a small catch. The issue is that $\Phi(P)$ is only an expectation, yet we want to have the guarantee that we go for $\Phi(P)$ steps with sufficiently high probability (this is particularly crucial if we want to prove a statement which parameterized by δ). However, this is fairly simple (if somewhat tedious) to overcome – the idea is that we partition the n challenges into n/λ groups of λ challenges. For every such group, we look at the initial configuration P when learning the first of the next λ challenges, and note that with sufficiently high probability (roughly $e^{-\Omega(\lambda^2)}$ by a Chernoff bound) there will be one challenge (among these λ ones) which is at least (say) $\Phi(P)/2$ away from the closest pebble. This allows us to reduce a strategy for the n-challenge pebbling game on L_n to a strategy for the (n/λ)-round expectation game. The value of λ can be chosen small enough not to affect the overall analysis.

Proof (Theorem 4). As the first step in the proof, we are going to reduce playing the game $\mathsf{pebble}(L_n, C = [n], n, T, P_{\mathsf{init}})$, for an arbitrary player T and initial pebbling configuration P_{init} ($|P_{\mathsf{init}}| \leq k$), to a simpler (and somewhat different) pebbling game, which refer to as the *expectation game*.

To this end, we introduce first the concept of a *potential* of a pebbling configuration. For a pebbling configuration $P = \{\ell_1, \ell_2, \ldots, \ell_m\} \subseteq [n]$, we define the

potential function $\Phi : 2^{[n]} \to \mathbb{N}$ as

$$\Phi(P) := \tfrac{m}{n} + \tfrac{1}{n} \sum_{i=0}^{m} (1 + \ldots + (l_{i+1} - l_i - 1))$$

$$= \tfrac{m}{n} + \tfrac{1}{2n} \sum_{i=0}^{m} (\ell_{i+1} - \ell_i) \cdot (\ell_{i+1} - \ell_i - 1) = \tfrac{1}{2n} \sum_{i=0}^{m} (\ell_{i+1} - \ell_i)^2 - \tfrac{n+1-2m}{2n}$$

where $m = |P|$, where we let $\ell_0 = 0$ and $\ell_{m+1} = n + 1$. Indeed, $\Phi(P)$ is the expected number of moves required (by an optimal strategy) to pebble a random challenge starting from the pebbling configuration P, where the expectation is over the choice of the random challenge. (Note in particular that $\Phi(P)$ requires to pay at least one move even if we have already a pebble on the challenge.) In other words, $\Phi(P)$ is exactly $\mathsf{time}(L_n, [n], 1, \mathtt{T}^*, P)$ for the optimal strategy $\mathtt{T}^*$ (which is $\mathsf{A}_{\mathsf{greedy}}$ from the proof of Lemma 3).

The potential is used in the definition of the game $\mathsf{expect}(n, t, \mathtt{T}, P_{\mathsf{init}})$ which is played by a pebbler $\mathtt{T}$ as follows.

$\mathsf{expect}(n, t, \mathtt{T}, P_{\mathsf{init}})$: The t-round expectation game of parameter n and an initial pebbling configuration $P_{\mathsf{init}} \subseteq V$ is played by challenger and player $\mathtt{T}$ as follows.

1. Initialize $\mathsf{cnt} := 0$, $\mathsf{round} := 0$, $P_{\mathsf{cnt}} := P_{\mathsf{init}}$ and $\mathsf{cost} := |P_{\mathsf{init}}|$.
2. Player $\mathtt{T}$ submits a sequence of **non-empty** pebbling configurations $(P_{\mathsf{round},1}, \ldots, P_{\mathsf{round},t_{\mathsf{round}}}) \subset [n]^{\times t_{\mathsf{round}}}$,
3. Let $P_{\mathsf{round},0} := P_{\mathsf{cnt}}$. Check if $t_{\mathsf{round}} \geq \Phi(P_{\mathsf{cnt}})$ and $\forall i \in [t_{\mathsf{round}}]$

$$\forall v \in P_{\mathsf{round},i} \setminus P_{\mathsf{round},i-1} \ : \ \mathsf{parent}(v) \in P_{\mathsf{round},i-1} \ .$$

   If check fails, output $\mathsf{cnt} = \mathsf{cost} = \infty$ and halt.
4. $\mathsf{cnt} := \mathsf{cnt} + t_{\mathsf{round}}$.
5. $\mathsf{cost} := \mathsf{cost} + \sum_{j=1}^{t_{\mathsf{round}}} |P_{\mathsf{round},j}|$.
6. $P_{\mathsf{cnt}} := P_{\mathsf{round},t_{\mathsf{round}}}$.
7. $\mathsf{round} := \mathsf{round} + 1$. If $\mathsf{round} < t$ go to step 2, otherwise if $\mathsf{round} = t$ the experiment is over, the output is the final count cnt and the cumulative cost cost.

For a (randomized) pebbler $\mathtt{T}$ and initial configuration P_{init}, we write $\mathsf{expect}_{n,t}(\mathtt{T}, P_{\mathsf{init}})$ for the output of the expectation game, noting that here this random variable only depends on the randomness of the pebbler and configuration P_{init}. We also define

$$\mathsf{cc}_\delta(\mathsf{expect}_{n,t}(\mathtt{T}, P_{\mathsf{init}})) := \inf \left\{ \gamma \ \middle| \ \Pr_{\mathsf{expect}(n,t,\mathtt{T},P_{\mathsf{init}})} [\mathsf{cost} \geq \gamma] \geq 1 - \epsilon \right\}$$

$$\mathsf{cc}_\delta(\mathsf{expect}_{n,t,k}) := \min_{\substack{\mathtt{T}, P_{\mathsf{init}} \subseteq V \\ |P_{\mathsf{init}}| \leq k}} \left\{ \mathsf{cc}_\delta(\mathsf{expect}_{n,t}(\mathtt{T}, P_{\mathsf{init}})) \right\}$$

The expectation game $\mathsf{expect}_{n,t,k}$ has however an important feature: Because the randomness is only over the pebbler's coins, these coins can be fixed

to their optimal choice without making the overall cc worse. This implies that $cc_\delta(\text{expect}_{n,t,k}) = cc_0(\text{expect}_{n,t,k})$ for all $\delta \geq 0$. In particular, we use the short-hand $cc(\text{expect}_{n,t,k})$ for the latter.

In the remainder of the proof, we are going to prove the following two lemmas. Below, we combine these two lemmas in the final statement, before turning to their proofs. (The proof of Lemma 5 is deferred to Appendix C.

Lemma 5 (Reduction to the expectation game). *For all n, t, k, λ, and any $\delta > 3\mu(t,\lambda)$, we have*

$$cc(\text{expect}_{n,t,k}) = cc_{\delta-3\mu(t,\lambda)}(\text{expect}_{n,t,k}) \leq 2 \cdot cc_\delta(L_n, C = [n], t \cdot \lambda, k) \ ,$$

where $\mu(t,\lambda) = t \cdot e^{-\lambda^2/8}$.

We note that in general, for every $\delta' \leq \delta$, we have $cc_{\delta'}(\text{expect}_{n,t,k}) \leq cc_\delta(\text{expect}_{n,t,k})$. This is because if a c is such that for all T and P_{init} we have $\mathbb{P}(\text{expect}_{n,t}(\text{T}, P_{\text{init}}) \geq c) \geq 1 - \delta'$, then also $\mathbb{P}(\text{expect}_{n,t}(\text{T}, P_{\text{init}}) \geq c) \geq 1 - \delta$. Thus the set from which we are taking the supremum only grows bigger as δ increases. In the specific case of Lemma 5, the $3\mu(t,\lambda)$ offset captures the loss of our reduction.

Lemma 6 (cc complexity of the expectation game). *For all $t, 0 \leq k \leq n$ and $\epsilon > 0$, we have*

$$cc(\text{expect}_{n,t,k}) \geq \left\lfloor \frac{\epsilon t}{2} \right\rfloor \cdot \frac{n^{1-\epsilon}}{6} \ .$$

To conclude the proof, before we turn to the proofs of the above two lemmas, note that we need to choose t, λ such that $t \cdot \lambda = n$, and $\mu(t,\lambda) = t \cdot e^{-\lambda^2/8} < \delta/3$. We also set $\epsilon = 0.5 \log\log(n)/\log(n)$, and note that in this case $n^{1-\epsilon} = n/\sqrt{\log(n)}$. In particular, we can set $\lambda = O(\sqrt{\log t})$, and can choose e.g. $t = n/\sqrt{\log n}$. Then, by Lemma 6,

$$cc(\text{expect}_{n,t,k}) \geq \left\lfloor \frac{\epsilon t}{2} \right\rfloor \cdot \frac{n^{1-\epsilon}}{6} = \Omega\left(\frac{n^2}{\log^2(n)}\right) \ .$$

This concludes the proof of Theorem 4.

Proof (Proof of Lemma 6). It is not difficult to see that if a pebbling configuration P has potential ϕ, then the size $|P|$ of the pebbling configuration (i.e., the number of vertices on which a pebble is placed) is at least $\frac{n}{6 \cdot \phi}$. We give formal proof for completeness.[11].

Lemma 7. *For every* **non-empty** *pebbling configuration $P \subseteq [n]$, we have*

$$\Phi(P) \cdot |P| \geq \frac{n}{6} \ .$$

---

[11] Note that the contra-positive is not necessarily true. A simple counter-example is when pebbles are placed on vertices $[0, n/2]$ of C_n (that is, $|P| = O(n)$). The expected number of moves in this case is still $\Omega(n)$.

Proof. Let $m = |P| \geq 1$, by definition:

$$\Phi(P) \cdot m = \left[\frac{1}{2n} \sum_{i=0}^{m} (\ell_{i+1} - \ell_i)^2 - \frac{n+1-2m}{2n} \right] \cdot m \ ,$$

where $\ell_0 = 0$ and $\ell_{m+1} = n+1$ are notational placeholders. Since configuration P is non-empty and $\Phi(P) \geq 1$, we have $m \geq \frac{m+1}{2}$ and: $\frac{n+1-2m}{2n} \leq \frac{1}{2} \leq \frac{1}{2} \cdot \Phi(P)$. Therefore

$$\frac{1}{2n} \sum_{i=0}^{m} (\ell_{i+1} - \ell_i)^2 \leq \frac{3}{2} \cdot \Phi(P) \ ,$$

and then

$$\Phi(P) \cdot m \geq \frac{2}{3} \left(\frac{1}{2n} \sum_{i=0}^{m} (\ell_{i+1} - \ell_i)^2 \right) \cdot \frac{m+1}{2} \ .$$

The fact that $\Phi(P) \cdot m \geq \frac{n}{6}$ follows because by the Cauchy-Schwarz Inequality,

$$\sum_{i=0}^{m} (\ell_{i+1} - \ell_i)^2 \cdot (m+1) \geq \left(\sum_{i=0}^{m} (\ell_{i+1} - \ell_i) \right)^2 \geq n^2 \ . \ \square$$

Also, the following claim provides an important property of the potential function.

Lemma 8. *In one iteration, the potential can decrease by at most one.*

Proof. Consider any arbitrary configuration $P = \{\ell_1, \ell_2, \ldots, \ell_m\} \subseteq [n]$. The best that a pebbling algorithm can do is to place new pebbles next to *all* the current pebbles – let's call the new configuration P'. That is,

$$P' = \{\ell_1, \ell_1 + 1, \ell_2, \ell_2 + 1, \ldots, \ell_m, \ell_m + 1\} \subseteq [n].$$

It's not difficult to see that that the potential of the new configuration is

$$\Phi(P') = \frac{1}{2n} \left(\ell_1^2 + \sum_{i=1}^{m} 1 + (\ell_{i+1} - (\ell_i + 1))^2 \right) - \frac{n+1-2|P'|}{2n} \tag{11}$$

$$\geq \frac{1}{2n} \left(m + \sum_{i=0}^{m} ((\ell_{i+1} - \ell_i)^2 - 2(\ell_{i+1} - \ell_i) + 1) \right) - \frac{n+1-2m}{2n} \tag{12}$$

$$\geq \Phi(P) + \frac{m}{n} - \frac{1}{n} \sum_{i=0}^{m} (\ell_{i+1} - \ell_i) \geq \Phi(P) - 1 \tag{13}$$

The first inequality holds because $|P'| \geq m$. $\qquad\square$

Assume without loss of generality the pebbler T is legal and deterministic. Consider a particular round $i \in [t]$ of the expectation game. Let P and P' denote the initial and final pebbling configurations in the i-th round, and let us denote $\phi_i = \Phi(P)$, the potential of the initial configuration in round i. Depending on the value of $\Phi(P')$, we classify the pebbling sequence from P to P' into three different categories:

<u>Type 1:</u> $\Phi(P') > \phi_i \cdot n^{\epsilon/2}$; or

<u>Type 2:</u> $\Phi(P') \leq \phi_i \cdot n^{\epsilon/2}$ – we have two sub-cases:

    <u>Type 2a:</u> the potential was *always* less than $\phi_i \cdot n^\epsilon$ for all the intermediate pebbling configurations; or

    <u>Type 2b:</u> the potential went above $\phi_i \cdot n^\epsilon$ for some intermediate configuration.

With each type, we associate a cost that the pebbling algorithm has to pay, which lower bounds the contribution to the cumulative complexity of the pebbling configurations generated during this stage. The pebbling algorithm can carry out pebbling of Type 1 for *free*[12] – however, the latter two have accompanying costs.

- For pebbling sequences of Type 2a, it is not difficult to see that the cumulative cost is at least $\phi_i \frac{n}{6n^\epsilon \cdot \phi_i} = \frac{1}{6} n^{1-\epsilon}$ since by **Lemma 7**, the size of the pebbling configuration is never less than $n/(6\phi_i n^\epsilon)$ during all intermediate iterations and in stage i valid pebbler must produce at least ϕ_i configurations.

- For sequences of Type 2b, by **Lemma 8**, it follows that in a Type 2b sequence it takes at least $\phi_i(n^\epsilon - n^{\epsilon/2})$ steps to decrease the potential from ϕn^ϵ to $\phi_i \cdot n^{\epsilon/2}$, and the size of the pebbling configuration is at least $\frac{n}{6\phi_i n^\epsilon}$ in every intermediate step by **Lemma 7**. Therefore, the cumulative cost is at least

$$\phi_i(n^\epsilon - n^{\epsilon/2}) \times \frac{n}{6\phi_i n^\epsilon} \geq \frac{n}{6} - \frac{n^{1-\epsilon/2}}{6} \geq \frac{1}{6} n^{1-\epsilon} \ ,$$

    where the last inequality follows for sufficiently large n.

To conclude the proof, we partition the $t \geq \lceil 2/\epsilon \rceil$ phases into groups of consecutive $\lceil 2/\epsilon \rceil$ phases. It is not difficult to see that any group *must* contain at least one sequence of Type 2: otherwise, with ϕ being the potential at the beginning of the first of theses $2/\epsilon$ phases, the potential at the end would be strictly larger than

$$\phi n^{\frac{\epsilon}{2} \frac{2}{\epsilon}} \geq \phi \cdot n > n/2$$

which cannot be, as the potential can be at most $\frac{n}{2}$. By the above, however, the cumulative complexity of each group of phases is at least $\frac{n^{1-\epsilon}}{6}$, and thus we get

$$cc(\mathsf{expect}_{n,t,k}) \geq \left\lfloor \frac{\epsilon t}{2} \right\rfloor \cdot \frac{n^{1-\epsilon}}{6} \ , \tag{14}$$

which concludes the proof of Lemma 6. $\qquad\qquad\qquad\qquad\qquad\qquad\qquad$ $\square$

As the second result, we show that the above theorem also holds for the entangled case.

---

[12] The cost is might be greater than zero, but setting it to zero doesn't affect the lower bound.

Theorem 9 (Entangled Pebbling Complexity of the Chain Graph). *For all $0 \leq k \leq n$ and constant $\delta > 0$,*

$$\mathsf{cc}_\delta^{\updownarrow}(L_n, C = [n], n, k) = \Omega\left(\frac{n^2}{\log^2 n}\right) .$$

Luckily, it will not be necessary to repeat the whole proof. We will give now a proof sketch showing that in essence, the proof follows by repeating the same format and arguments as the one for Theorem 4, using Lemma 3 as a tool.

Proof (Sketch). One can prove the theorem following exactly the same framework of Theorem 4, with a few differences. First off, we define a natural entangled version of the expectation game where, in addition to allowing entanglement in a pebbling configuration, we define the potential as

$$\Phi^{\updownarrow}(P) = \mathsf{time}^{\updownarrow}(L_n, C = [n], 1, \mathtt{T}^{*,\updownarrow}, P) ,$$

i.e., the expected time complexity for *one* challenge of an optimal entangled strategy $\mathtt{T}^{*,\updownarrow}$ starting from the (entangled) pebbling configuration P.

First off, a proof similar to the one of Lemma 5, based on a Chernoff bound, can be used to show that if we separate challenges in t chunks of λ challenges each, and we look at the configuration P at the beginning of each of the t chunks, then there exists at least one challenge (out of λ) which requires spending time $\Phi^{\updownarrow}(P)$ to be covered, except with small probability.

A lower bound on the cumulative complexity of the (entangled) expectaton game follow exactly the same proof as Lemma 6. This is because the following two facts (which correspond to the two lemmas in the proof of Lemma 6) are true also in the entanglement setting

- First off, for every P and $\mathtt{T}^{*,\updownarrow}$ such that $\Phi^{\updownarrow}(P) = \mathsf{time}^{\updownarrow}(L_n, C = [n], 1, \mathtt{T}^{*,\updownarrow}, P)$, Lemma 3 guarantees that there exist a (regular) pebbling strategy $\mathtt{T}'$ and a (regular) pebbling configuration P' such that $w(P) \geq |P'|$ and

$$\Phi^{\updownarrow}(P) = \mathsf{time}^{\updownarrow}(L_n, C = [n], 1, \mathtt{T}^{*,\updownarrow}, P)$$
$$\geq \mathsf{time}(L_n, C = [n], 1, \mathtt{T}', P') \geq \Phi(P') .$$

  Therefore, by Lemma 7,

$$w(P) \cdot \Phi^{\updownarrow}(P) \geq |P'| \cdot \Phi(P') \geq \frac{n}{6} .$$

- Second, the potential can decrease by at most one when making an arbitrary step from one configuration P to one configuration P'. This is by definition – assume it were not the case, and $\Phi^{\updownarrow}(P') < \Phi^{\updownarrow}(P) - 1$. Then, there exists a strategy to cover a random challenge starting from P wich first moves to P' in one step, and then applies the optimal strategy achieving expected time $\Phi^{\updownarrow}(P')$. The expected number of steps taken by this strategy is smaller than $\Phi^{\updownarrow}(P)$, contradicting the fact that $\Phi^{\updownarrow}(P)$ is the optimal number of steps required by any strategy. $\square$

4 From Pebbling to pROM

4.1 Trancscipts and Traces

We will now define the notion of a trace and transcript, which will allow us relate the computeLabel and pebble‡ experiments. For any possible sequence of challenges $c \in C^m$, let cnt_c denote the number of steps (i.e., the variable cnt) made in the computeLabel$(G, C, m, \mathtt{A}, \sigma_{\mathsf{init}}, \mathsf{h})$ experiment conditioned on the m challenges being c (note that once c is fixed, the entire experiment is deterministic, so cnt_c is well defined). Let $\tau_c = q_1|q_2|\ldots|q_{\mathsf{cnt}_c}$ be the *trace* of the computation: here $q_1 \subset [n]$ means that the first batch of parallel queries were the queries required to output the labels $\{\ell_i, i \in q_1\}$, etc..

For example $\tau_7 = 2|4, 5|7$ in the example below corresponds to a first query $\ell_2 = \mathsf{h}(2)$, then two parallel queries $\ell_4 = \mathsf{h}(4, \ell_1), \ell_5 = \mathsf{h}(5, \ell_2)$, and then the final query computing the label of the challenge $\ell_7 = \mathsf{h}(7, \ell_4, \ell_5, \ell_6)$.

A trace as a pebbling. We can think of a trace as a parallel pebbling, e.g., $\tau_7 = 2|4, 5|7$ means we pebble node 2 in the first step, nodes $4, 5$ in the second, and 7 in the last step. We say that an initial (entangled) pebbling configuration P_{init} is consistent with a trace τ if starting from P_{init}, τ is a valid pebbling sequence. E.g., consider again the traces $\tau_7 = 2|4, 5|7, \tau_8 = 3|6|8$ for the graph in Figure 1, then $P_{\mathsf{init}} = \{1, 5, 6\}$ is consistent with τ_7 and τ_8, and it's the smallest initial pebbling having this property. In the entangled case, $P_{\mathsf{init}}^{\ddagger} = \{\langle 1 \rangle_0, \langle 5, 6 \rangle_1\}$ is consistent with τ_7, τ_8. Note that in the entangled case we only need a pebbling configuration of weight 2, whereas the smallest pebbling configuration for the standard pebbling game has weight 3. In fact, there are traces where the gap between the smallest normal and entangled pebbling configuration consistent with all the traces can differ by a factor $\Theta(n)$.

Turning a trace into a transcript. We define the implications T_c of a trace $\tau_c = q_1|q_2|\ldots|q_{\mathsf{cnt}_c}$ as follows. For $i = 1, \ldots, \mathsf{cnt}_c$, we add the implication $(v_i) \to (f_i)$, where $v_i \subset [n]$ denotes all the vertices whose labels have appeared either as inputs or outputs in the experiment so far, and f_i denotes the labels contained in the inputs from this round which have never appeared before (if the guess for the challenge label in this round is non-empty, i.e., $\ell \neq \bot$, then we include ℓ in f_i).


Fig. 1

Example 10. Consider the graph from Figure 1 with $m = 1$ and challenge set $C = \{7, 8\}$ and the traces

$$\tau_7 = 2|4, 5|7 \quad \text{and} \quad \tau_8 = 3|6|8$$

We get

$$T_7 = \{(2) \to 1, (1, 2, 4, 5) \to 6\} \quad T_8 = \{(3, 6) \to 5\} \tag{15}$$

Where e.g. $(2) \to 1$ is in there as the first query is $\ell_2 = \mathsf{h}(2)$, and the second query is $\ell_4 = \mathsf{h}(4, \ell_1)$ and in parallel $\ell_5 = \mathsf{h}(5, \ell_2)$. At this point we so far only observed the label $v_2 = \{\ell_2\}$, so the label $f_2 = \{\ell_1\}$ used as input in this query is fresh, which means we add the implication $(2) \to 1$.

Above we formalised how to extract a transcript $T_{\boldsymbol{c}}$ from $(G, C, m, \mathsf{A}, \sigma_{\mathsf{init}}, \mathsf{h})$, with

$$T(G, C, m, \mathsf{A}, \sigma_{\mathsf{init}}, \mathsf{h}) = \cup_{\boldsymbol{c} \in C^m} T_{\boldsymbol{c}}$$

we denote the union of all $T_{\boldsymbol{c}}$'s.

4.2 Extractability and Coverability

In this section we introduce the notion of extractability and coverability of a transcript. Below we first give some intuition what these notions have to do with the computeLabel and pebble$^{\updownarrow}$ experiments.

Extractability intuition. Consider the experiment computeLabel$(G, C, m, \mathsf{A}, \sigma_{\mathsf{init}}, \mathsf{h})$. We can invoke A on some particular challenge sequence $\boldsymbol{c} \in C^m$, and if at some point A makes a query whose input contains a label ℓ_i which has not appeared before, we can "extract" this value from $(\mathsf{A}, \sigma_{\mathsf{init}})$ without explicit queuing h for it. More generally, we can run A on several challenge sequences scheduling queries in a way that will maximise the number of labels that can be extracted from $(\mathsf{A}, \sigma_{\mathsf{init}})$. To compute this number, we don't need to know the entire input/output behaviour of A for all possible challenge sequences, but the transcript $T = T(G, C, m, \mathsf{A}, \sigma_{\mathsf{init}}, \mathsf{h})$ is sufficient. Recall that T contains implication like $(1, 5, 6) \to 3$, which means that for some challenge sequence, there's some point in the experiment where A has already seen the labels ℓ_1, ℓ_5, ℓ_6, and at this point makes a query whose input contains a label ℓ_3 (that has not been observed before). Thus, given σ_{init} and ℓ_1, ℓ_5, ℓ_6 we can learn ℓ_3.

We denote with $ex(T)$ the number of labels that can be extracted using an optimal strategy. If the labels are uniformly random values in $\{0, 1\}^w$, then it follows that σ_{init} will almost certainly not be much smaller than $ex(T) \cdot w$, as otherwise we could compress $w \cdot ex(T)$ uniformly random bits (i.e., the extracted labels) to a string which is shorter than their length, but uniformly random values are not compressible.

Coverability intuition. In the following, we say that an entangled pebbling experiment $\mathsf{pebble}^{\ddagger}(G, C, m, \mathsf{P}, P_{\mathsf{init}}{}^{\ddagger})$ mimics the $\mathsf{computeLabel}(G, C, m, \mathtt{A}, \sigma_{\mathsf{init}}, \mathsf{h})$ experiment if for every challenge sequence the following is true: whenever $\mathtt{A}$ makes a query to compute some label $\ell_i = h(i, \ell_{p_1}, \dots, \ell_{p_t})$, P puts a (normal) pebble on i. For this $P_{\mathsf{init}}{}^{\ddagger}$ must contain (entangled) pebbles that allow to cover every implication in T (as defined above), e.g., if $(1, 5, 6) \to 3 \in T$, then from the initial pebbling $P_{\mathsf{init}}{}^{\ddagger}$ together with the pebbles $\langle 1 \rangle_0, \langle 5 \rangle_0, \langle 6 \rangle_0$ seen so far it must be possible derive $\langle 3 \rangle_0$, i.e., $\langle 3 \rangle_0 \in \mathsf{closure}(P_{\mathsf{init}} \cup \langle 1 \rangle_0, \langle 5 \rangle_0, \langle 6 \rangle_0)$. We say that such an initial state $P_{\mathsf{init}}{}^{\ddagger}$ covers T. We're interested in the maximum possible ratio of $\max_T \min_{P_{\mathsf{init}}{}^{\ddagger}, P_{\mathsf{init}}{}^{\ddagger} \text{ covers } T} \frac{w(P_{\mathsf{init}}{}^{\ddagger})}{ex(T)}$, and conjecture that it's a constant. That is, whenever T is k extractable, it can be covered by an initial pebbling $P_{\mathsf{init}}{}^{\ddagger}$ of weight $O(k)$. Assuming this conjecture we will be able to prove that pebbling time complexity implies the best possible pROM time complexity for any graph, and that cc complexity implies cumulative complexity in the pROM model for the scrypt graph. This section is self-contained (except for Definition 1) so people only interested in (dis)proving the conjecture just need to read the remaining part of this section.

Definitions. Let $n \in \mathbb{N}$, and $[n] = \{1, 2, \dots, n\}$. An "implication" $(\mathcal{X}) \to z$ given by a value $z \in [n]$ and a subset $\mathcal{X} \subset [n] \setminus z$ means that "knowing $\mathcal{X}$ gives z for free". We use $(\mathcal{X}) \to \mathcal{Z}$ as a shortcut for the set of implications $\{(\mathcal{X}) \to z \ : \ z \in \mathcal{Z}\}$.

A transcript is a set of of implications. Consider a transcript $T = \{\alpha_1, \dots, \alpha_\ell\}$, each α_i being an implication. We say that a transcript T is k $(0 \le k \le n)$ extractable if there exists an extractor E that makes at most $n - k$ queries in the following game:

- At any time E can query for a value in $[n]$.
- Assume E has values $\mathcal{L} \subset [n]$ and there exists an implication $(\mathcal{X}) \to z \in T$ where $\mathcal{X} \subset \mathcal{L}$, then E gets the value z "for free".
- The game is over when E has received all of $[n]$.

Every (even an empty) transcript T is 0 extractable as E can always simply ignore T and query for $1, 2, \dots, n$. Let

$$ex(T) = \max_k(T \text{ is } k\text{-extractable})$$

Example 11. Let $n = 5$ and consider the transcript

$$T = \{(1, 2) \to 3, (2, 3) \to 1, (3, 4) \to 2, (1) \to 4\} \tag{16}$$

This transcript is 2 but not 3 extractable. To see 2 extractability consider the E which first asks for 1, then gets 4 for free (due to $(1) \to 4$), next E asks for 2 and gets 3 for free (due to $(1, 2) \to 3$).

A set S of entangled pebbles covers an implication $(\mathcal{X}) \to z$ if $z \in \mathsf{closure}(S \cup \langle \mathcal{X} \rangle_0)$, with $\mathsf{closure}$ as defined in Definition 1.

Definition 12 (k-coverable). *We say that a transcript T is k-coverable if there exists a set of entangled pebbles S of total weight k such that every implication in T is covered by S. With $w(T)$ we denote the minimum weight of an S covering T:*

$$w(T) = \min_{S \text{ that covers } T} w(S)$$

Note that every transcript is trivially n coverable by using the pebble $\langle 1, \ldots, n \rangle_0$ of weight n which covers every possible implication. For the 2 extractable transcript from Example 11, a set of pebbles of total weight 2 covering it is

$$S = \{\langle 1, 2, 3 \rangle_2, \langle 1, 4 \rangle_1\} \tag{17}$$

For example $(3, 4) \to 2$ is covered as $2 \in \mathsf{closure}(\langle 1, 2, 3 \rangle_2, \langle 1, 4 \rangle_1, \langle 3, 4 \rangle_0) = \{1, 2, 3, 4\}$: we first can set $\Gamma = \{3, 4\}$ (using $\langle 3, 4 \rangle_0$), then $\Gamma = \{1, 3, 4\}$ using $\langle 1, 4 \rangle_1$, and then $\Gamma = \{1, 2, 3, 4\}$ using $\langle 1, 2, 3 \rangle_2$.

We will be interested in the size of the smallest cover for a transcript T. One could conjecture that every k-extractable transcript is k-coverable. Unfortunately this is not true, consider the transcript

$$T^* = \{(2, 5) \to 1, (1, 3) \to 2, (2, 4) \to 3, (3, 5) \to 4, (1, 4) \to 5\} \tag{18}$$

We have $ex(T^*) = 2$ (e.g. via query $2, 4, 5$ and extract $1, 3$ using $(2, 5) \to 1, (2, 4) \to 3$), but it's not 2-coverable (a cover of weight 3 is e.g. $\{\langle 5, 1 \rangle_1\}, \langle 2, 3, 4 \rangle_1\}$). With γ_n we denote the highest coverability vs extractability ration that a transcript over $[n]$ can have:

Conjecture 13. Let

$$\gamma_n = \max_{T \text{ over } [n]} \min_{S \text{ that covers } T} \frac{w(S)}{ex(T)} = \max_{T \text{ over } [n]} \frac{w(T)}{ex(T)}$$

then for some fixed constant γ, $\gamma_n \le \gamma$ for all $n \in \mathbb{N}$.

By the example eq.(18) eq.(18) above γ is at least $\gamma \ge \gamma_5 \ge 3/2$.

In Appendix B we'll introduce another parameter $shannon(w)$, which can give better lower bounds on the size of a state required to realize a given transcript in terms of Shannon entropy.

4.3 Pebbling time complexity implies pROM time complexity under Conjecture 13

We are ultimately interested in proving lower bounds on time or cumulative complexity in the parallel ROM model. We first show that pebbling time complexity implies time complexity in the ROM model, the reduction is optimal up to a factor γ_n. Under conjecture 13, this basically answers the main open problem left in the Proofs of Space paper [DFKP15]. In the theorem below we need the label length w to in the order of $m \log(n)$ to get a lower bound on $|\sigma_{\mathsf{init}}|$. For the proofs of space application, where $m = 1$, this is a very weak requirement, but for scrypt, where $m = n$, this means we require rather long labels (the number of queries q will be $\le n^2$, so the $\log(q)$ term can be ignored).

Theorem 14. *Consider any $G = (V, E), C \subseteq V, m \in \mathbb{N}, \epsilon \geq 0$ and algorithm* A. *Let $n = |V|$ and γ_n be as in Conjecture 13. Let $\mathcal{H}$ contain all functions $\{0,1\}^* \to \{0,1\}^w$, then with probability $1 - 2^{-\Delta}$ over the choice of $\mathsf{h} \leftarrow \mathcal{H}$ the following holds for every $\sigma_{\mathsf{init}} \in \{0,1\}^*$. Let q be an upper bound on the total number of h queries made by A and let*

$$k = \frac{|\sigma_{\mathsf{init}}| + \Delta}{(w - m \log(n) - \log(q))}$$

(so $|\sigma_{\mathsf{init}}| \approx k \cdot w$ for sufficiently large w), then

$$\mathsf{time}^{p\mathit{ROM}}(G, C, m, \mathsf{A}, \sigma_{\mathsf{init}}, \mathsf{h}) \geq \mathsf{time}(G, C, m, \lceil k \cdot \gamma_n \rceil)$$

and for every $1 > \epsilon \geq 0$

$$\mathsf{time}_\epsilon^{p\mathit{ROM}}(G, C, m, \mathsf{A}, \sigma_{\mathsf{init}}, \mathsf{h}) \geq \mathsf{time}_\epsilon(G, C, m, \lceil k \cdot \gamma_n \rceil)$$

In other words, if the initial state is roughly $k \cdot w$ bits large (i.e., it's sufficient to store k labels), then the pROM time complexity is as large as the pebbling time complexity of $\mathsf{pebble}(G, C, m)$ for any initial pebbling of size $k \cdot \gamma_n$ (which is $O(k)$ assuming our conjecture). Note that the above theorem is basically tight up to the factor γ_n: consider an experiment $\mathsf{time}(G, C, m, \mathsf{P}, P_{\mathsf{init}})$, then we can come up with a state σ_{init} of size $k \cdot w$, namely $\sigma_{\mathsf{init}} = \{\ell_i, i \in P_{\mathsf{init}}\}$, and define A to mimic P, which then implies

$$\mathsf{time}_\epsilon^{p\mathit{ROM}}(G, C, m, \mathsf{A}, \sigma_{\mathsf{init}}, \mathsf{h}) = \mathsf{time}_\epsilon(G, C, m, \mathsf{P}, P_{\mathsf{init}}) \quad \text{with} \quad |\sigma_{\mathsf{init}}| = k \cdot w$$

in particular, if we let $\mathsf{P}, P_{\mathsf{init}}$ be the strategy and initial pebbling of size k minimising time complexity we get

$$\mathsf{time}_\epsilon^{p\mathit{ROM}}(G, C, m, \mathsf{A}, \sigma_{\mathsf{init}}, \mathsf{h}) \geq \mathsf{time}_\epsilon(G, C, m, k) \quad \text{with} \quad |\sigma_{\mathsf{init}}| = k \cdot w$$

Wlog. we will assume that A is deterministic (if A is probabilistic we can always fix some "optimal" coins). We'll now prove two claims which will imply the theorem

Claim. With probability $1 - 2^{-\Delta}$ over the choice of $\mathsf{h} \leftarrow \mathcal{H}$. If the transcript $T(G, C, m, \mathsf{A}, \sigma_{\mathsf{init}}, \mathsf{h})$ is k-extractable, then

$$|\sigma_{\mathsf{init}}| \geq k \cdot (w - m \log(n) - \log(q)) - \Delta \tag{19}$$

where q is an upper bound on the total number of h queries made by A.

Proof. Let L be an upper bound on the length of queries made by A, so we can assume that the input domain of h is finite, i.e., $\mathsf{h} : \{0,1\}^{\leq L} \to \{0,1\}^w$. Let $|\mathsf{h}|$ denote the size of h's function table.

Let $\ell_{i_1}, \ldots, \ell_{i_k}$ be the indices of the k labels (these must not be unique) that can be "extracted", and let h^- denote the function table of h, but where the rows are in a different order (to be defined), and the rows corresponding to the queries that output the labels to be extracted are missing, so $|\mathsf{h}| - |\mathsf{h}^-| = k \cdot w$.

Given the state σ_{init}, the function table of h^- and some extra information α discussed below, we can reconstruct the entire function table of h. As this table is uniform, and a uniform string of length s cannot be compressed below $s - \Delta$ bits except with probability $2^{-\Delta}$, we get that with probability $1 - 2^{-\Delta}$ eq.(19) must hold, i.e.,

$$|\sigma_{\mathsf{init}}| + |\mathsf{h}^-| + |\alpha| \geq |\mathsf{h}| - \Delta$$

as $|\mathsf{h}| - |\mathsf{h}^-| = k \cdot w$ we get

$$|\sigma_{\mathsf{init}}| \geq k \cdot w - |\alpha| - \Delta$$

It remains to define α and the order in which the values in h^- are stored. For every label to be extracted, we specify on what challenge sequence to run the adversary A, and where exactly in this execution the label we want to extract appears (as part of a query made by A). This requires up to $m \log(n) + \log(q)$ bits for every label to be extracted, so

$$|\alpha| \leq k \cdot (m \cdot \log(n) + \log(q))$$

The first part of h^- now contains the outputs of h in the order in which they are requested by the extraction procedure just outlined (if a query is made twice, then we have to remember it and not simply use the next entry in h^-). Let us stress thaw we only store the w bit long outputs, not the inputs, this is not a problem as we learn the corresponding inputs during the extraction procedure. The entries of h which are not used in this process and are not extracted labels, make up the 2nd part of the h^- table. As we know for which inputs we're still missing the outputs, also here we just have to store the w bit long outputs such that the inputs are the still missing inputs in lexicographic order.

Let us mention that if A behaved nice in the sense that all its queries were the real requires required to compute labels, then we would only need $\log(n)$ bits extra information per label, namely the indices $i_1, \ldots, i_k$. But as A can be have arbitrarily, we can't tell when A actually uses real labels as inputs or some junk, and thus must exactly specify where the real labels to be extracted show up.

Claim. If the transcript $T = T(G, C, m, \mathsf{A}, \sigma_{\mathsf{init}}, \mathsf{h})$ is k-extractable (i.e., $ex(T) = k$), then

$$\mathsf{time}^{\mathtt{pROM}}(G, C, m, \mathsf{A}, \sigma_{\mathsf{init}}, \mathsf{h}) \geq \mathsf{time}(G, C, m, \lceil k \cdot \gamma_n \rceil) \qquad (20)$$

and for any $1 > \epsilon \geq 0$

$$\mathsf{time}^{\mathtt{pROM}}_\epsilon(G, C, m, \mathsf{A}, \sigma_{\mathsf{init}}, \mathsf{h}) \geq \mathsf{time}_\epsilon(G, C, m, \lceil k \cdot \gamma_n \rceil) \qquad (21)$$

Proof. We will only prove the first statement eq.(20). As T is k-extractable, there exist $(\mathsf{P}, P^{\updownarrow})$ where $P^{\updownarrow}$ is of weight $\leq \lceil k \cdot \gamma_n \rceil$ such that

$$\mathsf{time}^{\updownarrow}(G, C, m, \mathsf{P}, P^{\updownarrow}) = \mathsf{time}^{\mathtt{pROM}}(G, C, m, \mathsf{A}, \sigma_{\mathsf{init}}, \mathsf{h})$$

The claim now follows as

$$\mathsf{time}^{\updownarrow}(G, C, m, \mathsf{P}, P^{\updownarrow}) \geq \mathsf{time}^{\updownarrow}(G, C, m, \lceil k \cdot \gamma_n \rceil) = \mathsf{time}(G, C, m, \lceil k \cdot \gamma_n \rceil)$$

where the first inequality follows by definition (recall that $w(P^{\updownarrow}) \leq \lceil k \cdot \gamma_n \rceil$) and the second by Lemma 3 which states that for time complexity, entangled pebblings are not better than normal ones.

Theorem 14 follow directly from the two claims above.

4.4 The cumulative pROM complexity of the chain graph under Conjecture 13

Throughout this section $L_n = (V, E), V = [n], E = \{(i, i+1) \; : \; i \in [n-1]\}$ denotes the chain of length n, and the set of challenge nodes $C = [n]$ contains all verticies. In Section 3 we showed that – with overwhelming probability over the choice of a function $h : \{0,1\}^* \to \{0,1\}^w$ – the cumulative parallel entangled pebbling complexity for pebbling n challenges on a chain of length n is

$$cc^{\updownarrow}(L_n, C = [n], n, n) = \Omega\left(\frac{n^2}{\log^2(n)}\right)$$

this then implies a lower bound on he cumulative memory complexity in the pROM against the class $\mathcal{A}^{\updownarrow}$ of adversaries which are only allowed to store "encoding" of labels.

$$cmc_{\mathcal{A}^{\updownarrow}}^{pROM}(L_n, C = [n], n, n) = \Omega\left(w \cdot \frac{n^2}{\log^2(n)}\right)$$

This strengthens previous lower bounds which only proved lower bounds for cc complexity, which then implied security against pROM adversaries that could only store plain labels. In Section A we showed that $cc^{\updownarrow}$ can be strictly lower than cc, thus, at least for some graphs, the ability to store encodings, not just plain labels, can decrease the complexity.

In this section we'll show a lower bound on $cmc^{pROM}(G, C, m)$, i.e., without making any restrictions on the algorithm. Our bound will again depend on the parameter γ_n from Conjecture 13. We only sketch the proof as it basically follows the proof of Theorem 4.

Theorem 15. *For any $n \in \mathbb{N}$, let $L_n = (V = [n], E = \{(i, i+1) \; : \; i \in [n-1]\}$ be the chain of length n and γ_n be as in Conjecture 13, and the label length $w = \Omega(n \log n)$, then*

$$cmc^{pROM}(L_n, C = [n], n, \sigma_{\mathsf{init}}) = \Omega\left(w \cdot \frac{n^2}{\log^2(n) \cdot \gamma_n}\right)$$

and for every $\epsilon > 0$

$$cmc_{\epsilon}^{pROM}(L_n, C = [n], n, \sigma_{\mathsf{init}}) = \Omega_{\epsilon}\left(w \cdot \frac{n^2}{\log^2(n) \cdot \gamma_n}\right)$$

Proof (Proof sketch). We consider the experiment $\mathsf{computeLabel}(L_n, C, n, \mathsf{A}, \sigma_{\mathsf{init}}, \mathsf{h})$ for the A achieving the minimal $\mathsf{cmc}^{\mathsf{PROM}}$ complexity if h is chosen at random (we can assume A is deterministic). Let $(\mathsf{P}, P_{\mathsf{init}})$ be such that $\mathsf{pebble}^{\updownarrow}(L_n, C, n, \mathsf{P}, P_{\mathsf{init}})$ mimics (as defined above) this experiment. By Theorem 9, $\mathsf{cc}^{\updownarrow}(L_n, C = [n], n, n) = \Omega\left(n^2/\log^2(n)\right)$, unfortunately – unlike for time complexity – we don't see how this would directly imply a lower bound on $\mathsf{cmc}^{\mathsf{PROM}}$.

Fortunately, although Theorems 4 and 9 are about *cc* complexity, the proof is based on time complexity: At any timepoint the "potential" of the current state lower bounds the time required to pebble a random challenge, and if the potential is small, then the state has to be large (cf. Claim 7).

For any $0 \le i \le n$ and $\boldsymbol{c} \in C^i$ let $\sigma_{\boldsymbol{c}}$ denote the state in the experiment $\mathsf{computeLabel}(L_n, C, n, \mathsf{A}, \sigma_{\mathsf{init}} = \emptyset, \mathsf{h})$ right after the i'th label has been computed by A and conditioned on the first i challenges being $\boldsymbol{c}$ (as A is deterministic and we fixed the first i challenges, $\sigma_{\boldsymbol{c}}$ is well defined).

At this point, the remaining experiment is $\mathsf{computeLabel}(L_n, C, n-i, \mathsf{A}, \sigma_{\boldsymbol{c}}, \mathsf{h})$. Similarly, we let $P_{\boldsymbol{c}}$ denote the pebbling in the "mimicing" $\mathsf{pebble}^{\updownarrow}(L_n, C, n - i, \mathsf{P}, P_{\boldsymbol{c}})$ experiment after P has pebbled the challenge nodes $\boldsymbol{c}$. Let $P'_{\boldsymbol{c}}$ be the entangled pebbling of the smallest possible weight such that there exists a P' such that $\mathsf{pebble}^{\updownarrow}(L_n, C, n - i, \mathsf{P}, P_{\boldsymbol{c}})$ and $\mathsf{pebble}^{\updownarrow}(L_n, C, n - i, \mathsf{P}', P'_{\boldsymbol{c}})$ make the same queries on all possible challenges.

The expected time complexity to pebble the $i+1$'th challenge in $\mathsf{pebble}^{\updownarrow}(L_n, C, n - i, \mathsf{P}', P'_{\boldsymbol{c}})$ – and thus also in $\mathsf{computeLabel}(L_n, C, n - i, \mathsf{A}, \sigma_{\boldsymbol{c}}, \mathsf{h})$ – is at least $n/2 \cdot w(P'_{\boldsymbol{c}})$ by Claim 7. And by Theorem 14, we can lower bound the size of the state $\sigma_{\boldsymbol{c}}$ as (assuming w is sufficiently larger than $\log(n)$)

$$|\sigma_{\boldsymbol{c}}| \ge \Omega(w \cdot w(P'_{\boldsymbol{c}})/\gamma_n)$$

The *cc* cost of computing the next $(i + 1)$th label in $\mathsf{computeLabel}(L_n, C, n - i, \mathsf{A}, \sigma_{\boldsymbol{c}}, \mathsf{h})$ – *if we assume that the state remains of size at lest $|\sigma_{\boldsymbol{c}}|$ until this label is computed* – is

$$\frac{n}{2 \cdot w(P'_{\boldsymbol{c}})} \cdot |\sigma_{\boldsymbol{c}}| = \Omega\left(\frac{n}{w(P'_{\boldsymbol{c}})} \cdot \frac{w \cdot w(P'_{\boldsymbol{c}})}{\gamma_n}\right) = \Omega\left(\frac{n \cdot w}{\gamma_n}\right)$$

As there are n challenges, this would give an $\Omega(n^2 \cdot w/\gamma_n)$ bound on the overall cc complexity. Of course the above assumption that the state size never decreases is not true in general, we don't want to make any assumptions on A's behaviour.

In the above argument, we don't actually use the size $|\sigma_{\boldsymbol{c}}|$ of the current state, but rather the potential of the lightest pebbling $P'_{\boldsymbol{c}}$ necessary to mimic the remaining experiment. Following the same argument as in Theorem 4 (in particular, using Claim 8) on can show that the potential must remain almost constant for some of the n challenges. This argument will lose us a $1/\log^2(n)$ factor in the cc complexity, giving a $\Omega\left(w \cdot \frac{n^2}{\log^2(n) \cdot \gamma_n}\right)$ as claimed.

Acknowledgments

References

[AABS14] Leonardo C. Almeida, Ewerton R. Andrade, Paulo S. L. M. Barreto, and Marcos A. Simplicio Jr. Lyra: Password-based key derivation with tunable memory and processing costs. Cryptology ePrint Archive, Report 2014/030, 2014. http://eprint.iacr.org/2014/030.

[AS15] Joël Alwen and Vladimir Serbinenko. High parallel complexity graphs and memory-hard functions. In Rocco A. Servedio and Ronitt Rubinfeld, editors, 47th ACM STOC, pages 595–603. ACM Press, June 2015.

[BDK15] Alex Biryukov, Daniel Dinu, and Dmitry Khovratovich. Fast and tradeoff-resilient memory-hard functions for cryptocurrencies and password hashing. Cryptology ePrint Archive, Report 2015/430, 2015. http://eprint.iacr.org/2015/430.

[BK15] Alex Biryukov and Dmitry Khovratovich. Tradeoff cryptanalysis of memory-hard functions. Cryptology ePrint Archive, Report 2015/227, 2015. http://eprint.iacr.org/2015/227.

[CJM14] Time memory tradeoff analysis of graphs in password hashing constructions. pages 256–266, 2014.

[DFKP15] Stefan Dziembowski, Sebastian Faust, Vladimir Kolmogorov, and Krzysztof Pietrzak. Proofs of space. In Rosario Gennaro and Matthew J. B. Robshaw, editors, CRYPTO 2015, Part II, volume 9216 of LNCS, pages 585–605. Springer, Heidelberg, August 2015.

[DKW11] Stefan Dziembowski, Tomasz Kazana, and Daniel Wichs. One-time computable self-erasing functions. In Yuval Ishai, editor, Theory of Cryptography, volume 6597 of Lecture Notes in Computer Science, pages 125–143. Springer Berlin Heidelberg, 2011.

[DNW05] Cynthia Dwork, Moni Naor, and Hoeteck Wee. Pebbling and proofs of work. In Victor Shoup, editor, Advances in Cryptology – CRYPTO 2005, volume 3621 of Lecture Notes in Computer Science, pages 37–54. Springer Berlin Heidelberg, 2005.

[FLW13] Christian Forler, Stefan Lucks, and Jakob Wenzel. Catena: A memory-consuming password scrambler. Cryptology ePrint Archive, Report 2013/525, 2013. http://eprint.iacr.org/2013/525.

[Lee11] Lee, Charles. Litecoin, 2011.

[Mar]   Crypto-Currency Market Capitalizations. `http://coinmarketcap.com/`. Accessed: 2015-10-07.

[Per09] Colin Percival. Stronger key derivation via sequential memory-hard functions. 2009.

[PJ]   C. Percival and S. Josefsson. The scrypt password-based key derivation function. `https://datatracker.ietf.org/doc/draft-josefsson-scrypt-kdf/?include_text=1`. Accessed: 2015-10-07.

[PPA+15] Sunoo Park, Krzysztof Pietrzak, Joël Alwen, Georg Fuchsbauer, and Peter Peter Gaži. Spacecoin: A cryptocurrency based on proofs of space. Cryptology ePrint Archive, Report 2015/528, 2015. `http://eprint.iacr.org/2015/528`.

A Entanglement Can Improve Cumulative Complexity.

In this section we show that – unlike for time complexity – for cumulative complexity entanglement helps: For the graph in Figure 2 given challenge set $C = \{C_1, \ldots, C_8\}$ and letting $m = k = 1$ we get

$$\mathsf{cc}^{\updownarrow}(G, C, m = 1, k = 1, 0) = 3n + 4 \tag{22}$$

$$\mathsf{cc}(G, C, m = 1, k = 1, 0) \geq 3.5n \tag{23}$$

Thus, the cumulative complexity of the best standard pebbling (for $m = 1$ rounds with an initial pebbling of weight 1) is a factor $7/6$ more expensive than the best entangled pebbling with the same parameters. This constant can be improved arbitrarily close to 1.5. We leave it as an interesting open question whether it is possible to come up with a family of graphs where this gap is a superconstant factor.

To see eq.(22), consider the initial pebbling configuration $P_{\mathsf{init}} = \{\langle A, B\rangle_1\}$ and suppose we are given a challenge C_i with $i \in [n]$. We keep the entangled pebble $\langle A, B\rangle_1$, and if $i \leq 4$ then pebble the nodes n and $2n$ simultaneously (by pebbling $1, n + 1$, then $2, n + 2$, etc.). Otherwise, when $i \geq 5$, pebble nodes $3n$ and $4n$ simultaneously. In either case the is cost $3n$ (as it takes n steps where at any step the entire entangled pebbling configuration has weight 3). Once we have n and $2n$ pebbled (or $3n$ and $4n$ as the case may be), one more step gives us A (or B), and then from $\langle A, B\rangle_1$ we get B (or A). Thus, in the next step we can immediatly pebble the challenge node C_i. So, cumulative complexity (sum of the weight of the pebbling over all timesteps) is $3n + 4$.

To see eq.(23), one first must convince oneself that the best initial pebbling is to put a pebble on either A or B (the reason we have as many as 8 challenge nodes is to make sure that putting the pebble on any one of them will not be optimal). If we have a pebble on A, then we can pebble C_i for any $i \geq 5$ at a cost of $3n + O(1)$, while for $i \leq 4$ this pebble is useless, and the cost is $4n + O(1)$, thus, in expectation the cc cost will be $3.5n + O(1)$.
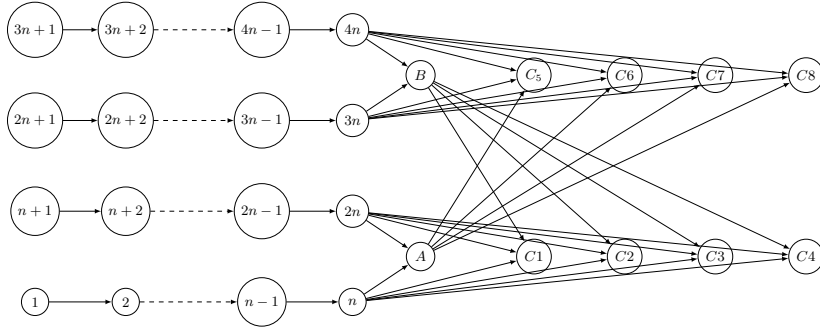
Fig. 2: A graph whose *entangled* pebbling complexity is lower than its classical pebbling complexity.

B   Neither Extractability nor Coverability are Optimal

As we'll show in Section 4.3, given a transcript $T := T(G, C, m, \mathsf{A}, \sigma_{\mathsf{init}}, \mathsf{h})$, the extractability $ex(T)$ implies a lower bound on the size of the state σ_{init} of $w \cdot ex(T)$ (w being the length of the labels, in this informal discussion we ignore additive terms of size $O(n \log n)$). Whereas the weight $w(T)$ gives a lower bound of $\mathsf{time}^{\updownarrow}(G, C, m, w(T))$ (i.e., in terms of lower bonds on time complexity of the entangled pebbling game) on the time complexity of $\mathsf{computeLabel}(G, C, m, \mathsf{A}, \sigma_{\mathsf{init}}, \mathsf{h})$. At the same time $w \cdot w(T)$ is an upper bound on the necessary size of the state σ_{init} in the sense that there exists an algorithm A' (which depends only on A) such that we always can find a σ_{init}', $|\sigma_{\mathsf{init}}'| \leq w \cdot w(T)$ where $T(G, C, m, \mathsf{A}', \sigma_{\mathsf{init}}', \mathsf{h}) = T$.

Thus, the bound we will prove is tight up to a factor $w(T)/ex(T)$, which we conjecture to be constant. For the transcript T^* from eq.(18) we showed $w(T)/ex(T) = 3/2 = 1.5$, which raises the question which of the two quantities is not "optimal". The answer is both, and T^* is an example for which neither quantity is optimal.

Using an automated prover for Shannon inequalities[13], we showed that any state which is sufficient to realize the implications form T^* must have size at least $2.5 \cdot w$. Concretely, if the labels are uniform of length w, which in particular implies their Shannon entropy is at least w, i.e.,

$$H(\ell_1) = H(\ell_2) = \ldots = H(\ell_5) = w$$

and the implications from T^* can be satisfied given some initial state σ_{init}, which means

$$H(\ell_1 | \ell_5, \ell_2, \sigma_{\mathsf{init}}) = H(\ell_2 | \ell_1, \ell_3, \sigma_{\mathsf{init}}) = \ldots = H(\ell_5 | \ell_4, \ell_1, \sigma_{\mathsf{init}}) = 0 \qquad (24)$$

---

[13] http://xitip.epfl.ch/

and moreover the ℓ_i are all independent, then the basic Shannon inequalities imply that the state must be of size at least

$$H(\sigma_{\mathsf{init}}) \geq 2.5 \cdot w$$

For even w, we can show that a state σ_{init} of length $2.5 \cdot w$ bits satisfying eq.(24) exists. Let $L_i \| R_i = \ell_i, L_i, R_i \in \{0,1\}^{w/2}$, then

$$\sigma_{\mathsf{init}} = \{L_1 \oplus L_2, R_2 \oplus R_3, L_3 \oplus L_4, R_4 \oplus R_5, L_5 \oplus R_1\}$$

satisfies eq.(24). For example $H(\ell_1 | \ell_5, \ell_2, \sigma_{\mathsf{init}}) = 0$ as $\ell_1 = R_1 \| L_1$ can be computed from the conditional part (and thus has no entropy) as $R_1 = (L_5 \oplus R_1) \oplus L_5$ and $L_1 = (L_1 \oplus L_2) \oplus L_2$.

For a transtcript T, let $shannon(T)$ denote the lower bound (divided by the label length w) on the size of a state required to satisfy the implications in T, so $shannon(T^*) = 2.5$, and in general $ex(T) \leq shannon(T) \leq w(T)$. We conjecture

Conjecture 16. Let

$$\Gamma_n = \max_{T \text{ over } [n]} \frac{w(T)}{shannon(T)}$$

then for some fixed constant Γ, $\Gamma_n \leq \Gamma$ for all $n \in \mathbb{N}$.

This conjecture is presumably weaker than Conjecture 13 as $shannon(T) \geq ex(T)$ for any T. Unfortunately Shannon entropy of a variable is the *expected* length of its shortest encoding, whereas $ex(T)$ gives us a lower bound of almost $w \cdot ex(T)$ with high probability. For his reason $shannon(T)$ seem much less convenient to work with than $ex(T)$, and we state our results in terms of γ_n from Conjecture 13, not Γ_n.

C Proof of Lemma 5

Proof (Lemma 5). For an arbitrary pebbler $\mathtt{T}'$ and initial pebbling configuration P_{init} (s.t. $|P_{\mathsf{init}}| \leq k$), we define $\mathsf{lc\text{-}pebble}_{n,t\cdot\lambda}(\mathtt{T}', P_{\mathsf{init}})$ to be the output cost of game $\mathsf{pebble}(L_n, C = [n], t \cdot \lambda, \mathtt{T}', P_{\mathsf{init}})$. In particular, let c be any value such that

$$\mathbb{P}(\mathsf{lc\text{-}pebble}_{n,t\cdot\lambda}(\mathtt{T}', P_{\mathsf{init}}) \geq c) \leq 1 - \delta .$$

Assume without loss of generality the pebbler $\mathtt{T}'$ is legal and deterministic. We are going to build a new pebbler $\mathtt{T} = \mathtt{T}_\lambda$ for $\mathsf{expect}_{n,t}$ such that

$$\mathbb{P}(\mathsf{expect}_{n,t}(\mathtt{T}, P_{\mathsf{init}}) \geq 2c) \leq 1 - \delta + 3\mu(\lambda) . \tag{25}$$

Note that this implies the theorem, since by taking the appropriate pebbler $\mathtt{T}'$ and configuration P_{init}, we can make c as close as we wish to $\mathsf{cc}_\delta(L_n, C = [n], t \cdot \lambda, k)$. But then, this implies in particular that $\mathsf{cc}_{\delta-3\mu(\lambda)}(\mathsf{expect}_{n,t,k}) \leq 2c$ by (25).

More concretely, the (randomized) pebbler $\mathtt{T}$ is going to internally simulate the $(t \cdot \lambda)$-round execution of game $\mathsf{lc\text{-}pebble}_{n,t\cdot\lambda}$ with pebbler $\mathtt{T}'$ and configuration P_{init}. Here, for simplicity of explanation, we denote $P'_{i,j}$ to be the jth pebbling configuration of round i in pebbling game; and $P_{i,j}$ to be the jth pebbling configuration of round i in expectation game. In each round $i \in [t]$, $\mathtt{T}$ simulates the λ consecutive rounds $r(i)+1, \ldots, r(i)+\lambda$ of the execution of $\mathtt{T}'$ with the challenger from $\mathsf{lc\text{-}pebble}_{n,t\cdot\lambda}$, where $r(i) = (i-1) \cdot \lambda$. This in particular entails sampling λ challenges $v_{r(i)+1}, \ldots, v_{r(i)+\lambda} \xleftarrow{\$} [n]$, and running $\mathtt{T}'$ (from its current state) being fed these challenges. Then, $\mathtt{T}$ simply checks that the returned sequence P' of configurations (which includes the configurations returned in all λ rounds)

$$P' = (P'_{r(i)+1,1}, \ldots, P'_{r(i)+1,t_{r(i)+1}}), \ldots, (P'_{r(i)+\lambda,1}, \ldots, P'_{r(i)+\lambda,t_{r(i)+\lambda}})$$

are made of at least $\phi(P'_{r(i)+1,0})$ sets, i.e., $t(i) := \sum_{j=1}^{\lambda} t_{r(i)+j} \geq \phi(P'_{r(i)+1,0})$. If so, $\mathtt{T}$ simply outputs $P_i = P'$ for the i-th round. Otherwise, $\mathtt{T}$ simply obtains P_i by extending P' to contain $\phi(P'_{r(i)+1,0})$ sets before outputting, for example by repeating every configuration in P' v times, for the smallest v such that $v \cdot t(i) \geq \phi(P'_{r(i)+1,0})$. Note that this increases by a factor v the cumulative cost of the pebbling output in this round by $\mathtt{T}$, and thus we want to keep v as small as possible, say $v \leq 2$. In particular, note that $t(i) \geq \lambda$, since $t_{r(i)+j} \geq 1$ for all $i \in [\lambda]$, thus this is guaranteed implicitly when $\phi(P'_{r(i)+1,0}) \leq 2\lambda$.

By construction, it is not hard to see that $\mathtt{T}$ is a legal pebbler for $\mathsf{expect}_{n,t}(\cdot, P_{\mathsf{init}})$, i.e., the game never output ∞. However, we now need to show that its cc does not grow too much compared to that of $\mathtt{T}'$ in $\mathsf{lc\text{-}pebble}_{n,t\cdot\lambda}(\mathtt{T}', P_{\mathsf{init}})$. To do this, we are going to prove that except with probability μ over the choice of the $t \cdot \lambda$ challenges, we have $t(i) \geq \phi(P'_{r(i)+1,0})/2$, or in other words, we can always pick $v \leq 2$ in all t rounds with probability at least $1 - \mu$.

Before we prove this fact, note that this is enough to conclude the proof, as in particular this implies that with probability $1 - \mu$ over the choices of the challenges (denote the corresponding event as G), we have $\mathsf{expect}_{n,t}(\mathtt{T}, P_{\mathsf{init}}) \leq 2 \cdot \mathsf{lc\text{-}pebble}_{n,t\cdot\lambda}(\mathtt{T}', P_{\mathsf{init}})$, and thus

$$\mathbb{P}(\mathsf{expect}_{n,t}(\mathtt{T}, P_{\mathsf{init}}) \geq 2c) \leq \mathbb{P}(\neg G) + \mathbb{P}(\mathsf{expect}_{n,t}(\mathtt{T}, P_{\mathsf{init}}) \geq 2c \mid G)$$

$$\leq \mu + \mathbb{P}(\mathsf{lc\text{-}pebble}_{n,t\cdot\lambda}(\mathtt{T}', P_{\mathsf{init}}) \geq c \mid G)$$

$$\leq \mu + \frac{1-\delta}{1-\mu} \leq \mu + (1-\delta)(1+2\mu) = 1 - \delta + 3\mu \ .$$

where we have used the fact that for any two events A and B, $\mathbb{P}(A \mid B) \leq \mathbb{P}(A)/\mathbb{P}(B)$.

We are now left with proving that $\mathbb{P}(\neg G) \leq \mu$. To this end, let $P'_{r(i)+1,0} = \{\ell_1, \ell_2, \ldots, \ell_m\}$ be the starting configuration in round i of $\mathtt{T}$, respectively in round $r(i)+1$ of the simulated $\mathtt{T}'$. Consider the sampling of the next λ challenges $v_{r(i)+1}, \ldots, v_{r(i)+\lambda}$. Since P' output by $\mathtt{T}'$ is always legal, all of $v_{r(i)+1}, \ldots, v_{r(i)+\lambda}$

must be contained in one of its configurations. For every challenge v_j, let $d_j = v_j - \ell_{j^*}$, where ℓ_{j^*} is the largest element in $P'_{r(i)+1,0}$ which is not larger than v_j. Then, it is clear that we must have

$$t(i) \geq \max_{j \in [\lambda]} d_{r(i)+j} \geq \frac{1}{\lambda} \sum_{i=1}^{\lambda} d_{r(i)+j} =: \overline{d}(i) ,$$

since we need to take at least as many pebbling steps to cover, starting in $P'_{r(i)+1,0}$, the farthest of the λ challenges from a position which has a pebble in $P'_{r(i)+1,0}$.

Note that the expectation of $\overline{d}(i)$ (over the choice of the λ challenges) is exactly,

$$\mathbb{E}(\overline{d}(i)) = \frac{m}{n} + \frac{1}{2n} \sum_{i=0}^{m} (\ell_{i+1} - \ell_i) \cdot (\ell_{i+1} - \ell_i - 1) = \phi(P'_{r(i)+1,0}) .$$

Moreover, $\mathbb{E}(d(i)) = \mathbb{E}(d_{r(i)+j})$ for all $j \in [\lambda]$. Now let us assume that $\phi(P'_{r(i)+1,0}) \geq \lambda$. (Otherwise, $t(i) \geq \phi(P'_{r(i)+1,0})/2$ already.) By the Chernoff bound, we have

$$\mathbb{P}(t(i) < \phi(P'_{r(i)+1,0})/2) \leq \mathbb{P}(\overline{d}(i) < \phi(P'_{r(i)+1,0})/2)$$

$$= \mathbb{P}(\sum_{i=1}^{\lambda} d_{r(i)+j} \leq \lambda \phi(P'_{r(i)+1,0})/2) \leq e^{-\lambda^2/8} .$$

The final bound on μ follows by the union bound and the fact that there are t rounds.