Storage Efficient Substring Searchable Symmetric Encryption

Iraklis Leontiadis, Ming Li

New Jersey Institute of Technology and University of Arizona, USA
{leontiad@njit.edu,lim@email.arizona.edu}

Abstract. We address the problem of substring searchable encryption. A single user produces a big stream of data and later on wants to learn the positions in the string that some patterns occur. Although current techniques exploit auxiliary data structures to achieve efficient substring search on the server side, the cost at the user side may be prohibitive. We revisit the work of substring searchable encryption in order to reduce the storage cost of auxiliary data structures. Our solution entails a suffix array based index design, which allows optimal storage cost $O(n)$ with small hidden factor at the size of the string n. We identify the leakages of the scheme following the work of Curtmola et al. [CCS06] and we analyze the security of the protocol in the real ideal framework. Moreover, we implemented our scheme and the state of the art protocol Chase and Shen [POPETS15] to demonstrate the performance advantage of our solution with precise benchmark results. We improved the storage overhead of the encrypted index by a factor of **1.8** and the computation time thereof **4** times on 10^6 character data streams.

1 Introduction

Nowadays, there is a flourish of protocols delegated to run by an untrusted coalition of servers, systems, services, called hereafter the cloud. Due to the untrusted nature of the cloud, users seek to protect the privacy and security of their data with cryptographic primitives. The cloud on the other hand offers an economy of scale with the impressive resources it acquires, ranging from software to hardware. Uploading encrypted data however, renders operations on it infeasible. Downloading, decrypting and running the operation on plaintext data, cancels the advantages, that the cloud offers for large storage and computational efficiency. Usually users need to perform a search on their data. Tailored protocols for secure searchable encryption have been proposed in the literature, whereby single or multiple users upload encrypted documents, with some auxiliary data structure called an index, allowing the cloud to correctly return documents containing a single, multiple or a boolean function of keywords, without compromising index, query, and documents privacy. Apart from their theoretical consideration in the literature, quite a few companies adopt this model to offer searchable encryption schemes over encrypted data [3, 10, 11, 19, 26, 29, 31].

While keyword based search protocols are quite common in a large range of applications, they cannot efficiently address all the possible queries a user submits to the cloud. Substring based queries have come to the forefront due to the ubiquitousness of devices and the progress in storage technology. Devices produce a big stream of data, which needs to be queried later on with *substring based queries*. Namely a substring query for a stream of data, consists of a substring of the stream and the result is the position of the substring in the big stream, or/and the number of occurrences of multiple substrings.

Applications. In a health-care application, data enclaves which hold giant stream of medical information such as DNA sequencing are asked to answer substring queries by medical labs. The possible position of a substring in the whole DNA sequence of a single person gives information about predisposition to diseases. As such, it is treated as personal sensitive information and should be protected. Nowadays, the sequencing process is possible thanks to the progress of computers. Online services offer DNA sequencing to institutions and individuals. In the logging systems scenario, companies, institutions and organizations produce log data of giant size. The logs are recorded and uploaded in a cloud infrastructure to take advantage of the cheap storage space. Log data are often searched to identify malicious substring patterns. The position of the suspicious searched string token will act as a bookmark to further download the logs data, which proceed and succeed that position for further investigation. Deep packet inspection (DPI) is another application whereby a gateway, firewall, or Intrusion Detection System (IDS) on behalf of a user is looking for prohibitive content on a bigger stream. In general, the vast amount of information renders substring queries a real challenge and reducing the storage cost of the encrypted index would increase the performance of such services.

Protecting the privacy of the data stream and the substring query, while allowing an untrusted cloud to correctly answer substring matching pattern efficiently and securely is not trivial. Following the searchable encryption approach, separating the data itself from the index, results in a prohibitive storage index cost $O(n^2)$, where n is the size of the stream. The index consists of all possible substrings of a stream of data of size n and the encrypted data are the positions of the substring. Recently Chase *et al.* [9] proposed a solution that asymptotically achieves $O(n)$ storage costs by exploiting the auxiliary data structure of the suffix tree. However the asymptotic costs of $O(n)$ hide a constant factor that can be roughly up to 20 [1, 5, 24] for the construction of the suffix tree due to the complexity of the tree and the extra pointers to traverse a tree.

Our goal, which launches our research is to reduce the increased size of index for substring queries, which has to be computed once and be kept at the cloud during the entire lifetime of the protocol : *Can we significantly reduce the expected storage cost of an encrypted index tailored for substring queries without sacrificing search time and privacy?* Even a small improvement would have big impact as the data to be indexed for substring queries span to million of elements. Moreover, besides the cloud side cost cuts, the client will also be positively affected as the smaller the size of the index it is outsourced, the less the charges clients commit to the cloud.

Idea. After encrypting the suffix tree of Chase *et al.* scheme [9], the encrypted structure leaks a lot information concerning the internal structure of the tree as number of leaves, children and double "touched" branches. The authors suggested a dummy node policy in order to hide as much information as possible. After constructing the suffix tree with N nodes, the suffix tree is filled up with $2n - N$ internal nodes nodes. To each node with less than σ children, where σ is the size of the vocabulary, up to σ dummy nodes are appended. Encrypting all these dummy blocks drastically increases the storage overhead and subsequently the communication cost of the protocol for index construction.

Our core idea lies at the properties of a suffix array based indexing. A suffix array contains information about the position of each suffix of a string and has constant size n for a string of n size. In contrast, the data structure of the suffix tree has no constant size and can acquire up to $2n$ nodes, with each node storing information about its edges, parent and children nodes, thus increasing the storage need. By choosing the suffix array we decrease the storage need for the construction of the index. The second factor which lends us to less storage and subsequently communication efficiency is the dummy blocks policy which is used to hide the structure of the suffix tree in [9]. Our dummy node policy to obscure the encrypted suffix array rests only on the frequency of the most frequent character. Namely, we fill up the original string with characters such that the frequency of each character is the same. However this approach raises a shortcoming when applied to data sets with skewed frequency distributions such as text based data sets. We overcome that limitation with a bucketization technique. Instead of building the index on a single character approach we explore the idea of grouping together characters consisting a bucket. Surprisingly our experimental evaluation showed that the skewed frequencies are diminished and the final storage overhead for the index is considerably smaller by a factor of **1.8** for text based datasets.

In this paper we design and analyze a storage efficient Substring Searchable Symmetric Encryption ($\mathsf{S^3E}$) protocol with variable size of substrings. We follow a different approach from existing techniques that allows us to achieve the efficiency, functional and security goals we want. In our technique we exploit a self-indexed data structure, which allows the cloud to search for substring queries. Its form resembles the suffix arrays with some additional extra steps. The main contributions in the paper are summarized as follows:

Contributions:

- *Storage efficient Substring Searchable Symmetric Encryption* ($\mathsf{S^3E}$): Thanks to the employment of the suffix array, which achieves a small hidden factor (≈ 4) in the $O(n)$ asymptotic complexity, compared to the bigger (≈ 20) hidden factor of the suffix tree, our design presents a storage efficient substring searchable symmetric encryption protocol.
- *Variable substring query length*: Our solution allows a dynamic issue of substring queries of variable size without the need of defining a fixed query size beforehand.
- *Provably secure.* Our scheme is provably secure in the real-ideal simulation paradigm with similar leakages as the state of the art scheme in [9].
- *Prototype implementation*: We implemented our protocol and the state of the art work in [9]. We performed a real world comparison of both schemes based on computation time and communication overhead to build the index and query for a substring. Our results show a performance advantage of **1.8** storage overhead on text based datasets as the enron email one.

Outline. In section 2 we introduce the problem this paper addresses. Afterwards, in section 3 we review similar cryptographic protocols for substring searchable symmetric encryption. We continue in section 4 with preliminaries of our solution and the basic building blocks. In section 5 we illustrate in more details the design components of our protocol. The full protocol description is presented in section 6. We then investigate the security and the costs of the proposed scheme for storage efficient substring searchable symmetric encryption. In section 8 we present our prototype implementation results. Finally we conclude in section 9.

2 Problem Statement

In this section we formalize the problem of string matching. We first start with the functional requirements of substring matching and afterwards we present the security requirements of the protocol.

2.1 Functional Requirements

Herewith pattern matching, string matching and substring matching are used interchangeably in this paper. We assume that a string S is modeled as an one dimension array $S[1...n]$. A substring is another array $T[1...m]$. The elements of each array are drawn from some finite alphanumerical alphabet Σ of size $\sigma = |\Sigma|$. We say that a substring T occurs in S if there exists $s : 1 \leq s \leq n - m$ and $S[s+1...s+m] = T[1...m]$, meaning that $S[s+j] = T[j], 1 \leq j \leq m$ (cf figure 1).


Fig. 1: Pattern matching

Naive algorithms for pattern matching achieve $O(n)$ on search time and 0 cost on preprocessing. The algorithm simply scans all the positions $i, 1 \leq i \leq n - m$ of the string S until it finds m consecutive matches at a position $j, 1 \leq j \leq n - m + 1$. Trading preprocessing efficiency for better search costs, Robin Karp algorithm [21] achieves $O(n - m + 1)$ search time and $\Theta(m)$ preprocessing amortized cost. In a similar trajectory Knuth-Moris-Pratt [22] has $\Theta(n)$ search complexity and $\Theta(m)$ preprocessing time. Boyer-Moore pattern matching technique [4] increases the preprocessing cost at $\Theta(m + \sigma)$ in order to have worst case search complexity $O(n)$. Following a different trajectory substring matching techniques achieve $O(m)$ search time by leveraging a more sophisticated preprocessing step, in which the suffixes of all substring are computed along with their positions in the string S, be it suffix tree[25, 30, 32] or suffix array [24]. Suffix tree though has a more expensive space efficiency due to the extra information the suffix tree has to keep [1, 5, 18, 24]. This cost is translated to a constant factor that approximates ≈ 20, which is hidden in the $O(n)$ asymptotic storage cost of the suffix tree construction. As a first step to relax this storage extra hidden cost we choose to build upon the suffix array string matching approach which has a much simpler storage cost which approximates $4n$ [1].

We redraw upon the *queryable encryption* syntactical definition of [9], since we believe it follows a deceptive abstraction. Namely, the functional definition claims to capture a generic framework for searchable encryption, in the sense that a query $\mathcal{F}$ can be any function keyword query, or substring query. However, an encrypted searchable encryption scheme is a more generic protocol, since it can be used to solve the substring searchable encryption problem with the encrypted inverted index technique as shown in the introduction. As such, searchable and substring encryption schemes cannot be addressed by the same definitional framework. Furthermore, the nature of the problem and the solution for substring queries drastically varies from keyword searchable encryption, since the index contains the data and there are not two separate objects, meaning that the index for substring queries is *self-indexed*, since from the index you can recover the underlying data structure. In contrast in encrypted searchable encryption, there is a clear distinction between the index, and the data structure that holds the data (files with keywords). For these reasons we rewrite the functional definitional framework for substring searchable encryption.

Definition 1. *A Substring Searchable Symmetric Encryption scheme (*$\mathsf{S^3E}$*) is a collection of four polynomial time algorithms (*$\mathsf{KeyGen}, \mathsf{PreProcess}, \mathsf{SrchToken}, \mathsf{Search}$*) defined as follows:*

- $\mathsf{k} \leftarrow \mathsf{KeyGen}(1^\lambda)$*: It is a probabilistic algorithm that takes as input the security parameter in the unary form* 1^λ *and outputs the secret substring search key* k.
- $\mathsf{SES} \leftarrow \mathsf{PreProcess}(\mathsf{k}, S)$*: This algorithm takes as inputs the stream* S *and the secret key* k *and outputs the substring encrypted data structure* SES.
- $\mathsf{tk}_{T,S} \leftarrow \mathsf{SrchToken}(\mathsf{k}, T[1...m])$*: It is a probabilistic algorithm that takes as input the secret substring search key* k*, a string* $T[1...m]$ *and outputs a trapdoor to search for the string* T *on data stream* S*, through* SES.
- $(s, \perp) \leftarrow \mathsf{Search}(\mathsf{tk}_{T,S}, \mathsf{SES})$*: It is a deterministic algorithm which takes as input a trapdoor* $\mathsf{tk}_{T,S}$ *and an substring encrypted structure* SES *and outputs the positions* s *in* S *that substring* T *occurs, or* $\perp$ *otherwise.*

A substring searchable encryption scheme is correct if $\forall \lambda \in \mathbb{N}, \forall S \in \Sigma, \forall \mathsf{k} \leftarrow \mathsf{KeyGen}(1^\lambda), \forall \mathsf{SES} \leftarrow \mathsf{PreProcess}(\mathsf{k}, S), \forall \mathsf{tk}_{T,S} \leftarrow \mathsf{SrchToken}(\mathsf{k}, T[1...m]), \mathsf{Search}(\mathsf{tk}_{T,S}, \mathsf{SES})$ always returns the correct positions s in the string S or $\perp$ otherwise.

2.2 Security Model

Intuitively the security guarantee we ask for is **1**) given a probabilistic polynomial time adversary $\mathcal{A}$ with access to a substring encrypted structure SES, $\mathcal{A}$ cannot gain more partial information about the underlying stream of data S and **2**) given a set of trapdoor tokens for an adaptively generated set of queries $\mathbf{q} = (q_1, q_2, q_3, \ldots, q_o)$ associated with set of tokens $\mathbf{t} = (\mathsf{tk}_1, \mathsf{tk}_2, \mathsf{tk}_3, \ldots, \mathsf{tk}_o)$ $\mathcal{A}$ cannot learn anything for $\mathbf{q}$ and $\mathbf{t}$. Following the symmetric searchable encryption paradigm we know it is impossible to achieve those two security guarantees without leaking some extra information as the observed in [7, 8, 12].

We express the security guarantees of the protocol in terms of simutability [23]. First a leakage function $\mathcal{L}$ is defined, which expresses the leakage of a $\mathsf{S^3E}$ scheme to an adversary $\mathcal{A}$, through the transcripts of the protocol. The simulation framework assumes two worlds. The $\mathbf{Real}^{\mathsf{S^3E}}_{\mathcal{A}(\lambda)}$ world, whereby adversaries can corrupt the parties they want and the $\mathbf{Ideal}^{\mathsf{S^3E}}_{\mathcal{A},\mathcal{S}(\lambda)}$ one in which there is only benign behavior of each party. The security analysis narrows down to the design of a simulator $\mathcal{S}$, who tries to simulate the malicious behavior in the $\mathbf{Ideal}^{\mathsf{S^3E}}_{\mathcal{A},\mathcal{S}(\lambda)}$ world, only through access to the leakage function $\mathcal{L}$. We say that a protocol is secure if $\mathcal{S}$ simulates indistinguishable views of the adversary $\mathcal{A}$ in the $\mathbf{Ideal}^{\mathsf{S^3E}}_{\mathcal{A},\mathcal{S}(\lambda)}$ world.

The adversary $\mathcal{A}$ plays the role of a semi-honest cloud and during the two world we assume a challenger $\mathcal{C}$ who interacts with $\mathcal{A}$. We describe the two world in algorithmic details in what is follows:
$\mathbf{Real}^{\mathsf{S^3E}}_{\mathcal{A}(\lambda)}$ world:

- $\mathcal{C}$ runs $\mathsf{KeyGen}(1^\lambda)$ to obtain k.
- $\mathcal{A}$ chooses a string $S \in \Sigma$, sends it to $\mathcal{C}$ and $\mathcal{C}$ replies with $\mathsf{SES} \leftarrow \mathsf{PreProcess}(\mathsf{k}, S)$ to $\mathcal{A}$.
- $\mathcal{A}$ issues a polynomial number of adaptively chosen queries $\mathbf{q} = (q_1, q_2, q_3, \ldots, q_o)$ and receives from $\mathcal{C}$ a set of tokens $\mathbf{t} = (\mathsf{tk}_1, \mathsf{tk}_2, \mathsf{tk}_3, \ldots, \mathsf{tk}_o)$.
- Finally $\mathcal{A}$ outputs $v = (\mathsf{SES}, \mathbf{t})$.

$\mathbf{Ideal}^{\mathsf{S^3E}}_{\mathcal{A},\mathcal{S}(\lambda)}$ world:

- $\mathcal{A}$ outputs a string stream S.
- The simulator $\mathcal{S}$ through the leakage $\mathcal{L}$ generates SES and forwards it to $\mathcal{A}$.
- $\mathcal{A}$ issues a polynomial number of queries $\mathbf{q} = (q_1, q_2, q_3, \ldots, q_o)$. $\mathcal{S}$ replies to each of the queries through the leakage function $\mathcal{L}$ with $\mathbf{t} = (\mathsf{tk}_1, \mathsf{tk}_2, \mathsf{tk}_3, \ldots, \mathsf{tk}_o)$.
- Finally $\mathcal{A}$ outputs $v = (\mathsf{SES}, \mathbf{t})$.

Definition 2. *A substring searchable encryption scheme scheme is adaptively* $\mathcal{L}$*-semantically secure against a probabilistic polynomial time adversary A if there is exists a polynomial Simulator* $\mathcal{S}$ *such that for all polynomial time distinguishers* $\mathcal{D}$:

$$| \Pr[\mathcal{D}(v) = 1 : v \leftarrow \mathbf{Real}^{\mathsf{S^3E}}_{\mathcal{A}(\lambda)}] -$$
$$\Pr[\mathcal{D}(v) = 1 : v \leftarrow \mathbf{Ideal}^{\mathsf{S^3E}}_{\mathcal{A},\mathcal{S}(\lambda)}]| \leq \mathsf{neg}(\lambda)$$

3 Related work

Tailored substring searchable encryptions schemes have been proposed in the literature [9], [14], [13]. Here we present a detailed analysis of the state of the art in substring searchable encryption protocols. Chase *et al.* [9] leverage the auxiliary data structure of the suffix tree. A suffix tree is a compressed suffix trie, can be computed in time $O(n)$ and allows for substring search in O(m) time on a substring of size m. Its amortized storage cost $O(n)$ hides a big constant factor, which can goes up to 20 [1, 5, 18, 24]. In [13] the authors extended the efficient SSE scheme for boolean queries from [6] in order to support substring matching. The idea is to build an index of overlapping k-grams, to prepend its relevant position and encrypt it. When a user needs to perform a substring query, the cloud performs a conjunctive keyword search for all the k-grams of the substring and returns the position. The disadvantage of the scheme comes at the need of storing all the overlapping k-grams at the cloud, which will represent substrings.

In [14] the authors follow a different approach. Instead of taking the *index-then-encrypt* approach with fast symmetric cryptographic primitives, they modify the subset sum problem, which is used to build public key encryption schemes, in a means such that the cloud can solve it. This contradicts the security definition of subset sum problem, which asks for impossibility of an adversary to find a solution to a specific instance. More specifically the user uploads a special instance $\mathbf{T}$ of a subset sum problem such that given a trapdoor R_i associated with a substring, the solution can be solved in time O(m) by the cloud; with the special property that the integers which sum to $\mathbf{T}$, parametrized by R_i are the positions of the substring in the string. This technique hides also the search pattern but comes at the cost of fixed size substrings, that must be defined in the beginning of the protocol. Moreover the substring should be substantially small with respect to the big stream. Our solution in contrast allows variable size of substring of any size.

Recently, Blass and Moataz [27] strengthen the security requirements by hiding the search and access patterns, following the ORAM approach. By leveraging the Path ORAM technique and the suffix array construction for substring queries, the authors manage to reduce the bandwidth, with a binary recursive tree above the position map. Each node in the tree represents a Path ORAM of the binary search tree for the suffix array. However, in order the cloud to be able to perform an oblivious binary search has to keep track of all the suffixes, which blows up the storage cost for the server. Furthermore the need for storing the suffixes cancels out the suffix array storage advantage over suffix tree. Finally, due to the Path ORAM technique the user has to store a state logarithmic on the length of the string–for the position map. The extra security guarantees of the tailored ORAM scheme do not allow for efficient storage cost both at the client and the cloud side, which is the goal for our work. Sanders *et al.* [28] presented a public key substring searchable encryption, where the underlying substring algorithm is a variant of Rabin-Karp algorithm, thus suffering from linear substring search time. It is noteworthy though to mention that their goal is to decouple keyword generation from the encryption of data in order to allow searches even when the underlying text is not known when the index is built.

4 Preliminaries

In order to reduce the storage cost for our Substring Searchable Symmetric encryption scheme ($\mathsf{S^3E}$) we first substitute the storage expensive suffix tree of the state of art work in [9] with a suffix array `SA`. A suffix array for a string S of size n constitutes of an integer array of size n, which has at each position a pointer to the start of the matched suffix $T[1...m]$ in the string S. `SA` is lexicographically sorted with respect to all the possible suffixes and can be computed in linear time on the size of the string S. In order to look for the position of a substring, a binary search in `SA` is performed, which is used as an index to the original string. Thus, the running time for a substring search is $O(m + logn)$. Let us now consider a concrete example to uncover its details. Suppose S=lalakis. The algorithm for the suffix array proceeds as follows:

1. Compute all the suffixes starting from the right-most position: s, is, kis, akis, lakis, alakis, lalakis.
2. Lexicographically sort the suffixes: akis, alakis, is, kis, lakis, lalakis, s.
3. Find the position in S of each suffix from step 2 and store them in an array $SA = [4, 2, 6, 5, 3, 1]$
4. Output `SA`.

However, plugging the `SA` for a substring searchable symmetric encryption scheme raises some difficulties. We assume that the suffix array is encrypted under a secret key of the user. In order for the cloud to retrieve the correct encrypted index position from `SA` ought to run a binary search obliviously without learning the underlying string S, query substring T, or any of the suffixes. A solution to the problem is to use the technique presented by Gentry *et al.* [16], which allows for a single ORAM query in order to perform a binary search over encrypted data. However, in order to adapt this approach it is required from the server apart from the encrypted suffix array, to store the tree of the encrypted data, which would be an extra burden for its storage complexity, plus there should be one extra round of communication due to the ORAM protocol in order to rebuild the specific path.

In this paper we are taking a different approach, which achieves storage efficiency. We take advantage of the *self-indexed* data structure Ferragina-Manzini index, called hereafter FM index [15]. Namely, from FM index the untrusted cloud can answer substring queries by leveraging the suffix array `SA`, without the need for an ORAM query. The neat property of the FM index is that it can reconstruct the original string S with some extra auxiliary data structures, thanks to its instantiation from the Burrows-Wheeler Transformation algorithm (BWT) [5]. For the reconstruction it employs the LF mapping technique, thus there is no need to store the encrypted stream S. The FM index can be derived from `SA`, as such its computational overhead is almost for free, after the computation of the suffix array. We describe the core building blocks of the FM index in what it follows.

4.1 Pattern matching

In this section we describe the compressed index FM, that will be used for the construction of our Secure Pattern Matching (S^3E) protocol. The design lies heavily on the BWT transformation for compression of bit-strings and on a special `LF` mapping for the reconstruction of the original string from BWT. The BWT, along with the `LF` mapping technique and some auxiliary information are the basic blocks of the compressed index for substring queries.

BWT Transformation The Burrows-Wheeler Transformation (BWT) transforms a stream of data by leveraging the entropy of each character. In a nutshell, the data stream S is transformed to an encoding W such that compression algorithms provide high rate of compression. For ease of completeness we show the steps to transform an original stream S to W with BWT in algorithm 1. First, the algorithm appends the terminating symbol \$ to the input string S. Then, it builds the matrix W by permuting the symbol \$. At each iteration the permutation is appended as a new row to the matrix W. Finally the rows of W are sorted lexicographically in an ascending way. An example is shown in figure 2 for string $S = lalakis$. The upper table of the figure shows the permutations and the final shorted matrix W is shown at the bottom matrix. The transformation is the first step towards compression with the `LF` mapping that is shown next.

LF Mapping The `LF` Mapping technique takes the first `F` and last `L` columns from the BWT transformation and through an iterative process (algorithm 2) reconstructs the original string S. Starting from the first elements of each column from `F` and `L`, the algorithm employees `L` as an index to the `F` column. Each time the element of the `L` column is appended to a LIFO stack. The value at the current position will be used as an index for the `F` column for the next loop. An example is presented in figure 3. At the first iteration the pointer indicates the first position in both columns `F`, `L`. For the next iteration the `L` character 's' indicates the index for the first column `F`, which can be found at its last position with `F[7]` =s. The current character at the `L` column is appended to a stack D. For the next iteration the current character at the `L` column indicates the next index for the `F` column. The character i is pushed to the stack D. The procedure halts when the position at `L` is \$. Then the algorithm pops all elements from D and the initial string S is fetched.

Algorithm 1: BWT transformation

Input: String S
Output: BWT$(S) = W$
l=length(S)+1;
S.append($);
i=0;
while $i¡l$ **do**
    r_i=rotate(s,$) // *The rotate algorithm permutes the characters of the original string and returns the permuted string*;
    W.addrow(r_i) // *It adds the permuted row from the previous step to the matrix W*;
    $i++$;
end
return Sorted.W;

Fig. 2: BWT Transformation. The upper table shows the cyclic permutations for the string S=lalakis. As a first step for the construction of the BWT matrix the end symbol $ is appended at the end of S, which lexicographically precedes all alphabetical symbols. Then at each row a permutation of $ around the string is shown. The result is matrix BWM[length(S)+1][length(S)+1]. As a second step the rows are sorted lexicographically at the second bottom table of the figure starting from the first character of each string. The light gray first column shows the order of the characters after sorting and the last column is the result of the transformation BWT(S)=sllkaa$s.

Fig. 3: The LF mapping process is used to reconstruct the original string S from the transformed one after applying the BWT operation. Starting with the $ sign from the F column, the mapping progressively reconstructs the entire string S. The last column L is used as a "ladder step" to find the next ith index in the F column, which in turn maps to the ith entry in the L column. The entire procedure halts when L[i]==$

FM Index Suffix array vanilla construction has $O(n^2 log n)$ asymptotic computational cost. This stems from the fact that the n suffixes are first sorted by performing $O(n log n)$ comparisons and each comparison has cost n. Linear time algorithms have been designed by first constructing a suffix tree and then traversing it with a depth first edge in lexicographical order. However, our goal is to be storage efficient, meaning we want to eliminate the storage cost of a suffix tree, which practically approximates a constant factor of $20n$ [1, 5, 18, 24]. We pick up the *skew* algorithm [20] which is a *divide and conquer* based algorithm and achieves linear time construction. The approach of the *skew* algorithm is to recursively divide the suffixes in three groups depending on the position **pos** of all suffices: **pos** mod $h, h \in \{1, 2, 3\}$ and then merge the result.

The FM consists of three column arrays. The first one is the F column from the LF mapping, the second one is the L column, which corresponds to the BWT(S) and the last one corresponds to the suffix array SA. SA contains at each row i, the position in the original string S of the substring which corresponds to the i^{th} row of the W matrix obtained after applying the BWT transformation. L = BWT(S) can

---

Algorithm 2: LF Mapping

Input: First (F), Last column (L) from BWT
Output: S
D=0 // Initialize the stack D;
l=length.(F) // the length of F equals the length of L;
i=0;
while $L[i]! = \$$ **do**
  D.push(L[i]);
  i=find.F[L[i]]// find.[] denotes the index number in array [] that the element is. For instance
    find.F['s']=7 ;
while $D! = \backslash'0'$ **do**
  S=S+D.pop;
return S;

---

be computed with the formula $\mathsf{BWT}(S)[i] = S[\mathsf{SA}[i] - 1]$ from the suffix array. Furthermore for the traversal of the LF mapping the unique ranking of each character in each F, L needs to be stored in $r_\mathsf{F}, r_\mathsf{L}$ accordingly. Finally $\mathsf{FM} = \{\mathsf{F}[i], \mathsf{L}[i], r_\mathsf{F}[i], r_\mathsf{L}[i], \mathsf{SA}[i]\}_{i=1}^n$.

The entire challenge is on how a user encrypts FM in such a way that a untrusted cloud can correctly reply with the encrypted position on substring queries. We give an intuition of our approach in the next section, we also highlight some shortcoming thereof and we demonstrate our solution to alleviate it.

5 Intuition

In this section we provide some intuition about $\mathsf{S}^3\mathsf{E}$ protocol before delving into its precise description in the follow up section. First we start with showing how the client encrypts the index to allow the cloud process fast encrypted substring queries. Our solution is based on the FM index described in the previous section. The FM index consists of three arrays, which keep track of the F, L columns and the encrypted SA suffix array with positions of substrings and not all the suffixes as in [27]. To recap, the user computes the suffix-array SA and the F, L columns through the BWT transformation.

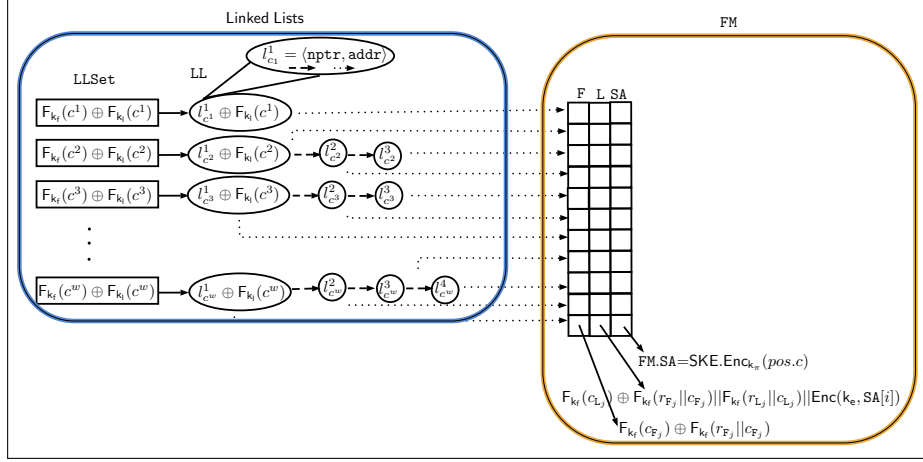|	
σ	Vocabulary size
Σ	Vocabulary
S	Original stream
T	Substring query
c	character
b	bucket
n	Size of S
m	Size of T
win	window size
SA	Suffix array
F	First Column of LF mapping
L	Last Column of LF mapping
LLSet	Hash map
LL	Linked List
$c_{\mathsf{F}_j}, b_{\mathsf{F}_j}$	j^{th} character,bucket from F column
$c_{\mathsf{L}_j}, b_{\mathsf{L}_j}$	j^{th} character,bucket from L column
r_{c_i}, r_{b_i}	Ranking of i^{th} character,bucket in the string S
r_{F_j}	Ranking of j^{th} character from F column
r_{L_j}	Ranking ofj^{th} character from L column
c^w	w^{th} character, bucket from the alphabet, $(1 \leq w \leq \sigma)$
b^z	z^{th} bucket,$(1 \leq z \leq n - \mathsf{win} + 1)$
λ	Security parameter
$\mathsf{F}(\cdot)$	Pseudorandom function (PRF)
$\Pi(\cdot)$	Psudorandom permutation (PRP)
$\mathsf{SKE} = \{\mathsf{Gen}, \mathsf{Enc}, \mathsf{Dec}\}$	Symmetric encryption
k_f	PRF key
k_l	PRF key
$k_{\pi_{1,2,3}}$	PRP keys

Table 1: Notations

Fig. 4: The encrypted FM index construction.

5.1 First Approach

We give an overview of our first approach. The protocol can be described in two phases: The encrypted index phase and the search phase.

Encrypted Index To facilitate the reader we split the encrypted index process in two steps (cf. figure 4) the linked list, which *bootstraps* the search procedure on the FM index by the cloud and 2) FM index itself. The notation used for the protocol is given in table 1. For the security of the scheme the user employs lightweight cryptographic primitives: a pseudorandom function $F(\cdot)$, a pseudorandom permutation $\Pi(\cdot)$ respectively and a symmetric encryption scheme $SKE = \{Gen, Enc, Dec\}$. The untrusted cloud, thanks to the LF mapping and the FM index computation does not need to store all the suffixes of a stream S (cf. figure 4).

Linked List The crux of the design is on how to allow fast indexing through a hash table, which means that there should be unique keys derived from the string with repetitive characters. We employ the ranking information r_c of each character along with the character itself. However, when a user is looking for a substring, it does not know the ranking of each character in the substring $T[1...m]$. We mitigate this deficiency by building a linked list LL_c for each character.

The user computes a hash table of linked lists $LLSet$, where each position $LLSet[F_{k_f}(c^w)], 1 \leq w \leq \sigma$ maps to the linked list LL_{c^w}. The number of linked lists equals the number of distinct elements c, denoted as σ in the data stream S, where each symbol $c^w, 1 \leq w \leq \sigma$ comes from an alphabet Σ. The hash table is used to fetch all the positions of a character in the stream S from the linked lists LL_{c^w}. Each linked list LL_{c^w} stores information concerning the retrieval of the position of c from S. More specifically each node in the list stores the following tuple: $\langle nptr, addr \rangle$, $nptr$ is a pointer to the next node of the current list and $addr$ is the address of the element c in the FM index.

The first node of each linked list is stored in the $LLSet$ hash table. In order to prevent the adversary who tries to correlate elements from the $LLSet$ with positions in the FM index, we further encrypt each key $F_{k_f}(c^w)$ in the $LLSet$ hash map with another key k_l as follows: $F_{k_f}(c^w) \oplus F_{k_l}(c^w)$. Thus the cloud cannot correlate associations from $LLSet$ to FM attack offline without observing any token. The key of the hash map $LLSet$ at $F_{k_f}(c^w) \oplus F_{k_l}(c^w)$ maps to the first element of the linked list LL_c, which is encrypted as $\langle nptr, addr \rangle \oplus F_{k_l}(c^w)$[1]. As such, the frequency of each character before a search query is hidden.

However, once the cloud receives queries, it can learn the frequency of encrypted characters in the linked list which represent characters of the string. In order to obfuscate frequency analysis on the

---

[1] Notice that even if we use a one time pad with the same key for two different elements: $a = \langle nptr, addr \rangle, b = F_{k_f}(c^w)$ an adversary by xoring the two ciphertexts encrypted under the same key $F_{k_l}(c^w)$, learns $ab = \langle nptr, addr \rangle \oplus F_{k_f}(c^w)$, which is a one time pad encryption of $\langle nptr, addr \rangle$ with key $F_{k_f}(c^w)$.

encrypted index from substring search queries, we pad the data stream with dummy blocks. These dummy blocks make all linked lists to appear with the same size. The core idea for padding is to produce dummy blocks from the vocabulary Σ depending on to the ranking of the most frequent character. E.g: Original stream=abbcd and Σ=abcd, then the dummy blocks equal dc={ a, c, d}. Finally the user chooses uniformly at random $dpos \xleftarrow{\$} \{|dc|\}$ and appends the original string S at position $dpos$ of dc. Following the previous example; if dpos=1 then S'=aacdbbcd. The cloud responds in the final round with the encrypted position pos. User accepts the result as correct if $dpos \leq pos \leq n - dpos$ and $pos + m < dpos + n$.

FM Index Encryption (figure 4). The second difficulty comes when the cloud tries to traverse the FM index through the LF mapping technique. The encrypted FM index contains unique digests of characters, while the cloud should identify matches from the token $\mathsf{tk}_{T,S}$, that encodes repetitive characters deterministically. In order to allow the cloud traverse the encrypted FM index, we encrypt the FM as a key value hash table where the key consists of $\mathsf{F}_{\mathsf{k_f}}(c_{\mathsf{F}_j}) \oplus \mathsf{F}_{\mathsf{k_f}}(r_{\mathsf{F}_j}||c_{\mathsf{F}_j})$ and the value is $\mathsf{F}_{\mathsf{k_f}}(c_{\mathsf{L}_j}) \oplus \mathsf{F}_{\mathsf{k_f}}(r_{\mathsf{F}_j}||c_{\mathsf{F}_j})||\mathsf{F}_{\mathsf{k_f}}(r_{\mathsf{L}c_j}||c_{\mathsf{L}_j}), \mathsf{E}(\mathsf{pos}_j)$. Finally the user permutes all the tuples with a secure permutation: $\Pi_{\mathsf{k}_\pi}(t_j)$.

Search During the Search phase on a substring query $\mathsf{tk}_{T,S} = \mathsf{F}_{\mathsf{k_f}}(T[1...m]) = \mathsf{F}_{\mathsf{k_f}}(T[1]), \mathsf{F}_{\mathsf{k_f}}(T[2]), \ldots \mathsf{F}_{\mathsf{k_f}}(T[m]), \mathsf{F}_{\mathsf{k_l}}(T[m])$ the cloud proceeds as follows:

Bootstrap. First it needs to *bootstrap* the search by finding the correct candidate positions in the encrypted FM table through the linked list, which correspond to all positions in the string S where the last character of the query possibly exists. From the LLSet hash table it looks for the value with key $\mathsf{F}_{\mathsf{k_f}}(T[m]) \oplus \mathsf{F}_{\mathsf{k_l}}(T[m])$. This value maps to a linked list LL_c, in which each node maps to the encrypted FM tuple $t_j = t_j^0, t_j^1, t_j^2 =$

$$\langle \underbrace{\mathsf{F}_{\mathsf{k_f}}(c_{\mathsf{F}_j}) \oplus \mathsf{F}_{\mathsf{k_f}}(r_{\mathsf{F}_j}||c_{\mathsf{F}_j})}_{\mathsf{F}},$$

$$\underbrace{\mathsf{F}_{\mathsf{k_f}}(c_{\mathsf{L}_j}) \oplus \mathsf{F}_{\mathsf{k_f}}(r_{\mathsf{F}_j}||c_{\mathsf{F}_j})||\mathsf{F}_{\mathsf{k_f}}(r_{\mathsf{L}c_j}||c_{\mathsf{L}_j})}_{\mathsf{L}}, \underbrace{\mathsf{E}(\mathsf{pos}_j)}_{\mathsf{SA}}\rangle_{j=1}^n.$$

In order to decrypt the first element of the linked list the cloud uses $\mathsf{F}_{\mathsf{k_l}}(T[m])$ as a key to decrypt $\langle \mathtt{nptr}, \mathtt{addr}\rangle \oplus \mathsf{F}_{\mathsf{k_l}}(c^w)$ and learns $\langle \mathtt{nptr}, \mathtt{addr}\rangle$. The cloud uses $\mathsf{F}_{\mathsf{k_f}}(T[m])$ and applies a xor operation on the F column at the ranges that it retrieved from the linked list of the c_m character LL_c and learns $\mathsf{F}_{\mathsf{k_f}}(r_{\mathsf{F}_j}||c_{\mathsf{F}_j})$.

Iteration. The cloud uses $\mathsf{F}_{\mathsf{k_f}}(r_{\mathsf{F}_j}||c_{\mathsf{F}_j})$ as a key to decrypt the first part of the L column element $\mathsf{F}_{\mathsf{k_f}}(c_{\mathsf{L}_j}) \oplus \mathsf{F}_{\mathsf{k_f}}(r_{\mathsf{F}_j}||c_{\mathsf{F}_j})$ and reveals $\mathsf{F}_{\mathsf{k_f}}(c_{\mathsf{L}_j})$. It then fetches the encrypted L column as $k = \mathsf{F}_{\mathsf{k_f}}(c_{\mathsf{L}_j}), b = \mathsf{F}_{\mathsf{k_f}}(r_{\mathsf{L}_j}||c_{\mathsf{L}_j})$ in which $\mathsf{F}_{\mathsf{k_f}}(c_{\mathsf{L}_j}) = \mathsf{F}_{\mathsf{k_f}}(T[m-1])$ and for all nodes from the linked list computes $k \oplus b$, which is used as a key for the F column. The procedure terminates when the processed substring character is the first one $\mathsf{F}_{\mathsf{k_f}}(T[1])$.

At this point the cloud returns to the user all the encrypted $\mathsf{E}(\mathsf{pos}_j)$ for the substrings. The user decrypts and accepts the result as long as the decrypted position is in the range of the size of original stream without padding.

Second round. From the per-character one way function $\mathsf{F}_{\mathsf{k_f}}$ evaluation of the substring query: $\mathsf{tk}_{T,S} = C[1], C[2], \cdots, C[m] \leftarrow \mathsf{F}_{\mathsf{k_f}}(T[1...m])$ and the LF mapping the protocol leaks to the cloud in cleartext the exact differences of the positions of two encrypted substring in the stream S as long there are unique matches. More specifically, the number of iterations in the LF mapping traversal (algorithm 2), reveals how many positions two substrings they differ, as long as there is *unique* match in S. Eventually, an untrusted cloud can decrypt the entire encrypted SA array, which contains encrypted positions of all substrings in S, since it knows its addresses. To circumvent the leakage we first use two different permutations to permute the tuples t_j: $\Pi_{\mathsf{k}_{\pi_1}}(t_j^0, t_j^1) = \pi_j^{0,1}, \Pi_{\mathsf{k}_{\pi_2}}(t_j^2) = \pi_j^2$. As such, after the permutation $\underbrace{\mathsf{F}_{\mathsf{k_f}}(c_{\mathsf{F}_j}) \oplus \mathsf{F}_{\mathsf{k_f}}(r_{\mathsf{F}_j}||c_{\mathsf{F}_j})}_{\mathsf{F}}, \underbrace{\mathsf{F}_{\mathsf{k_f}}(c_{\mathsf{L}_j}) \oplus \mathsf{F}_{\mathsf{k_f}}(r_{\mathsf{F}_j}||c_{\mathsf{F}_j})||\mathsf{F}_{\mathsf{k_f}}(r_{\mathsf{L}_j}||c_{\mathsf{L}_j})}_{\mathsf{L}}$ are stored in position $\pi_j^{0,1}$ at the FM array and $\underbrace{\mathsf{E}(\mathsf{pos}_j)}_{\mathsf{SA}}$ at position π_j^2. The cloud as traverses the token returns the permuted encrypted position of the substring token, the client applies the inverse permutation and fetches the correct cell from the FM array. By doing so the cloud cannot learn on its own, the encrypted position of a substring. The

second permutation prevents him to learn this information by stopping the traversal of the index at any substring of the original query at its choice. To perform that it needs the contribution of the user. In part the second permutation can be viewed as the *induced permutations* of *structured encryption* [8], which encrypt positions of items belonging a data structure, but in contrast with *structured encryption* a user in substring queries is interested to identify positions of sub-elements (substrings) of the original data sructure and not the entire elements (e.g: positions of elements in a matrix).

5.2 Improved scheme

Shortcoming Recall that during the encryption of the index, the user adds dummy blocks at the linked lists in order to alleviate frequency attacks. Namely after issuing a query token $\mathsf{tk}_{T,S}$ the client reveals to the cloud the key $\mathsf{F}_{\mathsf{k}_\mathsf{l}}(T[m])$ in order to locate the head of the list which corresponds to the key $\mathsf{F}_{\mathsf{k}_\mathsf{f}}(T[m]) \oplus \mathsf{F}_{\mathsf{k}_\mathsf{l}}(T[m])$. The cardinality of the list corresponds to the frequency of that character in the original string S. In our first approach we address this problem by adding dummy blocks in all linked lists in order their size to equal the size of the longest one.

The aforementioned technique protects the client from frequency attacks on the original string and imposes low overhead in case of a string drawn from a distribution with homogeneous frequency characters. However in a more *skewed* dataset with characters having broad frequencies, then the technique of dummy characters can drastically affect the efficiency of the system. More specifically the dummy characters may double the size of the final size of the index, thus degrading the storage overhead and the computational efficiency of the client.

Bucketization of characters In order to tackle the problem of increased storage due to the need of a large number of dummy characters to make the list of equally size, we flatten the frequencies as follows: we group together the characters into buckets of equal size and we build index on top of the buckets. By doing so the frequencies of buckets composed of consecutive characters of the string are close together and thus few dummy buckets are needed. User starts by organizing the string in buckets: It selects the size win of a window that traverses the string S. At each iteration z it extracts win consecutive characters, puts them in the bucket b and shifts the window one character forward at the left. Finally a set of buckets $\{\mathsf{b}_z\}_{z=1}^{n-\mathsf{win}+1}$ is extracted.

Index. The entire procedure is similar to the first approach but instead of operating on single characters everything operates on buckets. Now the F and L column arrays correspond to buckets and the input to the $\mathsf{F}_{\mathsf{k}_\mathsf{f}}$ is not a single character c but a bucket b. We use the same notation conventions with the first approach but instead of operating on characters we operate on buckets, i.e: we denote by b^1 the first bucket of string, $\mathsf{b}_{\mathsf{F}_j}$ is the j^{th} bucket of the F column, etc. The LLSet hashmap stores for each key $\mathsf{F}_{\mathsf{k}_\mathsf{f}}(\mathsf{b}^z) \oplus \mathsf{F}_{\mathsf{k}_\mathsf{l}}(\mathsf{b}^z)$, the head of lists, which correspond to buckets of the stream. The entire procedure to encrypt the index is identical with our basic approach and we omit a repetitive overview thereof.

Search. During the search phase if the size of the token, which consists of possible consecutive buckets of the original string S, is a multiple of the windows size win then the protocol protocol proceeds identically as with the *per-character* previous version of the search procedure.

However when the size of the token is not multiple length of win then there will be always a faulty mismatch. The possible match of the last bucket of the token will be inside the last non-matched bucket from the index. Recall that from the LF traversal on the FM index, the search starts from the last character-bucket and proceeds up to the first character-bucket of the search token. As such the cloud during the search phase it will always misfire a mismatch. We alleviate the correctness problem as follows. The cloud starts the search not from the last bucket of the token $\mathsf{F}_{\mathsf{k}_\mathsf{f}}(T[m])$ but for the previous one $\mathsf{F}_{\mathsf{k}_\mathsf{f}}(T[m-1])$. If there is a match for all $m-1$ blocks of the token then the client needs to decrypt the mth bucket in the string to verify matching of the first $e = n \bmod \mathsf{win}$ remaining characters which correspond to the last bucket of the query token. To do so the client has to encrypt and upload the stream S similarly with [9] in a per bucket fashion resulting to an array of encrypted buckets B. To avoid the leakage of the position the client permutes the buckets with a pseudorandom permutation Π, keyed by k_{π_3} : $B'[j] = B[\Pi_{\mathsf{k}_{\pi_3}}(j)]$ and instead of querying for the mth bucket it forwards a request for bucket number $B[\Pi_{\mathsf{k}_{\pi_3}}(m)]$. It then decrypts the bucket and compares it with the last bucket from the query to identify a matching happening at position m.

6 Protocol

6.1 S^3E Description

We are now ready to give the full details of our substring searchable symmetric encryption protocol, which alleviates the storage overhead shortcoming of the first approach with our bucketization technique:

– $k \leftarrow$ KeyGen(1^λ): This algorithm runs by the user takes as input the security parameter 1^λ and generates random keys $k = (k_f, k_l, k_{\pi_{1,2,3}}, k_{e,c})$ for a PRF $F_{k_f} : \{0,1\}^\lambda \times \{0,1\}^\nu \rightarrow \{0,1\}^\mu$, a PRP $\Pi_{k_\pi} : \{0,1\}^\lambda \times \{0,1\}^\nu \rightarrow \{0,1\}^\nu$ and a symmetric encryption algorithm SKE = $\{$Gen, Enc, Dec$\}$. Finally it outputs k to the user. For the generation of the keys we assume a source of randomness $\mathcal{R}$ and a pseudorandom generator G seeded with

$$s_f \xleftarrow{\$} \mathcal{R}, s_l \xleftarrow{\$} \mathcal{R}, s_{\pi_{1,2,3}} \xleftarrow{\$} \mathcal{R}, s_e \xleftarrow{\$} \mathcal{R} :$$
$$(k_f, k_l, k_{\pi_{1,2,3}}, k_{e,c}) \leftarrow G(s_f), G(s_l), G(s_{\pi_1}), G(s_{\pi_2}), G(s_e)$$

– SES $\leftarrow$ PreProcess(k, S): User owns a stream S, which contains characters $c \in \Sigma$. S has n characters. Let $\max_b$ be the cardinality of most frequent bucket and f_{b_i} the frequency of bucket b_i. User:
   1. Parses the string S as buckets: $\{b_z\}_{z=1}^{n-\text{win}+1}$, each of size win characters and k in total distinct buckets:
      for $z = 1; z++; z = n - \text{win}$
         $b_z = (z + \text{win} \leq |S|)?S[z...z + \text{win}] : S[z...n - z]$
   2. Chooses dummy buckets $dc = \sum_{j=1}^k \max_b - f_{b_i}$ that constitute a dummy stream. The user chooses uniformly at random $dpos \xleftarrow{\$} \{|dc|\}$ and appends the original string S at position $dpos$ of dc
   3. Computes the suffix array SA and the F, L columns on input the buckets B and stores them as the FM index: FM = F$||$L, SA.
   4. Encrypts all buckest $B = $SKE.E$(k_c, b_z), 1 \leq z \leq n + k - \text{win} + 1$
   5. Permutes $B'[j] = B[\Pi_{k_{\pi_3}}(j)]$
   6. Encrypts elements of SA array with SKE.E$(k_e, SA[i]), 1 \leq i \leq +k - \text{win} + 1$.
   7. Applies the PRF to each element of F as follows:

   $$F[i] = F_{k_f}(b_{F_i}) \oplus F_{k_f}(r_{F_i}||b_{F_i})$$

   8. Computes

   $$L[i] = F_{k_f}(b_{L_i}) \oplus F_{k_f}(r_{F_i}||b_{F_i})||F_{k_f}(r_{L_i}||b_{L_i})$$

   9. Applies a pseudorandom permutation Π_{k_π} to the tuples:

   $$t_0 = F_{k_f}(b_{F_i}) \oplus F_{k_f}(r_{F_i}||b_{F_i}), t_1 = SA[i]),$$
   $$t_1 = F_{k_f}(b_{L_i}) \oplus F_{k_f}(r_{F_i}||b_{F_i})||F_{k_f}(r_{L_i}||b_{L_i})$$

   using k_{π_1} and with k_{π_2} user permutes E$(k_e, SA[i])$, for $i = 1, ..., n + k - \text{win} + 1 : \Pi_{k_{\pi_1}}(t_i^0, t_i^1) = \pi_i^{0,1}, \Pi_{k_{\pi_2}}(t_i^2) = \pi_i^2 = FM'$.
   10. For every distinct bucket in $F[i] = F_{k_f}(b_{F_i}) \oplus F_{k_f}(r_{F_i}||b_{F_i})$ the user initiates a linked list LL_c and at each node stores $LL_c.$nptr for the next node of the list and $LL_c.$addr which points to the tuple t_i with a matching $F_{k_f}(b_{F_i})$. Finally it encrypts the first element of each linked list LL_c with $F_{k_l}(b_i) : \langle$nptr, addr$\rangle \oplus F_{k_l}(b_i)$.
   11. Stores the head pointers of the collections of all linked lists in a hash table LLSet with key $k = F_{k_f}(b_i) \oplus F_{k_l}(b_i)$ and value v a pointer to the head of the list LL_c, which stores information about the $F_{k_f}(b_i)$ character, meaning all its positions to the encrypted FM index.
   12. Finally outputs SES $= ($LLSet$, LL_c, FM', B')$ and keeps only the keys $k = (k_f, k_l, k_{\pi_{1,2,3}}, k_{e,c})$.

– $tk_{T,S} \leftarrow$ SrchToken($k, T[1...m]$): This algorithm takes as input the secret substring search key k, a string $T[1...m]$ and outputs a trapdoor to search for the string T on data stream S, through SES:
   1. Parse $T[1...m]$ to buckets $T_b = \{T_b^i\}_{i=1}^{m-\text{win}}$ of size win:
      for $i = 1; i++; i = m - \text{win}$
         $T_b^i = (i + \text{win} \leq |S|)?T[i...i + \text{win}] : T[i...m - i]$

2. User with his secret PRF key k_f computes $tk_{T,S} = C[1], C[2], \cdots, C[m - win] \leftarrow F_{k_f}(T_b = \{T_b^i\}_{i=1}^{m-win-1}), F_{k_l}(T_b^{m-win})$ and forwards $tk_{T,S}$ to the cloud.

- $(s, \perp) \leftarrow$ Search$(tk_{T,S}, \text{SES})$: The cloud parses the token query $tk_{T,S} = C[1], C[2], \cdots, C[m - win]$ and searches the position in S from the encrypted index SES as follows:

1. **if** $|tk_{T,S}|$ mod win $== 0$
   $x = m - \text{win}, \text{flag} = 0$
   else
   $x = m - \text{win} - 1, \text{flag} = 1$ // `Search for the last equal size bucket pattern before the last`
   `one.`

2. $u = find(\text{LLSet}, C[x])$//`find in dictionary LLSet the value` u `with key` $C[x]$`.u is a pointer`
   `to the head of a list LL, which stores pointers to all buckets` $T[m]$ `in` S

3. **if** $u == \perp$ **return** $\perp$

4. **while** $u \neq \perp$ **do**
   $K = K \cup u.\text{addr}$ //`traverse the list and store in the set` K `the addresses of the`
   `characters.`
   $u = u.\text{nptr}$

5. **for** $p = x - 1; p > 1; p = p - 2$
   for $i = 1; i < size(K); i++$
   if $\text{SES.L}[K[i]]^{(1)} == C[p - 1]$//`Store in the set` $KEYS$ `only the elements from the F`
   `column, whose associated L element equals the next bucket from` C `in a backword order.`
   $\text{SES.L}[K[i]]^{(1)}$ `maps to` $F_{k_f}(b_{L_i}) \oplus F_{k_f}(r_{F_i} || b_{F_i})$ `and` $\text{SES.L}[K[i]]^{(2)}$ `to` $F_{k_f}(r_{L_i})$`.`
   $KEYS = KEYS \cup \text{SES.L}[K[i]]$
   else $K = K - K[i]$ //`Remove all the non matched elements from the key set` K`.`
   if $K == \perp$ **return** $\perp$
   for $i = 1; i < size(KEYS); i++$
   $r_{F_i} = C[p] \oplus KEYS^{(1)}[i]$
   $z = r_{F_i} \oplus KEYS^{(2)}[i]$ //`Compute the key from the L column as` $KEYS^{(1)}[i] \oplus KEYS^{(2)}[i]$`,`
   `which corresponds to` $F_{k_f}(b_{F_i}) \oplus F_{k_f}(r_{F_i} || b_{F_i})$ `in the F column of the SES object.`
   if $\text{SES.F}[z] \neq \perp$
   continue
   else $KEYS = KEYS - KEYS[i]$ //`Remove all the non matched elements from the key`
   `set` $KEYS$`.`
   $K = KEYS$

6. **if** $K == \perp$ **return** $\perp$

7. The cloud sends to the user $\text{SES.FM}'[K]$. The client runs the inverse permutation Π_{π_2} to the K indexes and gets back $\{i'\}$ and asks the cloud for $\text{SES.FM}'[\{i'\}]$. After getting back the results the client decrypts $\text{pos} = \text{E}(k_e, \text{SA}[\{i'\}])$ with k_e and learns the position pos of the asked substring T in S. In case $\text{flag} == 0$ user accepts the result as correct if $dpos \leq pos \leq n - dpos$ and $\text{pos} + m < dpos + n$.

8. **if** $\text{flag} == 1$ recap that the LF traversal on the FM index starts at the last character and proceeds invertly. Notice that we excluded at step **1** during the **else** branch the remaining last bucket as it will always be a mismatch even if the first characters match the user search pattern. So the cloud returns also B'[pos]. The client runs the permutation $\text{pos}' = \Pi_{k_{\pi_3}}(\text{pos})$ and asks the cloud to return the bucket $b_{pos'}$. Then client checks if the first m mod win characters of the decrypred bucket equals the last m mod win characters of the search pattern and accepts the pos as valid otherwise discards the result.

7 Security Analysis

We illustrate the security of the scheme pertaining to definition 2. More specifically, we show the existence of a simulator $\mathcal{S}$, who has access to the leakage function $\mathcal{L}$ and produces indistinguishable views to an adversary $\mathcal{A}$. Conceptually the proof demonstrates that an adversary $\mathcal{A}$, who can be a semi-honest cloud cannot learn more information from what it can be leaked in an ideal work without malicious behaviors.

User splits the original data stream S in consecutive buckets of size win by moving each time the window to the right one character. During the encryption of the index users write to the LLSet linked

list and to a three dimensional array FM. FM consists of three columns: F, L from the BWT transformation, which are used to traverse the string, and the SA column, which consists of encrypted positions of the corresponding suffixes. To avoid leakage of frequencies users pad each linked list LLSet. In order to impede off-line traversal of the index, the user further encrypts the header of each linked list. Moreover, two different permutations $\Pi_{k_{\pi_{1,2}}}$ are applied first on all the rows of the first two columns F, L and then to the SA column to prevent an adversary from correlating encrypted positions to intermediate results after having obtained a token. I.e: without the permutation, $\mathcal{A}$ after obtaining a token can stop the search to a substring of the queried substring and learn the associated encrypted position from SA column. After $\mathcal{A}$ observing LLSet and FM it learns the size of the padded string n'.

During the the search phase the adversary $\mathcal{A}$ observes encrypted tokens. Between two encrypted tokens $\mathcal{A}$ can identify similarity in between two tokens: i.e: whether two identical token have been issued, or identical encrypted buckets in between the tokens. Then we point out two different cases:

1. **Size of token multiple of** win: The user leaks to the cloud how many matches exist in the string S for the queried token, the encrypted cells at SA, which correspond to the matched positions.
2. **Size of token not a multiple of** win: The user in case of a token, whose size is not a multiple of the size win of each bucket, needs to verify the correctness of the potential encrypted match(es). As such the cloud returns the permuted possible position pos. User decrypts and applies $\mathsf{pos}' = \Pi_{k_{\pi_3}}(\mathsf{pos})$ and asks the cloud to return the bucket at position pos'. The cloud learns the encrypted position of the last bucket of the token and its encrypted text.

Integrity guarantees of the data and the index are assured thanks to the use of authenticated symmetric cryptography. Before stating our theorem concerning the security of $\mathsf{S^3E}$ we formally define our leakage function $\mathcal{L}$ from our aforementioned analysis.

Definition 3. *A leakage function $\mathcal{L}$ for a $\mathsf{S^3E}$ scheme comprises the following three leakage functions:*

- PreProcess *Leakage: $\mathcal{L}_1$ includes the padded size of the data stream $n' >= |S|$.*
- SrchToken *Leakage: The SrchToken Leakage $\mathcal{L}_2$ reveals the length of the token $|\mathsf{tk}|$, how many common characters reside in it and similarity patterns between different tokens.*
- Search *Leakage: $\mathcal{L}_3$ leaks how many times a substring token tk exists in the padded string S with dummy blocks.*
- Intermediate*: Leakage $\mathcal{L}_4$ leaks intermediate addresses of the F and L columns of a query.*
- Access *pattern: Leakage $\mathcal{L}_5$ leaks the addresses of the SA cells for a fixed query.*

To eliminate $\mathcal{L}_2, \ldots, \mathcal{L}_5$ leakage profiles an expensive ORAM scheme can be used, however it is out of the scope of the paper, as it will increase communication and storage overhead drastically.

Theorem 1. *Let $\mathsf{F}_{k_f}, \Pi_{k_\pi}, \mathsf{SKE} = \{\mathsf{Gen}, \mathsf{Enc}, \mathsf{Dec}\}$ be a pseudorandom function, a pseudorandom permutation and a semantically secure symmetric encryption scheme respectively, then our substring searchable symmetric encryption scheme $\mathsf{S^3E}$ is adaptively $\mathcal{L}$-semantically secure.*

Proof. (Sketch)

In the $\mathbf{Real}^{\mathsf{S^3E}}_{\mathcal{A}(\lambda)}$ world (cf. 2.2) the adversary can obtain the encrypted index and encrypted keywords of its choice. In the beginning the Challenger selects the size of the buckets win and uniformly at random keys $\mathsf{k} = (\mathsf{k}_f, \mathsf{k}_l, \mathsf{k}_{\pi_{1,2,3}}, \mathsf{k}_{e,c})$ for a PRF $\mathsf{F}_k : \{0,1\}^\lambda \times \{0,1\}^\nu \to \{0,1\}^\mu$, a PRP $\Pi_{k_\pi} : \{0,1\}^\lambda \times \{0,1\}^\nu \to \{0,1\}^\nu$ and a symmetric encryption algorithm $\mathsf{SKE} = \{\mathsf{Gen}, \mathsf{Enc}, \mathsf{Dec}\}$. Upon receipt of a stream S of size n, the Challenger employs the $\mathsf{SES} \leftarrow \mathsf{PreProcess}(\mathsf{k}, S)$ as presented in section 6 and forwards SES to $\mathcal{A}$. We distinguish between matching q^m and non-matching queries $q^{nm} : \mathbf{q} = \bigcup q^{nm} q^m$. We assume for the ease of readability that adversary issues only matching queries q^m. During the Search phase, $\mathcal{C}$ has also access to $\mathcal{L}_2$=similarities between tokens, $\mathcal{L}_3$=# times a token exist in the substring. $\mathcal{C}$ upon receipt of $\mathbf{q}$ checks in a table QT whether $\mathbf{q} \in QT$. If so $\mathcal{C}$ fetches the corresponding token $\mathsf{tk}_{T,S}$ and forwards it to $\mathcal{A}$. If this is the frist time for $\mathbf{q}$ then $\mathcal{C}$ computes $\mathsf{tk}_{T,S} = C[1], C[2], \cdots, C[m] \leftarrow \mathsf{F}_{k_f}(T[1...m]), \mathsf{F}_r(T[m])$. Finally $\mathcal{A}$ receives $\mathbf{t} = (\mathsf{tk}_1, \mathsf{tk}_2, \mathsf{tk}_3, \ldots, \mathsf{tk}_o)$ for each substring query.

Within a sequence of hybrid games we show the indistinguishable transformation of $\mathbf{Real}^{\mathsf{S^3E}}_{\mathcal{A}(\lambda)}$ game to eventually the $\mathbf{Ideal}^{\mathsf{S^3E}}_{\mathcal{A},\mathcal{S}(\lambda)}$ game, which concludes the proof. The simulator $\mathcal{S}$ computes the simulated encrypted index $\mathsf{SES}^* = (\mathsf{LLSet}^*, \mathsf{LL}_c^*, \mathsf{FM}'^*)$ as follows:

- $\mathsf{Game_0}$: This game is equivalent with the $\mathbf{Real}_{\mathcal{A}(\lambda)}^{\mathsf{S^3E}}$ game.
- $\mathsf{Game_1}$: This game behaves as the $\mathbf{Real}_{\mathcal{A}(\lambda)}^{\mathsf{S^3E}}$ game with the difference that $\mathcal{S}$ does not have access to S. The simulator through the $\mathcal{L}_1$ leakage function builds the substring encrypted structure SES as follows: We assume the existence of an algorithm $S \leftarrow \mathsf{Build}(n', str)$, which takes as input $n' \in \mathbb{N}$ and the structure $str = \{\mathsf{b}\}_{i=1}^{n'}, \mathsf{b} \in \Sigma^*$ and outputs a bitstring of length n', from a vocabulary Σ^*. Notice that as in the real game the valid length of the original stream is not revealed and only the length of the string after the padding n' is leaked. $\mathcal{S}$ selects uniformly at random keys $\mathsf{k} = (\mathsf{k_f}, r, \mathsf{k_\pi}, \mathsf{k_e})$ for a PRF $\mathsf{F_{k_f}} : \{0,1\}^\lambda \times \{0,1\}^\nu \rightarrow \{0,1\}^\mu$, a PRP $\Pi_{\mathsf{k_\pi}} : \{0,1\}^\lambda \times \{0,1\}^\nu \rightarrow \{0,1\}^\nu$ and a symmetric encryption algorithm $\mathsf{SKE} = \{\mathsf{Gen}, \mathsf{Enc}, \mathsf{Dec}\}$ and runs $\mathsf{SES} \leftarrow \mathsf{PreProcess}(\mathsf{k}, \mathsf{Build}(\mathcal{L}_1))$. $\mathcal{S}$ uses $\mathsf{F_{k_f}}$ to evaluate bit strings of length c_n: $\mathcal{L}_2(q) = c_n$. For the simulation of the tokens and its responses, $\mathcal{S}$ uses the leakage obtained from $\mathcal{L}_2 \ldots \mathcal{L}_5$.
- $\mathsf{Game_2}$: This game behaves similarly with $\mathsf{Game_1}$, but we replace the $\mathsf{F_{k_f}}$ with a real random function which is evaluated through access to an oracle $\mathcal{O}^{\mathsf{RF}}(\lambda, \mu, \nu)$.
- $\mathsf{Game_3}$: This game behaves similarly with $\mathsf{Game_2}$, but we replace the $\Pi_{\mathsf{k_\pi}}$ with a real random permutation which is evaluated through access to an oracle $\mathcal{O}^{\mathsf{R\Pi}}(\lambda, \nu)$.
- $\mathsf{Game_4}$: In $\mathsf{Game_4}$ we replace the semantically secure $\mathsf{SKE} = \{\mathsf{Gen}, \mathsf{Enc}, \mathsf{Dec}\}$ with real random values by querying an oracle $\mathcal{O}^{\mathsf{RE}}(\lambda)$.

We write $\mathsf{Game_i} \approx \mathsf{Game_j}$ to denote that the view of probabilistic polynomial time adversary $\mathcal{A}$ is indistinguishable between the output of $\mathsf{Game_i}$ and $\mathsf{Game_j}$. $\mathsf{Game_0} = \mathbf{Real}_{\mathcal{A}(\lambda)}^{\mathsf{S^3E}}$ by definition, $\mathsf{Game_1} \approx \mathsf{Game_0}$ as long as no collisions happen to the evaluation of $\mathsf{F_{k_f}}$, $\Pi_{\mathsf{k_\pi}}$, $\mathsf{SKE} = \{\mathsf{Gen}, \mathsf{Enc}, \mathsf{Dec}\}$ or E, $\mathsf{Game_2} \approx \mathsf{Game_1}$ as long as $\mathsf{F_{k_f}}$ is indistinguishable from real random function, $\mathsf{Game_3} \approx \mathsf{Game_2}$ thanks to the indistinguishable output of $\Pi_{\mathsf{k_\pi}}$ from real random permutations, $\mathsf{Game_4} \approx \mathsf{Game_3}$ because of the semantically secure $\mathsf{SKE} = \{\mathsf{Gen}, \mathsf{Enc}, \mathsf{Dec}\}$ and finally $\mathsf{Game_5} = \mathbf{Ideal}_{\mathcal{A}, \mathcal{S}(\lambda)}^{\mathsf{S^3E}}$ by definition.

Game	Change	Indistinguishability Argument
$\mathsf{Game_0}$	$\mathsf{Game_0} = \mathbf{Real}_{\mathcal{A}(\lambda)}^{\mathsf{S^3E}}$	By definition
$\mathsf{Game_1}$	Replace $\mathsf{F_{k_f}}, \mathsf{F_{k_l}}$	Pseudorandomness of $\mathsf{F_{k_f}}$
$\mathsf{Game_2}$	Replace $\Pi_{\mathsf{k_{\pi_{1,2,3}}}},$	Pseudorandomness of $\Pi_{\mathsf{k_\pi}}$
$\mathsf{Game_3}$	Replace $\mathsf{SKE} = \{\mathsf{Gen}, \mathsf{Enc}, \mathsf{Dec}\}$	Semantically secure $\mathsf{SKE} = \{\mathsf{Gen}, \mathsf{Enc}, \mathsf{Dec}\}$
$\mathsf{Game_4}$	$\mathsf{Game4} = \mathbf{Ideal}_{\mathcal{A}, \mathcal{S}(\lambda)}^{\mathsf{S^3E}}$	By definition

Table 2: Hybrid games

7.1 Comparison

Protocol	Search	Index [PS—CS]		Query [FR—LR]		VLS	Rounds	SL
CS[9]	$O(m+k)$	$20n$	$4(n + \Sigma_{i=1}^{2n-(2+N)} \sigma - child(i))$	$\frac{m(m+1)}{2\theta}$	m+k	✓	3	SP+QPP+IIP+LIP
FJKNRS[13]	$O(n)$	-		m	0	✓	1	SP
FHV[14]	$O(n-m)$	-		m	0	✗	1	✗
$\mathsf{S^3E}$	$O(m+k)$	$4n$	$4(n + \sum_{j=1}^k \mathsf{max_b} - f_{\mathsf{b}_i})$	m	1	✓	2, 3	SP+QPP+IIP

Table 3: Comparison of existing substring searchable encryption protocols. Index space is further categorized in plaintext space storage space (PS) and ciphertext space (CS). The overhead of [13] and [14] is undefined as the schemes do not take advantage of any auxiliary data structure for efficient substring search. For the query complexity we analyzed its size in terms of two separated phases: at the first round (FR) of the protocol and the last one (LR), in case of multiple rounds protocols. **VLS** denotes variable length substring search and **SP** the search patternleakage: QPP: Query prefix pattern, IIP: Index intersection pattern, LIP: Leaf intersection pattern.

We perform a comparison of our $\mathsf{S^3E}$ with existing solutions (cf. table 3). We analyzed the search running time in asymptotic complexity, index space requirements both in the plaintext and in the ciphertext space, query size, variable length capability, rounds of communication and search leakage. Since our scheme competes mostly with [9] here we further elaborate its cost analysis from table 3.

Search. Thanks to the usage of encrypted dictionary the cost of searching a m length string is $O(m+k)$, where k denotes the number of occurrences. However, due to the extra dummy blocks the search cost is

increased to $O(m+k+\sum_{j=1}^{k} (\mathsf{max_b} - f_{\mathsf{b}_i}))$, where $\mathsf{max_b}$ is the most frequent bucket and f_{b_i} the frequency of bucket b_i.

Index. For the index space complexity, we analyzed the space requirement in the plaintext space and in the ciphertext space. For the plaintext space analysis we assume, that a pointer or integer requires 4 bytes. Recall that a suffix tree has n leaves, at most $n-1$ internal nodes and at most $2n-2$ edges. Thus, for a naive suffix tree implementation we need 2 pointers for each leaf: one for the parent node and one for its position to the original stream, resulting in $8n$ bytes. Four pointers for each internal node: one for the parent node, one for each leftmost child, one for the right sibling and one pointer for the suffix link, which reduces the search time during a substring query. The total storage cost for the internal nodes is $4 * 4n = 16n$. For each edge, suffix trees allocate one pointer for the beginning position of the substring in the stream and one for the end position of the substring in the stream increasing the space cost to $24n + 4 * 2 * 2n = 40n$. The space cost of the solution based on suffix trees [9] can be further reduced to $20n$ by eliminating the need to store suffix links and parent pointers. However, the extra dummy blocks further augment the storage overhead. Assuming a suffix tree with N internal, each node is further padded with dummy children nodes so as to each node has σ children, where σ is the size of the vocabulary. Furthermore, the internal nodes are padded with up to $2n-2$ nodes where n is the size of the string. Finally the size of the extra dummy nodes in [9] is: $\Sigma_{i=1}^{2n-(2+N)}(\sigma - child(i))$, where $child(i)$ equals the number of children for internal node i in the suffix tree. In contrast, in $\mathsf{S^3E}$ we replace the space expensive suffix trees with suffix arrays and as such the index space cost is reduced from $20n$ bytes to $4n$ bytes.

For the storage space computation during the encryption of the index, be it suffix tree or suffix array, we exclude a per byte comparison and we assume a ciphertext comparison. The encryption of the index is based on the translation of the suffix tree to an encrypted dictionary. Thus, all the extra pointers of the suffix tree are excluded. Following the protocol from [9], the user encrypts $2n$ substrings which equal to the number of edges of the suffix tree, n leaves and n characters of the original stream, resulting in $4n$ encryptions. In our solution thanks to the FM mapping the user sends the encrypted suffix array, plus two more n size arrays for the FM index construction; one for the F column of the index and one for the L column. In the end it uploads $4(n + \sum_{j=1}^{k} \mathsf{max_b} - f_{\mathsf{b}_i})$ encrypted values to the cloud, in total.

Query. For the query size, we assume a block cipher of size θ and a substring query of size m. In [9] the substring is encrypted incrementally: for the substring "abc" user encrypts separately E(a), E(ab), E(abc). As such, for big substring queries as in DNA queries, the number of ciphertexts exceeds the number of the substring m. The total number of encryptions equals $\frac{1}{\theta} + \frac{2}{\theta} + \cdots + \frac{m}{\theta} = \frac{m(m+1)}{2\theta}$ during the first round. At the last round the user asks for the positions of each character separately augmented by a factor of m the substring size. In $\mathsf{S^3E}$ the substring query has only per character encryptions of each character in the first round plus a ciphertext for the last round. Our solutions also allows variable size substring queries, since the size of the substring query is decoupled from the scheme and can be defined online during the query phase as in [9].

Rounds. For the rounds of communications, $\mathsf{S^3E}$ can return the substring search results in 2 rounds of communication or in 3 in case of a query size not a multiple of the bucket size win. During the first round, the client sends an encrypted substring query and the cloud responds with encrypted addresses of the corresponding suffix array positions. At the second round the client decrypts the permuted position of the suffix tree and queries the cloud for the unpermuted encrypted position. A third round is performed in case of a query not a multuple of the size of the bucket, whereby the client verifies the correctness of the possible match.

Security Leakage. Concerning the search leakage, Chase *et al.* [9] scheme leaks the search pattern, meaning an untrusted could can identify similarities between two or more substring search queries. Moreover, the scheme reveals the query prefix pattern, which leaks whether a node has been visited for a previous substring in the suffix tree, the index intersection pattern which allows the cloud to learn if the returned index position has already been asked and finally, the leaf index leaks when any of the returned positions of the tree leaves have been previously queried. Since in $\mathsf{S^3E}$ we avoid the use of a suffix tree, $\mathsf{S^3E}$ does not leak the leaf pattern. We inherit though, the index intersection pattern, which reveals if any returned index has been returned in a previous query. As in [9] our scheme reveals also the cloud differences of the indexes when a user asks for substrings that they do differ in one position and there is only a single position in the original stream S. Both schemes employ a padding policy to add dummy blocks in order to obfuscate the structure of the index and the stream. $\mathsf{S^3E}$ is also adaptively secure

16

under the real-ideal simulation paradigm. We also use an authenticated encryption scheme in order to assure the integrity of the messages.

8 Performance

In this section we present our implementation results. We demonstrate the practicality of $\mathsf{S^3E}$ with benchmark experiments, comparing our results with the scheme of Chase *et al.* [9], in order to validate the claims of our performance improvements. To accomplish the comparison we also implemented the suffix tree based construction of [9] called hereafter ST.

8.1 Implementation

For the benchmark experiments we used a machine running Ubuntu 14.04 with kernel version 3.19.0-29. The machine has 8GB RAM memory and is equipped with an INTEL Intel(R) Core(TM) i5-5200U CPU @ 2.20GHz processor with 4 cores. We instantiated the PRF, PRP using `Poly1305 + Salsa20`, `BLAKE2b` [2], which outperform their competitors in computational efficiency. The most computational heavy operation on the client side is the computation of the index, which becomes the bottleneck of the total performance. In our benchmarks, since the code will vary in different machines, we choose to isolate the code in different modules in order to evaluate their relevant performance. Thus, the first module consists of operations that take place in cleartext data, the second module entails the cryptographic primitives used to encrypt the `FM` index and the third module consists also of all `I/O` operations in order to serialize the encrypted `FM` index. For our benchmarks we used two real datasets: one DNA sequence from `https://www.ncbi.nlm.nih.gov/` and randomly extracted emails from the enron email data set `https://www.cs.cmu.edu/~./enron/`. We varied in either cases the size of the corpus in order to observe the feasibility results in these variations.

Our comparison is based on two metrics: **a)storage overhead** for the encryption of the index and the computation of a substring query as an encrypted token, **b)computation time** of each operation. The reported computation times for each experiment are taken as the average of 100 trials. As we implemented both schemes on the same machine, which simulates both the client and the server, we can derive accurate and fair observations about the performance of the protocols on real metrics. In the real world the server can be implemented in a more powerful machine, however this does not change the storage overhead or the computation performance fraction of both schemes.

Thanks to our encrypted suffix array construction, we achieve a storage improvement by a factor of **1.71** for the DNA sequence data stream and **1.57** for the enron email data. This occurs first because of the extra information a node of the suffix tree should keep (leaf nodes, parent nodes, auxiliary information for the substring of the path in the tree) and due to the dummy nodes policy, which increases the size of the tree. Subsequently that affects the computation time for the computation of the encrypted index with a $\approx$ **4x** blowup on average for data sets of size 10^6.

8.2 Benchmarks

Index In tables 4,5, we depict the storage and computational overhead incurred by the computation of the encrypted index using the suffix array in $\mathsf{S^3E}$ using different window sizes for the buckets and the suffix tree in ST scheme of [9].

Dataset	#Characters									
	10^2		10^3		10^4		10^5		10^6	
	Storage	Time	Storage	Time	Storage	Time	Storage	Time	Storage	Time
DNA	60KB	0.44s	567KB	1.07s	6.4MB	9.66s	63MB	89.50s	589MB	1382s
Enron	57KB	0.42s	562KB	0.97s	6.3MB	9.07s	62MB	85.74s	579MB	1268s

Table 4: [9] Index storage and computational overhead for variable size data sets.

We observed an increased overhead in the size of the encrypted index for the ST scheme [9], compared with ours as expected. On average, over all the the data sizes, for different data sets the gain of $\mathsf{S^3E}$

Dataset	#Characters									
	10^2		10^3		10^4		10^5		10^6	
	Storage	Time	Storage	Time	Storage	Time	Storage	Time	Storage	Time
DNA										
win = 2	37KB	0.26s	410KB	49s	4.2MB	3.50s	42MB	29.90s	456MB	728s
win = 2^3	38KB	0.15s	401KB	48s	4.1MB	3.46s	40MB	29.78s	410MB	683s
win = 2^5	29KB	0.11s	371KB	43s	3.7MB	2.34s	37MB	26.12s	373MB	640s
Enron										
win = 2	40KB	0.29s	402KB	50s	4.1MB	3.69s	41MB	30.65s	471MB	702s
win = 2^3	38KB	0.17s	396KB	46s	4MB	3.42s	40MB	29.23s	401MB	678s
win = 2^5	29KB	0.10s	385KB	42s	3.7MB	2.56s	37MB	27.50s	372MB	630s

Table 5: S^3E Index storage and computational overhead for variable size data sets and buckets window size.

over ST approximates a factor of **1.64**. However, for realistic big data streams consisting of 10^6 the gain reaches a factor of **2**. Even though the computation of the encrypted index happens only once, the storage overhead incurred by its encryption is of more crucial importance than its computation time. A limited storage device is not capable of computing the encrypted index if that comes at an increased communication overhead.

We also measured the computation time of the encrypted index in both schemes in tables 4, 5. The S^3E index construction time outperforms ST. Apart from the extra dummy blocks and the increased size of the suffix tree compared with that of a suffix array, the increased computation cost stems from the way ST encrypts the suffix tree.

Query Encryption We run experiments in order to compute the storage overhead during the SrchToken phase. The token consists of a sequence of buckets from a vocabulary: be it characters from emails or characters from DNA sequence. As the query encryption is not affected by the distribution of the underlying characters and for compactness we choose to present results only from the DNA sequence. We also observed tiny differences at the query encryption time and the size of the bucket. We observe a reasonable increase in both the size of the encrypted query (figures 5,6) and its computation time as the query size increases. S^3E outmatches in query computation time due to the way the query is encrypted in ST scheme: Namely for each character 2 PRF, and one block cipher is invoked, while in S^3E only one PRF is invoked. The storage overhead of ST also outgrows faster since the substring is encrypted recursively and not by character. That is, the token: $T[1], T[2], ..., T[m]$ is encrypted as $ct_1 = PRF_1(T[1]), k_1 = PRF_2(T[1]), ct_2 = PRF_1(T[1]T[2]), k_2 = PRF_2(T[1]T[2])$ and so on. Finally the client forwards to the cloud: $\{Enc_{k_i}(ct_i)\}_{i=1}^m$.
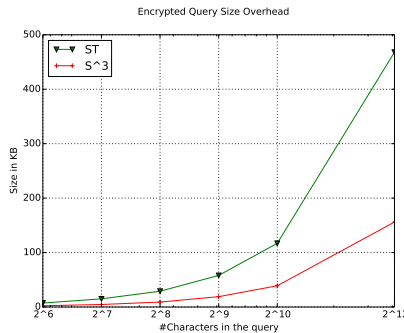

Fig. 5: Token storage overhead for both schemes in DNA streams. As the window size affects very epidemically the encryption of our query we use a fixed window size: win = 2^5

Response Overhead In figures 7, 8 we discern a slight outperformance of S^3E compared with ST in terms of substring response time. For the experiments we computed tokens of various lengths and perform a search on DNA streams of different sizes. The client computes and encrypts the index and uploads it to the cloud. The cloud simulated in the same machine runs the search algorithm, and we
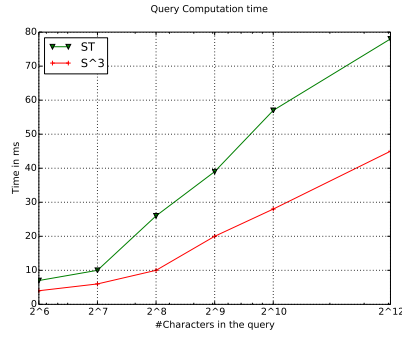
Fig. 6: Token computation time for both schemes in DNA streams. As the window size affects very epidemic the encryption of our query we use a fixed window size: $\mathsf{win} = 2^5$

computed the total search time. We perceived in both schemes, that for considerable smaller than 10^6 elements the running time tends to be independent on the size of the substring token. For a one million data stream there is a notable increased response time compared with the smaller data sets and there is a proportional increment in time with respect to the size of the token. In exact times, $\mathsf{S}^3\mathsf{E}$ surmounts ST [9] for the computation of the response at the cloud side. This outperformance is due to the increased size of the encrypted index in [9] with dummy blocks, based on a suffix tree data structure.


Fig. 7: Response time for $\mathsf{S}^3\mathsf{E}$ scheme in a DNA stream. Time is measured as the average over different bucket sizes: $2, 2^3, 2^5$


Fig. 8: Response time for ST [9] scheme

9 Conclusion

We designed and analyzed a substring searchable symmetric encryption protocol $\mathsf{S^3E}$, which achieves better storage performance than state of the art work [9]. The idea of our protocol is to leverage the self-indexing mechanism of $\mathtt{FM}$ index, which stores only n integer positions of its substrings. Our protocol is provably secure under the real-ideal paradigm. We also implemented our protocol and compared it with the state of the art work [9], showing its notable performance improvement in terms of storage overhead and computation time.

Bibliography

[1] M. I. Abouelhoda, E. Ohlebusch, and S. Kurtz. Optimal exact string matching based on suffix arrays. In *In Proceedings of the Ninth International Symposium on String Processing and Information Retrieval. Springer-Verlag, Lecture Notes in Computer Science*, 2002.

[2] J.-P. Aumasson, S. Neves, Z. Wilcox-O'Hearn, and C. Winnerlein. Blake2: Simpler, smaller, fast as md5. In *Proceedings of the 11th International Conference on Applied Cryptography and Network Security*, ACNS'13, pages 119–135, Berlin, Heidelberg, 2013. Springer-Verlag.

[3] Bitglass. `http://www.bitglass.com/company/news/press_releases/patentedencryption`.

[4] R. S. Boyer and J. S. Moore. A fast string searching algorithm. *Commun. ACM*, 20(10):762–772, Oct. 1977.

[5] M. Burrows and D. J. Wheeler. A block-sorting lossless data compression algorithm. Technical report, 1994.

[6] D. Cash, S. Jarecki, C. S. Jutla, H. Krawczyk, M. Rosu, and M. Steiner. Highly-scalable searchable symmetric encryption with support for boolean queries. In *Advances in Cryptology - CRYPTO 2013 - 33rd Annual Cryptology Conference, Santa Barbara, CA, USA, August 18-22, 2013. Proceedings, Part I*, pages 353–373, 2013.

[7] Y.-C. Chang and M. Mitzenmacher. Privacy preserving keyword searches on remote encrypted data. In *Proceedings of the Third International Conference on Applied Cryptography and Network Security*, ACNS'05, pages 442–455, Berlin, Heidelberg, 2005. Springer-Verlag.

[8] M. Chase and S. Kamara. Structured encryption and controlled disclosure. In *Advances in Cryptology - ASIACRYPT 2010 - 16th International Conference on the Theory and Application of Cryptology and Information Security, Singapore, December 5-9, 2010. Proceedings*, pages 577–594, 2010.

[9] M. Chase and E. Shen. Substring-searchable symmetric encryption. *PoPETs*, 2015(2):263–281, 2015.

[10] Ciphercloud. `http://www.ciphercloud.com/technologies/encryption/`.

[11] Ciphercloud. Q2-global-cloud-data-security-report. `http://pages.ciphercloud.com/rs/830-ILB-474/images/Q2-Global-Cloud-Data-Security-Report.pdf`.

[12] R. Curtmola, J. Garay, S. Kamara, and R. Ostrovsky. Searchable symmetric encryption: Improved definitions and efficient constructions. In *Proceedings of the 13th ACM Conference on Computer and Communications Security*, CCS '06, pages 79–88, New York, NY, USA, 2006. ACM.

[13] S. Faber, S. Jarecki, H. Krawczyk, Q. Nguyen, M. Rosu, and M. Steiner. Rich queries on encrypted data: Beyond exact matches. In *Computer Security - ESORICS 2015 - 20th European Symposium on Research in Computer Security, Vienna, Austria, September 21-25, 2015, Proceedings, Part II*, pages 123–145, 2015.

[14] S. Faust, C. Hazay, and D. Venturi. Outsourced pattern matching. In F. Fomin, R. Freivalds, M. Kwiatkowska, and D. Peleg, editors, *Automata, Languages, and Programming*, volume 7966 of *Lecture Notes in Computer Science*, pages 545–556. Springer Berlin Heidelberg, 2013.

[15] P. Ferragina and G. Manzini. Opportunistic data structures with applications. In *Proceedings of the 41st Annual Symposium on Foundations of Computer Science*, FOCS '00, pages 390–, Washington, DC, USA, 2000. IEEE Computer Society.

[16] C. Gentry, K. A. Goldman, S. Halevi, C. S. Jutla, M. Raykova, and D. Wichs. Optimizing ORAM and using it efficiently for secure computation. In *Privacy Enhancing Technologies - 13th International Symposium, PETS 2013, Bloomington, IN, USA, July 10-12, 2013. Proceedings*, pages 1–18, 2013.

[17] O. Goldreich, S. Goldwasser, and S. Micali. How to construct random functions. *J. ACM*, 33(4):792–807, Aug. 1986.

[18] G. H. Gonnet, R. A. Baeza-Yates, and T. Snider. Information retrieval. chapter New Indices for Text: PAT Trees and PAT Arrays, pages 66–82. Prentice-Hall, Inc., Upper Saddle River, NJ, USA, 1992.

[19] Hitachi. Searchable encryption: A technology for supporting secure application. `http://www.hitachi.com/rd/portal/contents/story/searchable_encryption/index.html`.

[20] J. Kärkkäinen and P. Sanders. Simple linear work suffix array construction. In *Automata, Languages and Programming, 30th International Colloquium, ICALP 2003, Eindhoven, The Netherlands, June 30 - July 4, 2003. Proceedings*, pages 943–955, 2003.

[21] R. M. Karp and M. O. Rabin. Efficient randomized pattern-matching algorithms. *IBM J. Res. Dev.*, 31(2):249–260, Mar. 1987.

[22] D. E. Knuth, J. H. M. Jr., and V. R. Pratt. Fast pattern matching in strings. *SIAM J. Comput.*, 6(2):323–350, 1977.

[23] Y. Lindell. How to simulate it - A tutorial on the simulation proof technique. *IACR Cryptology ePrint Archive*, 2016:46, 2016.

[24] U. Manber and G. Myers. Suffix arrays: A new method for on-line string searches. In *Proceedings of the First Annual ACM-SIAM Symposium on Discrete Algorithms*, SODA '90, pages 319–327, Philadelphia, PA, USA, 1990. Society for Industrial and Applied Mathematics.

[25] E. M. McCreight. A space-economical suffix tree construction algorithm. *J. ACM*, 23(2):262–272, Apr. 1976.

[26] Mitsubishielectric. http://www.mitsubishielectric.com/news/2013/0703.html.

[27] T. Moataz and E.-O. Blass. Oblivious substring search with updates. Cryptology ePrint Archive, Report 2015/722, 2015. http://eprint.iacr.org/2015/722.

[28] O. Sanders, C. Onete, and P.-A. Fouque. Pattern matching on encrypted streams: Applications to dpi and searches on genomic data. Cryptology ePrint Archive, Report 2017/148, 2017. http://eprint.iacr.org/2017/148.

[29] Skyhighnetworks. https://www.skyhighnetworks.com/cloud-encryption/.

[30] E. Ukkonen. On-line construction of suffix trees. *Algorithmica*, 14(3):249–260.

[31] Virtualworks. http://www.virtualworks.com/viaworks-enterprise-search/.

[32] P. Weiner. Linear pattern matching algorithms. In *Switching and Automata Theory, 1973. SWAT '08. IEEE Conference Record of 14th Annual Symposium on*, pages 1–11, Oct 1973.

A   Cryptographic Primitives

A.1   Pseudorandom functions (PRF)

Let the family of all functions in the universe from a domain X to a range Y to be $Func[X, Y]$. A truly random function $f \xleftarrow{\$} Func[X, Y]$ is chosen randomly from the set of $Func$. The set of all these functions is $|Y|^{|X|}$ (gigantic number). It is true that for any random function f with range size L chosen randomly from $Func[X, Y]$, $\Pr[f(x) = y] = 2^{-L}$. The randomness is not parametrized neither by the size of X and Y nor by the size of the domain. We define a pseudorandom function $f_k : X \to Y$ as a function from the set of all functions from X to Y as soon as a particular key k is fixed.

Definition 4. *Let Func=$\{F : X \to Y\}$ be a function family for all functions F that map elements from the domain X to the range R. Then a $PRF = \{f_k : X' \to Y'\} \subseteq Func$ for $k \xleftarrow{\$} K$, where K is the key space.*

The security of a PRF is modeled with a game which is known as *real or random* security game[17]. Intuitively, an adversary $\mathcal{A}$ is given access to an oracle that on input x from a domain X, flips a coin $b \xleftarrow{\$} \{0, 1\}$ and if $b = 0$ then it outputs $y = f(x)$, for $f \in Func[X, Y]$, otherwise it outputs $y = f_k(x)$. $\mathcal{A}$ issues queries to the oracle polynomially many times on input of the security parameter λ. Finally $\mathcal{A}$ outputs a guess b' for the bit b.

The advantage of a probabilistic polynomial time algorithm $\mathcal{A}$ in the PRF game is

$$\mathbf{Adv}_{\mathcal{A}}^{PRF} = \Pr[b \xleftarrow{\$} \{0, 1\}; b' \leftarrow \mathcal{A}(y) : b' = b]$$

Definition 5. *A PRF is computationally secure if all probabilistic polynomially time algorithms $\mathcal{A}$ have advantage in the PRF game: $\frac{1}{2} + \epsilon(\lambda)$, for a negligible function ϵ on the security parameter λ.*

A.2   Pseudorandom permutations (PRP)

A permutation is a bijective function where the domain and the range are equal. Similarly with the random functions, let $Perm[X]$ to be the set of all permutations for the domain X. Then a pseudorandom permutation (PRP) is a randomly chosen permutation from the set $Perm[X]$, keyed under a secret key k.

The advantage of a probabilistic polynomial time algorithm $\mathcal{A}$ in the PRP game is

$$\mathbf{Adv}_{\mathcal{A}}^{PRP} = \Pr[b \xleftarrow{\$} \{0, 1\}; b' \leftarrow \mathcal{A}(y) : b' = b]$$

Definition 6. *A PRP is computationally secure if all probabilistic polynomially time algorithms $\mathcal{A}$ have advantage in the PRP game: $\frac{1}{2} + \epsilon(\lambda)$, for a negligible function ϵ on the security parameter λ.*

A.3 Symmetric Key Encryption

A symmetric key encryption scheme $\mathsf{SKE} = \{\mathsf{Gen}, \mathsf{Enc}, \mathsf{Dec}\}$ consists of three algorithms. Gen takes as input a security parameter λ and outputs the secret key sk. The probabilistic encryption algorithm E takes as input the secret key sk and a plaintext x form the plaintext space $\mathcal{P}$ and outputs the ciphertext c. The decryption algorithm $\mathsf{SKE.Dec}$ takes as input a ciphertect form the ciphertext space $\mathcal{C}$ and the secret decryption key sk and outputs the plaintext $x \in \mathcal{P}$. Correctness follows $\iff \forall \mathsf{sk} \leftarrow \mathsf{Gen}(1^\lambda), \mathsf{SKE.Dec}((\mathsf{E}(\mathsf{sk}, x))) = x, \forall x \in \mathcal{P}$. Security is modeled with the standard game based indistinguishability experiment for polynomial probabilistic time adversary $\mathcal{A}$.

$\underline{\mathsf{PrivK}_{\mathcal{A},\mathsf{SKE}}(\lambda):}$

- $\mathcal{A}$ has access to the security parameter 1^λ.
- A key $\mathsf{sk} \leftarrow \mathsf{Gen}(1^\lambda)$ is generated and $\mathcal{A}$ can learn encryptions of x of its choice $x \in \mathcal{S} \subset \mathcal{P}$.
- Eventually $\mathcal{A}$ outputs x_0, x_1 where $|x_0| = |x_1|$. $b \xleftarrow{\$}$ and $\mathsf{E}(x_b, \mathsf{sk})$ is returned to $\mathcal{A}$.
- $\mathcal{A}$ outputs its guess for b, b'.

If $b' = b$ $\mathcal{A}$ succeeds and the experiment $\mathsf{PrivK}_{\mathcal{A},\mathsf{SKE}}(\lambda) = 1$.

Definition 7. *A symmetric encryption scheme SEK has indistinguishable encryptions if the probabilities* $\Pr[\mathsf{PrivK}_{\mathcal{A},\mathsf{SKE}}(\lambda) = 1] \leqslant \frac{1}{2} + \mathsf{neg}(\lambda)$.